

Using Local and Global Knowledge in Wireless Sensor Networks

**A thesis submitted in partial fulfilment
of the requirement for the degree of Doctor of Philosophy**

Christopher Gwilliams

14th May 2015

**Cardiff University
School of Computer Science & Informatics**

For Charlotte Luke,

**Thank you for helping me finish this chapter
and I am so glad to be writing the next one with you**

Love always,

Chris Gwilliams

Declaration

This work has not previously been accepted in substance for any degree and is not concurrently submitted in candidature for any degree.

Signed (candidate)

Date

Statement 1

This thesis is being submitted in partial fulfilment of the requirements for the degree of PhD.

Signed (candidate)

Date

Statement 2

This thesis is the result of my own independent work/investigation, except where otherwise stated. Other sources are acknowledged by explicit references.

Signed (candidate)

Date

Statement 3

I hereby give consent for my thesis, if accepted, to be available for photocopying and for inter-library loan, and for the title and summary to be made available to outside organisations.

Signed (candidate)

Date

Copyright © 2014 Christopher Gwilliams.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.2 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. A copy of the license is included in the section entitled “GNU Free Documentation License”.

A copy of this document in various transparent and opaque machine-readable formats and related software is available at <http://github.com/encima>.

Abstract

Wireless sensor networks (WSNs) have advanced rapidly in recent years and the volume of raw data received at an endpoint can be huge. We believe that the use of local knowledge, acquired from sources such as the surrounding environment, users and previously sensed data, can improve the efficiency of a WSN and automate the classification of sensed data. We define local knowledge as knowledge about an area that has been gained through experience or experimentation. With this in mind, we have developed a three-tiered architecture for WSNs that uses differing knowledge-processing capabilities at each tier, called the Knowledge-based Hierarchical Architecture for Sensing (K-HAS). A novel aligning ontology has been created to support K-HAS, joining widely used, domain-specific ontologies from the sensing and observation domains. We have shown that, as knowledge-processing capabilities are pushed further out into the network, the profit - defined as the value of sensed data - is increased; where the profit is defined as the value of the sensed data received by the end user.

Collaborating with Cardiff University School of Biosciences, we have deployed a variation of K-HAS in the Malaysian rainforest to capture images of endangered wildlife, as well as to automate the collection and classification of these images. Technological limitations prevented a complete implementation of K-HAS and an amalgamation of tiers was made to create the Local knowledge Ontology-based Remote-sensing Informatics System (LORIS). A two week deployment in Malaysia suggested that the architecture was viable and that, even using local knowledge at the endpoint of a WSN, improved the efficiency of the network. A simulation was implemented to model K-

HAS and this indicated that the network became more efficient as knowledge was pushed further out towards the edge, by allowing nodes to prioritise sensed data based on inferences about its content.

Acknowledgements

My eternal gratitude goes to Alun Preece and Alex Hardisty for their guidance, assistance, support, enthusiasm and faith in me throughout the research and writing of this thesis. Their encouragement of testing new ideas, questioning everything and ensuring I step out of the box that it seems all too easy to stay inside. I am also grateful for the skills they have pushed me to gain, from teaching to presenting to professional writing. None of this work would have been possible without you both.

Benoit Goosens, Mike Bruford, Danica Stark and all of those involved with Danau Girang Field Centre have been fountains of knowledge, support and inspiration. At the very least, I thank you for your hospitality each year and to the staff members that put up with a computer scientist that is clearly not well adapted to life in the rainforest with limited connectivity to the outside world.

I am grateful to the School of Computer Science & Informatics, staff and students have provided methods of growth, new skills and support throughout my time at Cardiff University. My thanks go to Matt Williams, for challenging my work from the early stages right through to completion and providing relief through films, and Will Webberley, a friend gained through the PhD that provided support, both professionally and personally, and I do not believe I could have completed this work without him. Special mentions must go to Martin Chorley, Jonathan Quinn, Gualtiero Colombo, Konrad Borowiecki, Diego Pizzocaro, Przemyslaw Woznowski, Rich Coombs, Liam Turner, Mark Greenwood and Matthew John for being there in the pub, as well as in academia.

I would like to dedicate this work to my father, who passed away in 2009, there is no way I would be anywhere close to this point without his input, teaching and encouragement pushing me to go further than I ever imagined I could go. Words cannot express the love and gratitude I have for Robert and Gina Gwilliams, two role models that truly set the bar high.

Contents

Abstract	iv
Acknowledgements	vi
Contents	viii
List of Publications	xv
List of Figures	xvi
List of Tables	xix
Listings	xx
List of Acronyms	xxi
1 Introduction	1
1.1 Motivation	2
1.2 Research Contributions	4
1.3 Thesis Structure	6

2 Background	7
2.1 Wireless Sensor Network Issues	8
2.1.1 Fault Tolerance	8
2.1.2 Hardware Constraints	8
2.1.3 Energy Constraints	9
2.1.4 Transmission Medium	10
2.1.5 Environment	11
2.2 Routing Protocols	12
2.2.1 Flat	12
2.2.2 Data-centric	13
2.2.3 Hierarchical	16
2.2.4 Location-based	17
2.2.5 Conclusion	19
2.3 Sensor Middleware	19
2.3.1 Issues	20
2.3.2 Existing Middleware	22
2.4 Biodiversity and Environmental Monitoring Sensor Networks	30
2.5 Conclusion	33
3 Exploratory Experiments in a Rainforest Environment	35
3.1 Danau Girang Field Centre	36
3.2 Hardware	40

3.2.1	Pandaboard	42
3.2.2	IGEP v2	42
3.2.3	Wasp mote	42
3.2.4	Discussion	43
3.3	Transmission Media	44
3.3.1	Wi-Fi	44
3.3.2	DigiMesh	47
3.3.3	Adapting and Optimising for Harsh Environments	48
3.4	Software	50
3.4.1	Triton	50
3.5	Conclusion	56
4	Wireless Sensor Network Architecture Design	57
4.1	The Local Knowledge Problem	58
4.2	K-HAS	60
4.2.1	Data Collection	61
4.2.2	Data Processing	63
4.2.3	Data Aggregation	65
4.2.4	Sequence Diagram	68
4.3	Technological Components	70
4.3.1	Data Standard	70
4.3.2	Middleware	75

4.3.3	Knowledge Capture	76
4.3.4	Knowledge Base	78
4.3.5	Routing Protocol	80
4.4	Walkthrough	82
4.4.1	Scenario	83
4.4.2	Data Collection	84
4.4.3	Data Processing	86
4.4.4	Data Aggregation	88
4.5	Conclusion	90
5	An Ontology for Knowledge-Based Wireless Sensor Networks	92
5.1	Background	92
5.1.1	Observation-Centric Ontologies	94
5.1.2	Sensor-Centric Ontologies	98
5.1.3	Combined Sensor and Observation Ontologies	99
5.1.4	Findings	101
5.2	Method	102
5.2.1	Identify Purpose	103
5.2.2	Build the Ontology	104
5.2.3	Evaluation	105
5.2.4	Documentation	106
5.3	Results	106

5.3.1	Identify Purpose	106
5.3.2	Build the Ontology - Capture	107
5.3.3	Build the Ontology - Code	109
5.3.4	Build the Ontology - Integrate Existing Ontologies	109
5.3.5	Extend the Ontology	111
5.4	Evaluation	113
5.4.1	Using the Ontology in an Application	113
5.4.2	Competency Questions	114
5.4.3	Comparing the Ontology with a Collection of Documents . . .	116
5.5	Conclusion	117
6	Deployment in the Field	120
6.1	Deployment Architecture	121
6.1.1	Data Collection	122
6.1.2	Data Processing and Aggregation	124
6.2	Deployment	125
6.2.1	Direct Connection	126
6.2.2	1-Hop Network	127
6.3	Results	128
6.4	Rules	130
6.4.1	Data Processing and Data Aggregation	130
6.4.2	Existing Data	132
6.5	Conclusion	136

7 Simulation Experiments and Evaluation of the Architecture	138
7.1 Simulation Environment	139
7.2 Simulation Scenarios	140
7.3 Network Setup	144
7.3.1 Parameters	144
7.3.2 Routing Protocol	148
7.3.3 Observation Transmission	150
7.3.4 Processing	151
7.4 Results	153
7.4.1 Normal Network	155
7.4.2 Saturated Network	158
7.5 Conclusion	162
8 Conclusion and Future Work	165
8.1 Summary of Contributions	166
8.1.1 K-HAS	167
8.1.2 Ontology	167
8.1.3 LORIS	167
8.1.4 Simulations	168
8.2 Future Work	168
Bibliography	172
Appendices	188

A Example Darwin Core Archive	189
B Darwin-SW Occurrence Example	191
C K-HAS Ontology	193
D Example Drools Rules Used in LORIS	196
E Interview Transcript: Gill Bolongan	197
F Qualitative Analysis Extract	200
G Class Diagram for Repast Simulation	201
H Repast Parameters File Example	203

List of Publications

The work introduced in this thesis is based on the following publications.

- [51] - Gwilliams, C., Preece, A., Hardisty, A., Goossens, B., & Ambu, L. N. (2012). Local and Global Knowledge to Improve the Quality of Sensed Data. International Journal of Digital Information and Wireless Communications (IJDWC), 2(2), 164–180.
- [52] - Gwilliams, C., Preece, A. D., & Hardisty, A. R. (2012). Poster: using local and global knowledge in wireless sensor networks. In Proceedings of the 10th international conference on Mobile systems, applications, and services (pp. 515–516).
- [49] - Gwilliams, C., Preece, A., Hardisty, A., Goossens, B., & Ambu, L. N. (2012). K-HAS: An Architecture for Using Local and Global Knowledge in Wireless Sensor Networks. In The International Conference on Informatics and Applications (ICIA2012) (pp. 331–341).

List of Figures

1.1	Map of Danau Girang in relation to Sabah Malaysia	3
2.1	GSN Architecture [12]	23
2.2	FACTS Architecture [110]	25
3.1	The Main Buidling at DGFC	36
3.2	Reconyx HC500 Camera	37
3.3	Sensor Node Hardware	41
3.4	Signal-to-Noise Ratio for Wi-Fi	46
3.5	Pelican Waterproof Case	49
3.6	Sequence of 3 Images Captured on a Reconyx Trigger	52
3.7	Resulting Processed Image from Original Set of 3	53
4.1	Overview of the K-HAS Architecture	58
4.2	Sequence Diagram for K-HAS upon Capture of an Observation	69
4.3	A Darwin Core Star Archive [119]	72
4.4	EXIF Tags from an Reconyx Camera	74

4.5	Knowledge Hierarchy for K-HAS	77
4.6	Web Interface for Observations	78
4.7	Web Interface for Classified Sensed Data	78
4.8	Web Interface for Drools Operations	80
4.9	Activity Diagram for DC Nodes during the Discovery Phase	82
4.10	Section of Proposed Topology for K-HAS in Danau Girang	84
4.11	Clouded leopard Image Capture	85
4.12	Processed Image of clouded leopard	85
5.1	The Core of an OBOE Observation [71]	94
5.2	The OGC O&M Ontology [88]	96
5.3	Darwin-SW Entity-Relationship Model [19]	97
5.4	SSN Ontology	100
5.5	K-HAS Base Ontology	109
5.6	K-HAS Ontology	112
6.1	Buckeye X7R and Battery Pack	123
6.2	LORIS Web Interface	125
6.3	Camera Locations Around Danau Girang	126
6.4	First Image of Lizard by Buckeye 3	128
6.5	Second Image of Lizard by Buckeye 3	129
7.1	Simulation Example in RePast Simulator	143

7.2	Mean Transmission Times for Delivered Observations for a Normal Network	157
7.3	Delivered vs Undelivered Observations for a Normal Network	158
7.4	Interesting vs Empty Delivered Observations for a Normal Network	158
7.5	Mean Transmission Times for Delivered Observations for a Saturated Network	160
7.6	Delivered vs Undelivered Observations for a Normal Network	161
7.7	Interesting vs Empty Delivered Observations for a Normal Network	162

List of Tables

3.1	Classification Results	54
3.2	Random Classifier Results	55
5.1	Competency Questions	107
5.2	K-HAS Concepts	108
5.3	Competency Questions	115
5.4	Evaluation of K-HAS against Darwin Core Occurrence Terms	117
7.1	Mapping of K-HAS to Simulation Terminology	143
7.2	Parameters Used to Configure Repast Simulations	145
7.3	All Observations Captured and Delivered	154

Listings

2.1 Example Virtual Sensor	24
2.2 Coverage Algorithm in FACTS Rules	26
4.1 Pseudocode DC Node Rules	63
4.2 Darwin Core Ecological Metadata File Fragment	71
4.3 Darwin Core Descriptor File	72
4.4 Darwin Core Occurrence Data	73
5.1 An Observation of a Vehicle encoded in O&M	95
5.2 Darwin-SW Representation of an Identification	98
6.1 Pseudocode of a Drools Rules File	131
6.2 Example Rule created from Existing Data	135
7.1 Processing Function for DA and DP Nodes	152
7.2 Processing Functions for Nodes	152
A.1 EML File	189
A.2 Metadata File	190
A.3 Example Set File (set.csv)	190
A.4 File Describing Image Locations (images.csv)	190
B.1 Darwin-SW Representation of a Living Specimen	191
C.1 K-HAS Ontology Source Code	193
D.1 Example Drools File	196
E.1 Sample Interview Transcript	197
H.1 Example params.xml file for a Repast simulation	203

List of Acronyms

API Application Programming Interface

BLE Bluetooth Low Energy

DA Data Aggregation

Darwin-SW Darwin Core Semantic Web

DC Data Collection

DGFC Danau Girang Field Centre

DP Data Processing

DwC Darwin Core

DwC-A Darwin Core Archive

EML Ecological Metadata Language

EPO Extension Plug-in Ontologies

EXIF EXchangeable Image Format

FN False Negative

FP False Positive

GBIF Global Biodiversity Information Facility

GEAR Geographical Energy Aware Routing

GSN Global Sensor Networks

HK High Knowledge

IFTTT If This Then That

INSIGHT INternet-Sensor InteGration for HabitaT monitoring

IoT Internet of Things

IR Infrared

JVM Java Virtual Machine

K-HAS Knowledge-based Hierarchical Architecture for Sensing

LKWS Lower Kinabatangan Wildlife Sanctuary

LOS Line of Sight

LORIS Local knowledge Ontology-based Remote-sensing Informatics System

MCFA Minimum Cost Forwarding Algorithm

MK Minimal Knowledge

NK No Knowledge

O&M Observations & Measurements

OBOE Extensible Observation Ontology

OGC Open Geospatial Consortium

OpenCV Open Computer Vision

OWL Web Ontology Language

QoS Quality of Service

ROI Region of Interest

RDL Ruleset Definition Language

SBC Single Board Compuetr

SDO Sensor Data Ontology

SensorML Sensor Model Language

SHO Sensor Hierarchy Ontology

SOS Sensor Observation Service

SPIN Sensor Protocol for Information via Negotiation

SPS Sensor Planning Ontology

SSN Semantic Sensor Network

SUMO Suggested Upper Merged Ontology

SWE Sensor Web Enablement

TEEN Threshold sensitive Energy Efficient sensor Network

TN True Negative

TOVE Toronto Virtual Enterprise

TP True Positive

Wi-Fi Wireless Fidelity

WSN Wireless Sensor Network

Chapter 1

Introduction

A wireless sensor network (WSN) consists of a collection of nodes with sensing and, typically, wireless communication capabilities. These sensing nodes can be complex and powerful devices with the ability to sense multiple phenomena simultaneously [77, 81, 114], or they can be simple motes with limited processing power [63, 76, 106]. A node has one or more sensors attached to it, with a wireless radio that is used to transmit the sensed data to an endpoint, where an endpoint is the base station of the network. Most WSNs use a single base station, providing a single endpoint, but there are networks with multiple endpoints; either used for different types of sensed data or to ensure that the load of the network is spread out.

Upon deployment, these nodes use their wireless capabilities to form communication links with their neighbours, where a neighbour is any node that is within transmission range. The way that nodes discover, and communicate with their neighbours is defined by their routing protocol. Routing protocols vary based on the purpose of the WSN, the requirements of data transmission as well as the characteristics of the nodes. Communication between nodes is expensive and drains the available power faster than any other action that a node performs [92]. For example, if a WSN is deployed in a building with constant power availability, then the routing protocol would not need to be modified to ensure nodes sleep to conserve battery or disable their radios for a period of time. However, not every WSN has unlimited resources at their disposal and these protocols, as well as the underlying structure of the network, are used to ensure the network is able to perform well for as long as possible.

Each WSN is different and each will have different constraints, a WSN that monitors traffic along a busy road may experience memory limitations, whereas a WSN that is deployed in the middle of a desert may experience power issues. Typically, however, all WSNs do the same thing: sense one or more characteristics of their environment and forward that data on to a specified endpoint; sometimes called the base station.

1.1 Motivation

Throughout this thesis, we focus on a scenario motivated by our collaboration with Cardiff University School of Biosciences, who run a research centre in the Malaysian rainforest, in Sabah, known as Danau Girang Field Centre (DGFC) [2]. Located on the banks of the Kinabatangan river, DGFC is shown in Figure 1.1 and has been running for more than six years and holds Masters, PhD and Undergraduate students from around the world, studying the ecology and biodiversity of the unique region.

The rainforest surrounding the Kinabatangan river is unique because the area was heavily logged until the late 1970s and the river now serves as a corridor, between large palm oil plantations, connecting two separate rainforest lots. The area is now secondary rainforest (rainforest that has grown since being destroyed) and is experiencing a large variety of wildlife using the area as a habitat, or as a path. Some of this wildlife is unique to this area of the world and DGFC has had sightings of animals that have not been seen in many years [17].

There is a variety of research projects currently underway in the field centre, looking into fish population, crocodile attacks, hornbill habitats or the movement patterns of small mammals. One project that has been running almost since DGFC opened, is the *corridor monitoring programme*, a programme that consists of dozens of wildlife cameras deployed in various areas around DGFC and a set of images are taken whenever an animal triggers a break in their infrared (IR) sensor.

The Kinabatangan is a protected wildlife reserve, with thick and humid forest, making

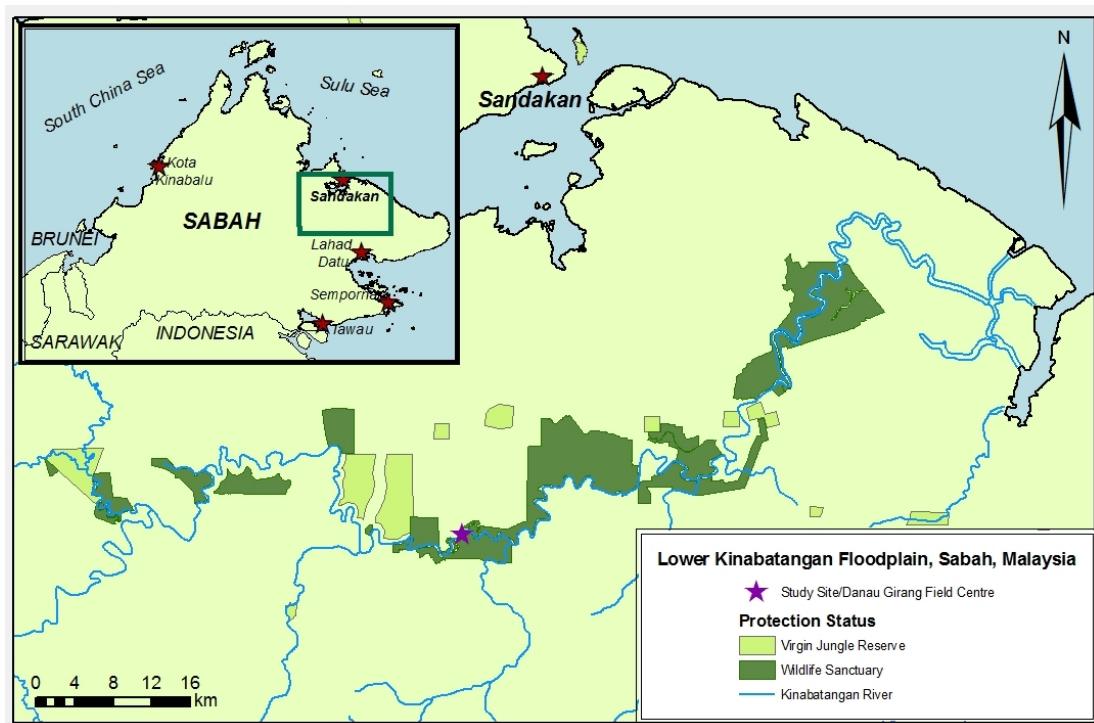


Figure 1.1: Map of Danau Girang in relation to Sabah Malaysia

it very difficult to walk through and even more difficult for hardware to survive the conditions. Cameras are placed along the river and up to 1km into the forest, capturing images when triggered and saving them onto SD cards. These SD cards are collected and stored at the field centre, where the images are manually collated and processed. The cameras are designed to have a battery life of three months but, due to the humidity, a battery life of three weeks is more realistic. In 2010, twenty cameras were deployed and half of them were inspected every two weeks, on a rotating basis. In that time, each camera can record more than a thousand pictures and the dynamic nature of the rainforest, such as the sun through leaves, falling trees and reflections in the water can cause the camera to trigger when an animal has not walked past. The events are known as *false triggers*, and they can make up to 70% of the images on an SD card and each of these must be manually processed.

Each trigger of the camera results in a set of three images being captured, this number has been chosen by the researchers at DGFC after some experimentation. Three images

allows for more than one image of slow-moving animals to be captured and, typically, allows for the middle image in the sequence to capture the body of fast-moving animals; with the first and last capturing the head and rear respectively.

We have used this scenario to test our hypothesis and implement a WSN that automates the collection, transmission, processing and storage of images, using local knowledge to classify the data and prioritise the flow of information through the network, making more efficient use of the limited power and bandwidth available.

1.2 Research Contributions

Here we outline the main research contributions explained in this thesis. We believe that two main contributions have been made and our experimental results serve to support these contributions. Our primary contribution is that we propose a **novel tiered network architecture where each subsequent tier possesses increased knowledge-processing ability**, K-HAS, that utilises the local knowledge of its surrounding environment, users and previously sensed data to process observations within the network and prioritise the data according to the inferred classification.

Knowledge is pushed out to the edge of the network to allow the nodes that capture observations to prioritise the data based on its content. The knowledge processing capabilities increase with each tier as the data moves toward the centre of the network, allowing more detailed inferences to be made and data to be given a higher priority. This smart utilisation of bandwidth allows data to be delivered in an order that is more useful than chronological, delivering the most valuable observations first and using human feedback to learn how important these observations were.

In addition, K-HAS uses a feedback loop to dynamically update the knowledge base on every node throughout the deployment so it is able to react to changes within the data recorded, and the network, in near real-time.

Experimental results from the Malaysian rainforest showed that current long-range wireless sensor technology is not yet at the point where nodes could be deployed in a harsh environment and left without human intervention for months at a time. A modified K-HAS system LORIS (Local knowledge Remote-sensing Ontology-based Informatics System) combined tiers that use knowledge-processing at the centre of the network, with commercially available, wireless cameras. This showed that knowledge-processing automates the handling of sensed data when it is received and can be used to infer patterns for future observations. LORIS was designed because we could not find robust, open sensing nodes that can send data over long distances and handle high humidity and extreme temperatures. Because of this, we used off-the-shelf (OTS) wireless cameras that ran a closed system, while providing some access through proprietary software. We then combined two of K-HAS' tiers so that processing only took place at the base station, matching a more traditional WSN topology; while still making use of local knowledge gained from previously sensed data.

We model the ideal implementation of K-HAS in a simulation environment, along with variations on the knowledge processing capabilities of each node. We use these simulations to show that local and global knowledge can prioritise sensed data effectively and that this prioritisation increases the efficiency of the network by delivering data based on its importance, rather than chronologically.

Our second contribution is an **ontology that combines, and extends, ontologies used in multiple domains to create an extension ontology for WSNs and ecological observations**. To formally define the structure of K-HAS and the data standard used, we developed an ontology that combined existing ontologies for sensor networks and ecological observations. This ontology can be used as a whole or in parts for WSNs that use some, or all, of K-HAS' architecture. On top of this, we extended these ontologies by adding specific terms relevant to the different node types and users within the K-HAS architecture.

We apply this ontology to our proposed architecture, and simulations, in order to define

a data standard for sensed data sent within the network, as well as terms to use in observations as well as identify the roles of both humans and sensors within the network.

1.3 Thesis Structure

The rest of this thesis is structured as follows. Chapter 2 provides background on wireless sensor networks and the use of knowledge-based technologies in that context. Chapter 3 explains technical decisions we made and the findings when running experiments in the Malaysian rainforest. Chapter 4 introduces the K-HAS architecture we have proposed and explains the purpose of each tier. Chapter 5 details the ontology we have proposed to support K-HAS and shows how current ontologies do not sufficiently cover all of the concepts involved with a *scientific observation*. Chapter 6 details our implementation in the Malaysian Rainforest and the changes we had to make in order for it to be feasible. Chapter 7 describes how we modelled K-HAS, and other scenarios, in a simulation environment. The results are then explained and we explain how different scenarios are suited to WSNs, in a multitude of environments, with different requirements. Chapter 8 then concludes this thesis and summarises our contributions and findings, as well as highlighting work that could be undertaken to take this project further.

Chapter 2

Background

Using knowledge in a WSN is related to existing research into sensor networks that utilise context-awareness in order to improve their effectiveness or adapt their sampling rate. An example of a context-aware sensor is an accelerometer attached to a node that is able to determine what the readings of the accelerometer indicate. For example, a smart phone with an accelerometer may use context-awareness to determine if the user is running, walking or going upstairs. Some of these actions may be more important than others and, thus, they can be prioritised. However, researching context-aware WSNs alone would limit our knowledge of WSNs in general and affect decisions we make when designing our own architecture. There is research that is valuable for all types of sensor network, such as: hardware design, routing protocol, transmission medium choice or middleware used.

This chapter is split into the following sections. Section 2.1 outlines the issues surrounding WSN design and deployment. Section 2.2 details relevant existing routing protocols for sensor networks that are used for a number of different purposes, from extending network lifetime by scheduling sleep patterns to storing sensed data on nodes and responding to queries. Section 2.3 highlights commonly-used sensor middleware. Section 2.4 shows some examples of existing WSNs that are related to our motivating scenario and Section 2.5 summarises the results of our research.

2.1 Wireless Sensor Network Issues

WSNs have been used in a number of domains, for a range of different purposes, from habitat monitoring [107] to military purposes [86] and healthcare [85]. While these applications are different, the technology behind each is similar. Each requires the use of nodes with sensors attached and each node requires a power source and storage devices.

According to [15], there are at least eight factors that affect the design of sensor networks, but we focus on a subset that are the most relevant to our research problem. Those we have given less consideration to are: production costs, scalability and topology. While these are important, production costs are not a concern for us in the research stage and scalability has been considered when studying existing networks. We have considered the following points in greater detail:

2.1.1 Fault Tolerance

WSNs typically contain a large number of nodes and any node can fail for various reasons, from a lack of power, filling its storage capacity, to factors in the environment causing the hardware to fail. While the hardware architecture of sensor nodes is typically similar, the variation between each deployment means that the device itself must be adapted to its environment. For example, [73] used a custom protective casing for their nodes so that they were able to survive being in the open while ensuring that the transmission range was not affected.

2.1.2 Hardware Constraints

A sensor node typically consists of: a platform that contains the memory and processing power, a sensor (or sensors) and a transceiver that uses a wireless standard, such as Wi-Fi or Zigbee. Cost and size are the most significant considerations when

designing a WSN. According to [58], the expectation of a sensor node is a matchbox-sized form factor. However, ten years ago, there was much research focussed on *smart dust* [62]. Smart dust is small, inexpensive, disposable nodes that can transmit until their power reserve is depleted, [16] highlights that it is a requirement for the nodes to cost less than USD10. In [32], a decade on from the first WSN papers, smart dust has not been realised and the focus has instead been on larger, more powerful nodes that have reduced in cost and grown in power. The Gartner Hype Cycle for 2013 [43] showed that smart dust is still in early innovation stages and may not be fully commercialised for another ten years. To counter this, research has been focussed on using software solutions to maximise the battery life in these more powerful, more expensive nodes, accompanied by the use of renewable energy sources.

2.1.3 Energy Constraints

Commonly, sensor nodes do not have access to a constant power supply and must run on a battery that is, generally, a similar size to the node. Larger batteries must usually be contained in a separate enclosure to the node and this makes nodes less compact and their deployment more difficult. This means that the nodes must be as efficient as possible, knowing when to transmit data and when to sleep. The lifetime of a sensor network is highly dependent on the battery life of each node and, unlike other mobile devices, they cannot typically be recharged [15]. Much work has been done on power efficient routing protocols, as well as the control of which attached devices are active [57, 100, 101].

The limited resources on the nodes mean that the sensing devices, and transceivers, attached must consume as little power as possible. Some routing protocols implement turning off wireless radios and scheduling a wakeup across the network [122], but the cost of turning off a device can waste just as much energy as leaving it on and sampling at a lower rate, if not more [38].

The use of energy in a node is dependent on how active the sensor(s) are, how much it transmits and receives, the transmission medium used as well as the environment it is in.

2.1.4 Transmission Medium

Widely used, general purpose transmission media, such as Wi-Fi, are viable solutions in WSNs when a high data rate is required and power is readily available. However, research has shown that Wi-Fi is extremely power-hungry and [65] shows that Wi-Fi consumes almost 9 times more energy, while transmitting, than other standards, such as Zigbee. Bluetooth is a more power-efficient standard that is becoming increasingly popular for sensing devices that are part of the ‘Quantified Self’ movement [105], with wearable devices that report measurements, such as heart rate, steps taken and calories burned. With the advent of the new low-power Bluetooth 4.0 , also known as Bluetooth Low Energy (BLE), this standard is supposed to allow months of continuous use on a coin-cell battery [45]. However, the theoretical max range is 100m and, using the same frequency, as Wi-Fi (2.4GHz) means that it is as susceptible to path loss and reduced transfer rates. [127] shows that a 2.4GHz Wi-Fi antenna is capable of transmitting up to 350m, while a considerably lower frequency of 41MHz was able to achieve links of 10km. The use of 2.4GHz frequencies in wet conditions have been shown to reduce the performance by up to 28% [75], while humidity in the air can reduce the performance by up to 78% [41]. New low-power, low-frequency standards have emerged in recent years and allow for a considerably longer range and increased battery life, at the cost of transmission speeds. Digimesh is an example of this and, while it can achieve 250kb/s using 2.4GHz, it has much slower speeds of 125kbps when using the 900MHz spectrum. However, it does offer a range of, up to, 64Km [20].

2.1.5 Environment

The environment that a node is deployed in can have a great impact on almost all aspects of a WSN, such as range or the operational lifetime of the node itself. Harsh environments that are not easily accessible make it difficult to place nodes and protected environments may limit where nodes can be placed. In [76], nodes were deployed within glaciers and had to survive extreme temperatures, lasting without human intervention, for months at a time. [69] attached collars to Zebras that had to withstand high speed movement, impact, dust and high temperatures. The deployment of any node requires extensive research as to the environment that it will be deployed in and adjustments must be made to ensure it is able to survive for extended periods without continued maintenance. Section 2.1.4 also shows that environment does not simply affect the hardware, but humidity can reduce the transmission range significantly, as well as moisture collecting on wireless antenna can reduce the range for days at a time.

Limited range in a WSN can be addressed by using *intermediate nodes* - nodes tasked with forwarding data from other sensing nodes that would otherwise be out of range - between disconnected sensing nodes. While they do not need to sense the environment directly, they can be vital in ensuring data from all areas of the network are delivered. In the case of [69], a car was used to drive near to zebras that had not passed close enough to the base station in order to transmit their data. The car acted as the intermediate node by collecting data from the zebra's collar and relaying it back to the base station. In a normal WSN, this has the downside of requiring more nodes, which increases the cost and the chance of a node failing, but it can connect clusters in a geographical region that would otherwise be unable to send sensed data, or they can provide multiple routes through the network to help preserve battery life.

2.2 Routing Protocols

Routing protocols specify how nodes in a WSN are organised, as well as how they transmit data throughout the network. In [14], the more popular routing protocols are surveyed and split into three of the main identified categories: data-centric, hierarchical and location-based. We use the aforementioned categories, as well as flat, to highlight some of the key protocols that are relevant to our work. Flat routing protocols are for WSNs that involve many nodes deployed over a large area, all sending data to a single endpoint. Data-centric is a protocol that focusses on sensed-data, where nodes advertise their contents and query other nodes to fulfill requests made by a user. Location-based protocols use the physical region that nodes are deployed in to request and send data. Hierarchical protocols are for heterogeneous networks that perform more than one task. For example, a network spread of over hundreds of miles could be split into individual networks, or it could follow a hierarchical structure where clusters of nodes send data to an gateway node (also known as a *cluster head*); a node that may not directly sense the environment but acts as an endpoint for a group of nodes [14]. The data then hops across other nodes in order to reach a final endpoint. The protocols have the task of ensuring that a network is performing at its best, providing the best lifetime and ensuring reliable and consistent delivery of data. This must reduce *flooding*, where nodes send every message to every link, aside from themselves, effectively flooding the network with unnecessary messages, and find a way to deliver data to the endpoint using the most efficient path possible.

2.2.1 Flat

Initially, this was the most common structure for a WSN, dozens of nodes spread out over a geographical area, with one or more neighbours, sending observations to a single endpoint.

MCFA

The Minimum Cost Forwarding Algorithm (MCFA) is a flat routing protocol that works by assigning costs to each node, based on how many hops they are from the endpoint [124].

Each node has a path-estimate of the cost of transmission from itself to the base station. The base station sends out a broadcast message and it is received by all nodes in range. The message contains a cost from the base station (initially zero) while every node has their cost set to infinity. The cost is stored at the node, incremented and sent on to all nodes in range. If the received cost is less than the current cost stored on the node, then the cost (and the neighbour) is updated and passed on.

Each message has a cost associated with it, which is based on the hops it has completed so far. A node that receives the message forwards it only if its cost matches the sum of the source node's cost (contained within the message) and the message's current cost. This ensures that all messages are sent through the minimum cost path, without storing explicit path information on each node.

This approach allows for dynamic reconfiguration of the network, as well as a reduced overhead due to not having to maintain a global routing table on each node. The assumption with MCFA, however, is that the direction of routing is always towards a fixed endpoint.

2.2.2 Data-centric

Data-centric routing protocols are not like traditional WSNs where nodes are given addresses; they use a method that involves the advertising, or querying, of the data that has been sensed and those with the relevant data can respond to the request.

SPIN

Sensor Protocols for Information via Negotiation (SPIN) is one of the first data-centric protocols and attempts to address the issue of flooding the network whenever new data is sensed by addressing the data through metadata [56]. SPIN works on three messages passed between nodes:

1. ADV - A message sent by a node when it has sensed new data, advertising what it has recorded.
2. REQ - Sent by nodes that received an ADV to request the data.
3. DATA - Message containing the sensed data.

When a node has sensed data, it sends an ADV message to all nodes within range. If any of those nodes are interested in the data, then they respond with a REQ message, at which point the DATA message is sent to nodes that responded.

SPIN eliminates the need for a global view of the network topology, as nodes only need to know their single hop neighbours. However, SPIN does not guarantee equal diffusion of data throughout the network as a node may be interested in the data sensed at the other edge of the network, with only nodes that are not interested in between. This would mean that those nodes would not request the data or pass it on.

SPIN-IT

An extension to SPIN, SPIN-IT uses a slightly different approach to receiving data and was developed solely for the transfer of images [120], using metadata to fulfill requests.

Nodes use the existing message structure of SPIN, but REQ messages are used as queries, sent to all nodes in transmission range. The receiving nodes keep these requests and generate a new REQ message, thus allowing nodes to store temporal paths. When a

REQ reaches a node that has the desired data, it responds with a ROUTE-REPLY message. This message is used because images are large and resource-constrained WSNs would have a much shorter lifetime if a lot of unnecessary transmissions were made. The ROUTE-REPLY is used in case multiple nodes, in range of the requesting node, have the data it has requested and it can then choose the optimal route. As each node keeps a history of REQ messages, these can be used to trace the requested data back through the network, to the originating node, without the overhead of maintaining a global routing table.

COUGAR

A slightly different data-centric approach is the proposed COUGAR protocol, viewing the network as a distributed database. Cougar is similar to SPIN because it does not forward data as soon as it is sensed, instead COUGAR uses a query language that abstracts the underlying network structure and uses received queries to generate a plan that utilises in-network processing to provide an answer based on the sensed data stored on all deployed nodes [123].

For example, in a building monitoring network, a user could query a base station for offices that are unoccupied. The base station then sends that query to all nodes in range and it is then dispersed throughout the network. Nodes that have data that can satisfy the query send back their results and these are combined as they move back through the network to the base station. The user then receives that data, along with the nodes that have provided it.

Within the network, a *leader* is selected and this node is used to aggregate the data from nodes that were able to fulfil all, or some, of the query. At risk of failure, each query should result in a *leader* being dynamically selected and it must have sufficient resources to be able to satisfy the request. This protocol was only proposed, and much of the technical detail has yet to be completed, but the concept of treating the network as a distributed database is a novel idea and this is one of the first protocols to suggest

the use of a query language that could be used by people without specific technical knowledge, which, in this case, is technical knowledge of query languages and how the routing protocol works.

2.2.3 Hierarchical

Hierarchical networks are WSNs that contain nodes of different classes; nodes at the edge of the network are typically clustered into groups and served by a gateway node. This gateway could be in charge of aggregating the data, processing the data, or simply forwarding it to an endpoint. Clusters of nodes allow the network to be spread out over a wider geographical area and gateway nodes can use a different transmission method to provide long distance links to the base station. Gateway nodes serving a cluster of nodes means that the network can scale easily as well, simply by adding a new cluster to the network.

TEEN

The Threshold sensitive Energy Efficient sensor Network (TEEN) protocol is designed for reactive sensor networks, networks that require instant reactions to changes sensed in their environment [74]. TEEN recognises that transmission is the most power hungry action for a node so each node is coded with a hard and soft threshold. The hard threshold is a value that makes nodes transmit the reading to their cluster head. Similarly, the soft threshold is a small change in the value of the sensed attribute that causes further transmissions.

During the initialisation of the network, the base station sends information about the thresholds and sensing attributes to all cluster heads in the network; the cluster heads then forward this on to all nodes in their cluster. When a node senses data over the hard threshold, it transmits to the cluster and only transmits again when new sensed values

are greater than the hard threshold and the difference between the current sensed value and the previous is greater than the soft threshold [74].

Clusters are assigned for a period of time and then new clusters are selected by the base station, at which point new attributes and thresholds are broadcast to all nodes. This kind of protocol allows the network to be dynamic after deployment and allows user input based on the data that has been sensed in the previous cluster times.

For example, a network could be tasked with sensing humidity in a rainforest but the thresholds have been set such that nodes are transmitting readings that are not of interest. A user can change these thresholds and they will be pushed out to the nodes at the time that the next clusters are chosen, without any need to visit the node or configure them individually.

2.2.4 Location-based

Instead of using the physical addresses of nodes, or the data they store, location-based protocols are based on the region that nodes are deployed in. Location-based routing relies on the fact that each node is aware of its own location and is also aware of the destination's location.

Span

Span is a protocol where nodes are selected as *coordinators* based on their positions. A node can decide to be a coordinator based on the amount of energy it has and the number of neighbouring nodes it would benefit if they were able to use it as a bridge [26].

An example of this would be node B placed between node A and C. C and A are unable to communicate directly so, when node B wakes up, it decides whether it should become a coordinator. It knows that it has sufficient energy levels and it can provide

connectivity for a previously disconnected area of the network, so it chooses to become a coordinator, staying awake and routing sensed data to other coordinators, which form the backbone of the network.

Results showed that using Span, in a system that transmits using 802.11, provides an network lifetime increase of more than a factor of 2 over networks that just use the 802.11 protocol.

GEAR

The Geographic Energy-Aware Routing protocol (GEAR) is similar to SPAN in that it makes routing choices based on both energy-awareness and location. Each node maintains an *estimated cost* and a *learning cost* of forwarding a packet through its neighbours. The estimated cost is calculated using the distance to the packet destination and the energy remaining on the node whereas the learning cost is the estimated cost that takes holes in the network into consideration [125].

GEAR is designed to perform in two phases: forwarding a packet towards a region and disseminating a packet within a region. When sending a packet towards a destination, GEAR either sends a packet on to the node in range that is closest to the destination or, if such a node does not exist, then a hole is identified. If a hole is identified then the node that minimises a cost is selected.

To disseminate a packet within a geographic area, GEAR uses algorithms based on the density of the network. Recursive geographic forwarding is typically used but this can result in an endless loop if the density of the network means that the region is unable to contact the destination. In that case, restrictive flooding is used.

2.2.5 Conclusion

Routing protocols can define the topology of a network, how data is sent and have an impact on the network lifetime by determining when nodes should sleep, when they should transmit sensed data and when (or if) they should request updates on the network topology. Researching these types of routing protocol allows us to determine the situations in which each category would be used and whether they fit our requirements. The routing protocol must be considered when developing a network architecture, so we have researched the four main categories and the more commonly-used protocols within those categories. If we chose a flat network structure, then MCFA would be the better choice but our architecture changes significantly if we choose to use a data-centric protocol.

Using this background knowledge, we can select a protocol that fits with our needs, or combine useful aspects of many. Data-centric protocols are useful for battery conservation and networks that do not require real-time reports. Location-centric is useful as it does not require specific node addressing but location-aware nodes are more expensive and power-hungry. Hierarchical, however, allows for in-network data-aggregation or processing but requires the use of a heterogenous network with a more rigid topology of nodes split into clusters.

With the knowledge of the requirements of a WSN architecture, we can pick the routing protocol that best suits these, or combine traits from multiple protocols in order to develop a hybrid protocol that fits our requirements exactly.

2.3 Sensor Middleware

Acting as a bridge between the hardware and the user, sensor middleware is software that abstracts the underlying network from the user and provides a means of accessing sensed data and administrating how the network performs [53]. These middlewares

must not be specific to a single network and provide support for as many different sensor nodes as possible. In [126], a middleware is said to provide standardised services to many applications and perform operations that make effective use of limited system resources.

In [116], a middleware should include four major components: programming abstraction, system services, runtime support and Quality of Service (QoS) mechanisms. In this section, we will discuss the challenges surrounding middlewares for WSNs and highlight some existing middleware that are particularly relevant to our research problem and motivating scenario.

2.3.1 Issues

WSNs present a range of new challenges to existing middleware, due to their resource constraints, deployment environments and more. However, there has been research into the key issues that must be addressed in order for middleware to be considered suitable. While there have been a number of surveys into these challenges [53, 94, 126], we will detail those that we believe to be most relevant to our work. Some of those that, while considered, have not been a primary issue are: dynamic network organisation, security and application knowledge.

Energy Constraints

It is rare that nodes in a WSN would have a constant power source, unlimited memory and a casing that can survive a harsh environment without decaying. In order to ensure that the lifetime of nodes is maximised, middleware needs to offer a power scheduling system that makes efficient use of the hardware on the node, generally disabling the radio after a set interval has elapsed, or use a combination of lower power sensors to provide sensed data of a similar quality [55].

Ideally, a middleware will be able to coordinate nodes through wireless communication, making efficient use of transmissions and dynamically modifying sleep schedules based on the power remaining.

Heterogeneity

Not every node in a network will have the same capabilities, manufacturer or hardware. WSN middleware needs to provide a standard interface to the applications that are making use of it, whilst at the same time accommodating differences arising from hardware coming from different manufacturers and perhaps having significantly different capabilities. Some middleware have been built for a specific set of hardware [109], however this homogeneity can provide an increase in the performance and efficiency of the network by only supporting a limited number of devices.

Real-world Integration

WSNs are often tasked with recording phenomena that are time-crucial, so a sensor middleware should provide a real-time interface to the data that it has sensed [12]. Ideally this data would be available outside of the network, though the use of an API.

Quality of Service

This issue is perhaps the most complex as QoS could apply to almost all aspects of the networks, such as efficiently using bandwidth, 100% uptime for nodes, guaranteed packet delivery or access to data stores. Some of these requirements are managed by the implementation of the routing protocol, the middleware should be able to monitor deployed nodes and report on their current status, as well as identify failures.

2.3.2 Existing Middleware

In this section, we identify existing middleware, explain how they address the issues highlighted in Section 2.3.1 and highlight how they relate to our research. While there are a lot of existing middleware, our research did not show any that used information abouts its environment or knowledge from previously sensed data to process data that is currently being sensed. We did, however, find some middleware solutions that utilise context and rules to administrate the network.

GSN

The Global Sensor Networks middleware (GSN) has been developed to manage heterogeneous sensor networks and be suitable for those without any technical knowledge [12].

GSN provides hardware abstraction through the use of *virtual sensors*, a data stream that abstracts implementation details from the actual sensed data. A virtual sensor can be comprised of many streams and it can even consist of many virtual sensors.

Virtual sensors are described using XML, with tags that consist of metadata for the sensor, the structure of the incoming data stream, SQL queries for processing the incoming data and querying times. What makes GSN stand out is that virtual sensors do not have to be sensors deployed within your network, or even sensors at all, some examples of GSN show virtual sensors being added that read in data from the weather websites. This also means that the underlying structure of the network is irrelevant to GSN, as well as the physical locations of the nodes. Unlike some middlewares that have an expectation of how data will be routed, GSN is decoupled from the routing protocol, allowing them to act independently.

GSN is completely open source and the Java code can be modified to suit a specific deployment.



Figure 2.1: GSN Architecture [12]

Figure 2.1 outlines the architecture of GSN, showing that the virtual sensors are stored on a central node and their inputs are managed and stored. GSN also comes bundled with a web interface to show all active sensors and their most recent recordings, as well as the implementation of web services to access the data outside of the interface.

Data from virtual sensors pass through the virtual sensor manager to the storage layer. Once the data has been stored, the query manager is invoked and queries are loaded from the repository and executed by the manager. The results of the queries are then handled by the notification manager and also made available to the web interface. No-

tifications can be extended to support many different forms of communication, such as SMS, email or web services.

Virtual sensors do not natively support all hardware, although new virtual sensors can be described using XML, and there may be a need to implement an entirely new virtual sensor. In this case, technical knowledge is required, and new sensors can be implemented through use of the Java programming language. This provides more control over the use of XML and allows users to specify how sensed data is stored in a database, use external libraries to receive proprietary data, specify processing workflows before the data is stored or implement new notification methods for the users of the network.

To show the simplicity of a basic virtual sensor, [11] describes a temperature sensor that we reproduce here in Listing 2.1. The file is human-readable and does not require specialist knowledge when compared with programming languages, with tags that have self-explanatory names. In this example, the output structure shows that only a temperature reading is received and that the data should be stored permanently. The stream source specifies the content of the stream and the query details the standard query that should be used to extract data from GSN.

```

1 <life-cycle pool-size="10" />
2 <output-structure>
3   <field name="TEMPERATURE" type="integer"/>
4 </output-structure>
5 <storage permanent-storage="true" size="10s" />
6 <input-stream name="dummy" rate="100" >
7 <stream-source alias="src1" sampling-rate="1" storage-size="1h">
8 <address wrapper="remote"> <predicate key="type" val="temperature" /> <
9   predicate key="location" val="bc143" /> </address>
10 <query>select avg(temperature) from WRAPPER</query>
11 <query>select * from src1</query>
12 </input-stream>
```

Listing 2.1: Example Virtual Sensor

The modularity and flexibility of GSN makes it different to existing middleware as it has not been designed for any specific hardware and modules of the middleware can be replaced, such as the database.

FACTS

One such example is the FACTS middleware, an approach that uses a fact repository to coordinate nodes. Rules can then be implemented to process sensed data and fired when certain conditions are met [110]. More traditional sensor middleware controls the network and manage sensed data but this rule based approach allows for more flexibility, where rules can control the transmissions and process the data upon receipt.



Figure 2.2: FACTS Architecture [110]

Figure 2.2 shows the FACTS architecture, with the middleware holding the rule-sets and a distributed fact repository. Data within the network is stored as facts, providing a standard data format throughout the network and hardware abstraction. When new facts are received, usually because of new sensing data, the rule engine checks to determine whether any rules should be fired. The rule-set definition language (RDL) is used here and each rule-set contains a group of relevant rules. Each rule is given a priority so that, if more than one rule is triggered by a fact, then the higher priority rules are fired first.

Listing 2.2 shows a FACTS rule, written in Haskell, that determines which geographic areas are covered by nodes. The name of each rule has no prefix, statements are prefixed with '`->`' and conditions are prefixed with '`<-`'. The `sendRange` rule runs after a timer expires and removes that timer. Lines 4 to 12 set a `rangeFact` fact that contains the range it expects to be able to cover. The rest of the rule sends the fact to all nodes within range and sets a new fact to show that it has sent its range information.

The (xyMinCovered) is a simplified rule that, upon receipt of range information from neighbouring nodes, checks whether the range of the neighbouring node overlaps with its own range and the final *determineCoverage* rule then stores that coverage information in a fact as knowledge of whether or not it is the only node to cover a particular geographic region. This can then be used to inform future routing and sleeping decisions.

```

1  sendRange
2  <- Exists Timer.expiredSlot
3  -> Retract Timer.expiredSlot
4  -> Define "rangeFact"
5  -> Set ("rangeFact" "xMin")
6  (posXSlot - System.txRadiusSlot)
7  -> Set ("rangeFact" "xMax")
8  (posXSlot + System.txRadiusSlot)
9  -> Set ("rangeFact" "yMin")
10 (posYSlot - System.txRadiusSlot)
11 -> Set ("rangeFact" "yMax")
12 (posYSlot + System.txRadiusSlot)
13 -> Send 0 System.txPowerSlot
14 ("rangeFact" [((("rangeFact" "owner")
15 == nodeIDSlot)])
16 -> Define "rangeSendFact"
17
18 xyMinCovered
19 <- Exists "rangeSendFact"
20 <- Eval ((posXSlot - System.txRadiusSlot)
21 < ("rangeFact" "xMin"))
22 <- Eval ((posYSlot - System.txRadiusSlot)
23 < ("rangeFact" "yMin"))
24 -> Define "xyMinCoveredFact"
25
26 determineCoverage
27 <- Exists xyMinCoveredFact"
28 -> Define "coveredFact"

```

Listing 2.2: Coverage Algorithm in FACTS Rules

While FACTS itself does not utilise any local knowledge, the repository is used as a source for all previously sensed data and would be an excellent source of knowledge to assist with the classification of future readings. Also, the ability to add new rule-sets, without technical knowledge of the hardware of each node means that users of the network have the ability to add knowledge in the form of less technical, high-level rules.

ITA Sensor Fabric

The ITA Sensor Fabric is a collaboration project between IBM, the US Army and the UK Ministry of Defence. Sensor Fabric, or Fabric, is a two-way messaging bus and set

of middleware services connecting network assets to each other and users [121].

The core difference between the Fabric middleware and others is that not every node is sensing all of the time, sensor nodes are tasked when there is a requirement and they stop as soon as that task has been fulfilled. Similar to sinks in a traditional WSN, Fabric utilises Fabric nodes, which run the following three pieces of software:

1. Message Broker - Provides the communication infrastructure.
2. Fabric Registry - Holds information about the current deployment, such as all nodes deployed, all assets, routing information and tasks. Deployed in the form of a database.
3. Fabric Manager - The main service on the node to track the status of connected sensors, establish communication channels, provide a container for processing, plug-ins and to extends the capabilities of the Fabric.

Fabric runs on a Publish/Subscribe model, a sensing requirement is sent to a messaging broker as a subscription and this is distributed through all Fabric nodes and, thus, all sensor nodes. Sensor nodes then publish their data and the relevant data is sent to all applications that have subscribed to the data.

The plugin structure of Fabric makes it stand out from existing middlewares, allowing its functionality to be extended through web interfaces.

Because Fabric has been developed for military purposes that cross countries, policy enforcement has been implemented to restrict access to the granularity of sensed data but these access levels do not simply apply to a military context. Using our motivating scenario, researchers and professors should see animal images whereas the Sabah Wildlife Department should see images of hunters and people in the forest.

Sensor Web Enablement

The Open Geospatial Consortium's (OGC) Sensor Web Enablement (SWE) is a set of standards to allow developers to make sensors and sensed data repositories accessible via the Internet [96]. While this is not a middleware, these standards can work with existing middleware in order to publish their sensed data to the Internet. The SWE framework consists of:

- Observations & Measurements (O&M) - General models and XML encodings for observations and measurements.
- Sensor Model Language (SensorML) - Standard models and XML schema for describing the process within sensor systems.
- PUCK - Defines a protocol to retrieve a SensorML description and other information from a device, enabling automatic sensor installation, configuration and operation.
- Sensor Observation Service (SOS) - Open interface for a web service to obtain observations and sensor descriptions for one or more sensors.
- Sensor Planning Service (SPS) - An open interface for a web service that a user can determine the feasibility of collecting data from one or more sensors and submit collection requests.

These components can be integrated into existing middleware, clients and sensors and not all need to be used, some middleware, such as MuffIN [113], have just used SOS and O&M but 52North have developed an open source implementations of all components of SWE [24].

Internet of Things

In [44], the Internet of Things (IoT) is described as the “concept of pervasive things... which, through unique addressing schemes, are able to interact with each other and cooperate with their neighbours to achieve common goals.” The number of devices that are part of the IoT grows daily, from thermostats and smoke detectors [8] to activity monitors [4]. Some of these devices have their data siloed away and it is only accessible through a web interface or a mobile app, but most have an Application Programming Interface (API) that allows the raw data to be accessed and queried. Because of this, sites like If This Then That (IFTTT) [6] have been created to link together various services and use simple if statements to act on changes in the data. IFTTT uses *recipes* that follow the rule of: `IF Service1 changes THEN perform action on Service2.` For example, `IF I take a picture using my phone THEN upload it to Twitter OR IF a new item appears on my RSS feed THEN add it to my reading list.` While, IFTTT is not a traditional middleware, it does monitor user’s connected services, some of which may be sensors around the home or attached to their person.

A more traditional middleware for the IoT is Xively, a cloud platform that allows devices to upload their sensor readings to an endpoint and view readings from others [10]. Feeds can be public or private and Xively can be used with almost any Internet-enabled devices, such as webcams, temperature sensors or alarms. Other sites, such as Dweet.io [3] and Open Sen.se [9], have also been in operation for a few years that provide a similar service to Xively. Open Sen.se combines the recipes of IFTTT with the storage of Xively to create a middleware that supports actions being performed when certain sensed data is received. For example, a temperature sensor could trigger an alert on Open Sen.se that would alert a user via e-mail, or through social media. This combination of devices with online services and applications allows a middleware to do more than simply store data, it can execute rules over the data, provide visualisations and create links between devices separated by hundreds, or thousands, of miles.

2.4 Biodiversity and Environmental Monitoring Sensor Networks

In this section we will cover existing WSNs that are related to our motivating scenario or, more specifically, biodiversity focussed WSNs that have been deployed to monitor wildlife and/or the environment. WSNs for habitat, and wildlife, monitoring are especially important because these are areas that often need to be untouched by humans. Areas with high human disturbance can influence the abundance of species and some habitats, i.e. underground burrows, may be impossible to monitor without destruction.

One of the most well known WSNs to monitor habitat is the network deployed on Great Duck Island, an island off the coast of Maine, USA. A hierarchical network of 32 nodes was deployed to monitor a bird, known as the leach's storm petrel [73]. This network used a clustering approach for groups of nodes to send data to a gateway node, which would then route it back to the base station. The base station, located a few kilometres away on the island, has internet access and uploads the data to allow users to browse and process the data.

A multihop approach was used here as they found that, for sufficient coverage, single hop connectivity would not cover all of the island. Acrylic enclosures were developed to ensure the nodes were weatherproofed for the conditions of the island, while maintaining the functionality of each sensor and not impeding transmission range. While the nodes, their casing and their sensors have been designed specifically for the deployment on Great Duck Island, the success of the network, running for 123 days in the early stages of WSN research [108], shows that this approach can be used elsewhere with similar effects; allowing hard to monitor and/or inaccessible areas to be continuously monitored.

On a smaller scale, INternet-Sensor InteGration for HabitaT monitoring (INSIGHT) is a single-hop WSN that allows remote access for data and reconfiguring of nodes [34]. Using off the shelf hardware, their findings show that their nodes could survive for 160

days on a single battery, supporting their claim that a single hop network allows for a longer network lifetime.

The key feature of this network is the ability for humans to remotely set reporting thresholds for sensor nodes. This means a user can prolong the lifetime of nodes by limiting the threshold they report on, as well as the fact that these thresholds are a way for users to add knowledge, albeit primitive, into a network.

While there is research on cameras used to monitor animals [63, 13], these networks are generally cameras deployed with their memory cards manually retrieved and processed. In recent years, however, the use of wireless technologies and image-based WSNs has increased, [42] uses wireless cameras to monitor the movement of animals between roads. Using commercial hardware and controlled sleep scheduling, this solution employs the use of nodes to detect movement and wake up more power-hungry camera nodes. While the nodes are wireless, the distance of the network from civilisation means that the data does still need to be collected manually and uploaded to a computer.

Due to the advent of smartphones and tablets, as well as the improvements in 3G technology, projects taking advantage of more modern technologies have grown in popularity. Using 3G enabled cameras, [129] have deployed a number of devices in locations all over the world, such as: Kenya, Indonesia and the USA. The images captured are transmitted to a server and a website allows the general public to not only see the images in near real-time, but to classify the images as well. This crowd-sourcing of collective knowledge lets people, that may not have domain knowledge, vote on an image and those votes are used to make classification easier.

Over the past fifteen years, WSNs have grown from a concept to a real solution for monitoring the habitats, movements and eating habits of wildlife all over the world. Whether it is using GPS collars to monitor the movement of cattle [47], monitoring animal habitats on a remote island or using cameras to capture the animals themselves, the popularity of these networks has grown considerably and advances in technology

have allowed these networks to be deployed in places that humans cannot access with ease.

There are a number of situations where we may want to record data in an environment that is not safe for humans, such as a volcano [118], or that may be difficult for electronics to survive. In these cases, special considerations must be made during the design and deployment of the WSN, in order to ensure maximum network lifetime with reliable readings.

GLACSWEB

Glacsweb is a sensor network to monitor the rate at which glaciers are melting [76]. Deployed in Norway, specially designed sensors have been drilled into glaciers to monitor pressure, temperature, orientation and strain. Due to the high pressure and exposure to moisture, a polyester casing was used so that, once bonded, the node inside would be protected from its environment, but also preventing it from being recoverable.

Reventador

Volcano Reventador is in northern Ecuador and a WSN has been deployed on there to monitor eruptions. Similar to Glacsweb, weatherproof enclosures were used to prevent ash and moisture from breaking the sensors and long-range external antenna mounted to a pole was used to achieve communications over large distances when wireless communications have proven to be difficult [118].

Rainforest

There have been many studies on how the rainforest affects the range and quality of wireless links [41, 117, 93]. In [41], the humidity was shown to reduce 802.11 range

by up to 78% and [117] explains that periods of rainfall reduce the link quality up to 100% in some cases, resulting in the loss of a hop and this could prevent data from reaching the endpoint.

These examples highlight some of the difficulties of deploying WSNs in harsh environments. Environments vary from freezing glaciers to humid rainforests to dry deserts with extreme temperature variation and the sensors cannot disturb their surroundings, be too conspicuous or require regular human attention that would disturb the area and affect the wildlife.

2.5 Conclusion

In this chapter we have explored the components that make a WSN, as well as some existing deployments that are relevant to our motivating scenario. Sensor middleware have come a long way in the past decade and increased capabilities for sensor nodes have allowed for more intense processing to be carried out before any data is seen by a human. Context is now used on sensor nodes to infer the activity they are undertaking or even to determine when it should wake to sample.

Local knowledge is not a new concept and it has been used for many years to extract information from indigenous people and in industry to adapt their processes to a local area, such as farming. However, we believe that our work is the first to apply local knowledge in the context of WSNs. Context-aware sensor networks have gone some way to support our hypothesis that some knowledge can increase the functionality of a network and enhance the quality of the sensed data, but we believe that the addition of local and global knowledge will only increase this functionality further.

The two primary components of a WSN that could be injected with local knowledge is the middleware and the routing protocol, each providing different benefits. Existing work has shown that the use of context-awareness in sensor middleware allows them to make dynamic, global changes to the network (Section 2.3), such as power

management, whereas routing protocols affect the data that is sampled and sent to the endpoint(s) of the network, such as adaptive thresholds.

One issue in proving this hypothesis is the deployment of a network that utilises local knowledge. The deployment environment of our motivating scenario is not only interdisciplinary, it is in a region that is humid, dynamic and dense. Existing research has shown that the deployment of a WSN in these conditions means that range will be greatly reduced [41] and changes in humidity can prevent communication altogether, as well as moisture affecting the hardware itself. To deploy a network, hardware must be adapted and the right medium must be chosen in order to maximise link quality and minimise dropped connections.

In Chapter 3, we describe how we used this research, along with our own findings, to choose suitable hardware for our network. As well as this, we used the results from previous rainforest range tests in the literature, to design experiments that would aid us in choosing, and verifying, our choice of transmission medium.

Chapter 3

Exploratory Experiments in a Rainforest Environment

In this chapter, we explain our motivating scenario in more detail, explore the sensor hardware and software that we researched and outline the results of experiments we undertook in the Malaysian rainforest. As highlighted in Section 1.1, we have been working with Cardiff University School of Biosciences to design and deploy a WSN that utilises local and global knowledge, using an area of rainforest in Malaysia owned by the Sabah Wildlife Department, called Danau Girang.

The structure of this chapter is as follows. Section 3.1 explains what we aimed to deploy in Danau Girang and what our considerations were. Section 3.2 introduces sensor hardware that is in use today and details the choices we made. Section 3.3 details the transmission medium choices we tried and also shows the results of experiments performed in both the UK and Malaysia. Section 3.4 explains some of the software applications we developed in order to process sensed data on the nodes. Finally, Section 3.5 summarises our findings and explains the choices we made for the sensor nodes we used in Danau Girang Field Centre (DGFC).

3.1 Danau Girang Field Centre

Based in Sabah, Malaysia, Danau Girang Field Centre is located in Lot 6 of the Lower Kinabatangan Wildlife Sanctuary (LKWS), surrounded by secondary rainforest that had been logged up until the 1970s. Experiencing a wet season from October to February and a dry season from March to September, the LKWS can receive more than 500mm of rainfall during the rainy season, dropping to lows of around 150mm during the dry seasons [115], and up to 100% humidity all year round.



Figure 3.1: The Main Building at DGFC

Danau Girang is in a unique location, situated in a rainforest corridor (shown earlier in Figure 1.1 in Section 1.1) that joins two areas of rainforest together, with the corridor surrounded by palm oil on each side. Because of this, animals use the corridor to move between the rainforest regions and some use it to enter the palm oil plantation for new feeding grounds. This gives the DGFC insight into the movement patterns of these

animals in the corridor as well as in the rainforest itself, with a wide variety of species that are not commonly seen in other tropical regions of the world. Due to the remote nature of the centre, power is provided by a set of diesel generators which, typically, provide power from 10 am to 1 pm and 5pm to 11pm daily. Wireless Internet access is provided by satellite with speeds approximating 56kbps, although the upload speeds are considerably faster than downloads.

The corridor monitoring programme is a scheme that has been in place for more than five years, using wildlife cameras to track the movement of animals through the rainforest corridor and to capture species that are rare or unique to South-East Asia, such as the Bornean clouded leopard. Currently, Reconyx Hyperfire HC500 (Figure 3.2) cameras are being used [1]. These are standalone cameras that capture a user-defined number of images which are saved to an SD card. A trigger is caused by an infrared (IR) motion sensor when an object causes the IR beam to break.



Figure 3.2: Reconyx HC500 Camera

Images must be manually collected every two weeks from the cameras and the batteries

are changed at that point as well, although a typical charge should last three months. In the field, the batteries last a matter of weeks shortly after the cameras are deployed, this is most likely because, although the cameras are equipped with watertight casing, they are still exposed to moisture when the case is opened every two weeks and 100% humidity seems to have an impact as well. Silica gel is used to prevent moisture inside the camera. We believe that humidity also reduces the battery life, as the charge drops from three months to around three weeks within the first few months of usage. However, the lack of constant power availability could also affect the charge of the batteries. Because power is not available all day, the batteries are charged for a period of a few hours at a time. Not only does this reduce the number of charge cycles that these batteries can experience but it also means that they are never fully charged within one time period of power availability. Solar recharging is being researched to counteract this. Each camera is secured to a tree and, when placed in higher risk locations, such as known elephant paths, have protective cases as well.

In 2010, twenty cameras were deployed for six month periods and then relocated based on the needs of the projects at that time. Currently, there are now ninety cameras with a view to expand and dozens of projects within the field centre use the images gathered from these cameras. Initially, it was the job of visiting research students to collect the images but, since the number of deployed cameras has grown, full-time staff have been taken on to maintain them.

Images that have been collected are stored on an external hard drive in the field centre and a spreadsheet is updated with information about each image, such as date, time, camera, filename and location. Research students and, later, staff, manually inspected each image and added their species classification to the spreadsheet. Because so many images came through the field centre each week, this was a long process and it was not uncommon for it to be a few weeks before the images were classified. Research students in charge of the cameras were usually on projects lasting 6 months so, until staff members were recently employed, their methods of cataloguing were quite different.

Sometimes, each student created a different spreadsheet for their time at DGFC and ignored the classification of images that did not contain the animal their project was focussed on. While this has improved now, staff members are tasked with combining these spreadsheets from the past few years and classifying all the images that have been left empty.

With image processing, we can automate the classification process. Using common tools and widely accepted methods, such as OpenCV [21] and object detection libraries [82], we can use previously classified images as templates. This allows us to determine if a new image contains an object of interest and extract templates of that object to match to new images and classify them to the species level, once a human has classified the original template. Humans can then check the classification and provide feedback on its validity. This removes the wait of weeks for a classification and automates the cataloguing, classification and storing of the images, when received at the field centre. This method should also allow us to alert researchers when our system believes an image of particular interest, such as a rare animal, has been taken. Without humans to classify images down to the species level, we can still use image processing to determine if an image is interesting or empty.

Our belief is that we can use the Lower Kinabtangan wildlife sanctuary (LKWS), and the locations of the existing cameras, to deploy a WSN that uses local knowledge gained from the researchers at DGFC to automate the collection of images, improve the battery life by not exposing the internals of the camera to the elements so often and, most important, prioritise the flow of data through the network by in-network processing.

Three annual visits, in 2011, 2012 and 2013, each lasting three weeks, have been made to DGFC to test out hardware, software and wireless choices, in an effort to optimise the network. These visits have also been used to extract local knowledge from the area and researchers, by semi-structured interviews and watching them work.

3.2 Hardware

Before any visits were made to DGFC, meetings with staff members of the field centre were held in order to gain a better understanding of the environment and the project. This is where we were alerted to the humidity of the region and the fact that the failure rate of the Reconyx cameras has been as high as 30%.

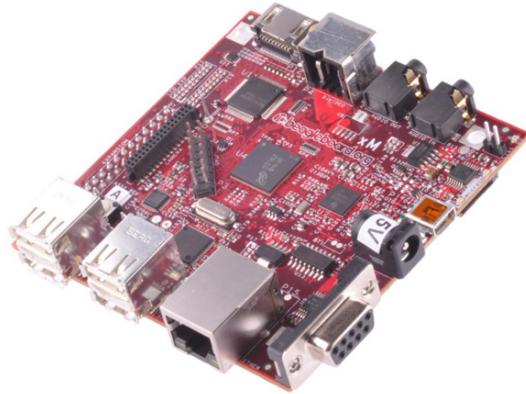
Reconyx cameras have no external interface support and the only way to access the images is through the removable SD card. Because of this, there is no way of attaching external sensor hardware to the existing cameras. In-network processing is an important requirement for our WSN and this did limit our choices to nodes that are computationally capable, e.g. the Pandaboard [7], than more common sensors, e.g. the IMote 2 [81].

In this section, we detail our research into suitable sensor hardware that met the following requirements:

1. Able to perform processing of images and metadata
2. Common interface availability (Serial, USB)
3. Wireless enabled
4. Battery-powered
5. Expandable memory

It should be noted that there are many devices out there that met the above criteria but we selected devices that had a large development community. At the time of writing, there are now many more, but this research was carried out in 2010, when the market for low-cost, high-power, small form factor single board computers (SBCs) was starting to grow. While we did research on many more devices, the three listed below

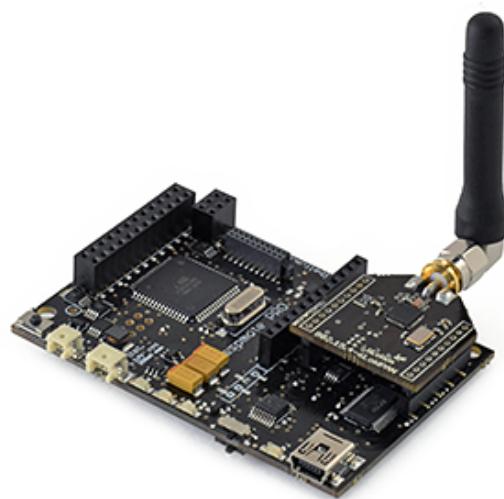
are the devices we physically trialled and tested. This section also details the modifications we made to any devices we tested in order to ensure they would survive in a humid environment.



(a) Pandaboard



(b) IGEP v2



(c) Wasp mote

Figure 3.3: Sensor Node Hardware

3.2.1 Pandaboard

Texas Instruments supported the development of a reference Single Board Computer (SBC) that had specifications similar to that of a modern smartphone (Figure 3.3a) and was capable of running desktop Linux, known as the Pandaboard [7]. A dual core 1Ghz ARM processor with 1GB of RAM, support for external storage, expansion ports, USB, Wi-Fi and Bluetooth in a board the size of two credit cards can be a powerful addition to a data heavy WSN, especially one that deals with images.

There is no mention of Pandabards in the literature being used in WSNs but the low power of the system and advanced capabilities make it suitable for processing and transmitting data simultaneously.

3.2.2 IGEP v2

The IGEP v2 (Figure 3.3b) is another ARM based SBC that uses a 1Ghz single core processor with 512MB RAM and similar connectivity features to the Pandaboard, but around half the size. This does result in a reduced power draw and the device is still capable of running desktop Linux.

Due to the smaller size, and easier commercial availability, the IGEP has been used as a sensor node to record, process and send readings from multiple devices, such as air temperature, GPS and oxygen saturation as part of environmental monitoring [97]. In their research, they found that the IGEP achieved 9.1 hours of uptime using a 4000mAH battery, a capacity used in many modern smartphones.

3.2.3 Wasp mote

The Wasp mote (Figure 3.3c) is a general purpose sensing board that is designed to allow plug-and-play connectivity for multiple sensor modules. The node can be programmed with C++, using the Standard Development Kit (SDK) developed by the

manufacturer and uploaded via USB. The processing power is not comparable to the more powerful SBCs, but the 600mAh battery is reported to last three months and the size is much smaller [68].

One of the primary benefits of the Wasp mote is that they are commercially available with an actively maintained programming environment.

3.2.4 Discussion

The capabilities of each of these boards vary greatly from running a subset of C on an embedded OS to running full desktop Linux. Each visit we made to Danai Girang allowed us to test these devices and determine their eligibility for nodes that would survive in the Malaysian rainforest and discover usable radio range data, which we discuss in the next section.

The trade-off with these devices is that the Wasp mote nodes, with limiting processing capabilities, is reported to yield a battery life of 6 months. Our tests yielded approximately 3 months, with minimal sleep scheduling, so this number does seem reasonable. However, the IGEP and Pandaboard are able to run any software that has been compiled for an ARM processor and do so within an unrestricted operating system. While this does allow them the flexibility to process data with the same power as a five year old desktop computer, our battery life tests drained the board within 6 hours when using 4 D cell batteries. We managed to increase this number by loading a power-efficient modification of the Arch Linux distro, removing any support for displays or running of a GUI, disabling unused ports and disabling the radios for as much of the uptime as possible. With these changes, we were able to achieve approximately three weeks of running with sleep periods.

Most SBCs, while having a reduced battery life, contain common I/O ports, such as: USB, VGA, Ethernet and Serial. This extensibility enables us to add additional wireless capabilities. For example, sensor nodes may require the use of a long range wire-

less medium for inter-node communication but the endpoint of a network uses Wi-Fi to allow users to connect to it and upload sensed data to the Internet; this can be done by adding a long-range wireless radio to an I/O port and providing software support.

3.3 Transmission Media

In this section, we explain the experiments we carried out to test the performance of different transmission media in the UK as well as the Malaysian rainforest. While range is the most important feature, a data rate that can handle hundreds of large readings in a day is a requirement, as wildlife cameras can generate ten images each time they are triggered and a high definition image can be up to 1MB in size. Our motivating scenario is focussed on the transmission of sets of 3 images (each around 800KB) for every trigger, where a sensor can trigger hundreds of times in a day.

3.3.1 Wi-Fi

Wi-Fi was already available on our initial test platforms and the high data rate made it suitable for sending a large volume of images in a short period. We knew that current cameras deployed in DGFC were up to 1km apart and we did not expect to cover that range completely, but we did expect to achieve that coverage by using intermediate nodes.

Research, outlined in Section 2.1.4, showed that the rainforest could reduce the range by up to 78% and the ideal maximum range of 2.4Ghz Wi-Fi is 100m [36].

We tested Wi-Fi range using two IGEP boards. We selected IGEP boards because they could be powered by 4 D Cell batteries and run a lightweight Linux operating system that allowed us to access information at the network layer, while still running higher level applications. The IGEP nodes we used did not have any additional hardware and the nodes were tested without the use of an external antenna. A Java application was

written to periodically scan for available networks and store those results in a text file. One IGEP board was set as the base station and attached to a tree, at the same height it would be if it was attached to a camera, and another was walked to specified points around the base station at defined locations. These locations were chosen to include as many distances as possible and as many different forms of obstacle between the searching node and the base station, such as: line of sight (LOS), medium vegetation or thick trees.

This experiment was run in a wooded area in the UK and in the rainforest at DGFC. The specified maximum range of 802.11g is 120m. When considering attenuation and obstacles we were expecting the signal to be reduced by up to 50% in the UK. However, we found that we received a maximum range of 30m, with LOS. Figure 3.4 shows the results we experienced, while testing in the UK, some of the drops in signal can be attributed to dense foliage and readings that were not LOS, but a maximum range of 31m, with an SNR of 29.5 dBm, is less than we expected, as the UK does not experience high humidity often.

The graph shows a drop at 22m: this was due to unusually dense foliage that restricted the LOS between the base station and the receiving node. With five runs of this test we observed the same results. The primary aim of this experiment was to show the viability of Wi-Fi and to ensure our application functioned as intended, which it did. Further experiments could have been run to remove the anomaly but the results of the experiments in Danau Girang were the focus.

Despite the poor range from the tests in the UK, it was consistent with other studies reporting signal degradation of up to 78% in areas with dense foliage. We visited Danau Girang in 2011 to gather the requirements of the network and ensure the hardware is able to survive the humidity. Range experiments were run in the rainforest to see if a more humid environment impacts range any further, Figure 3.4 shows this.

Figure 3.4 shows that the maximum distance to receive a signal is approximately the same in Malaysia as it is in the UK. There are more signal drops but this seems to

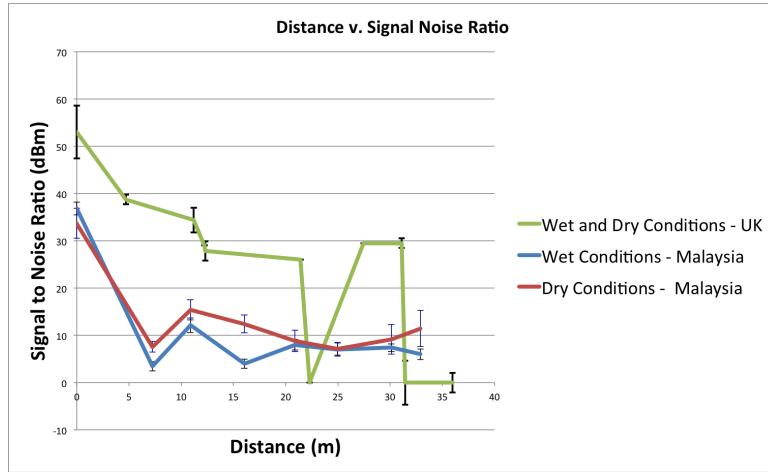


Figure 3.4: Signal-to-Noise Ratio for Wi-Fi

be due to denser foliage blocking the line of sight. However, it does suggest that the humid environment of the rainforest does not have a significant impact on the received signal. It is clear that the denser rainforest does impact the signal-to-noise ratio in a much shorter distance from the base station but a link is still made, allowing for a successful transmission of data.

Poor Wi-Fi performance led us to research alternative methods to increase the range without impacting the environment the network is to be deployed in. We considered using intermediate nodes, not attached to cameras, to account for the lack of range but, because some cameras can be up to 1km apart, we would need more than 30 nodes to create a connection between two locations.

We also researched wireless technologies that are more common in sensor networks. This does mean that the data rate is not as high as Wi-Fi and error correction in packet streams is not always as robust, but it is more suited to sensor networks, using less power and providing longer range.

Finally, we considered using the researchers or animals at Danau Girang, as ‘data mules’, creating temporary links between nodes while they are in the forest. However, the trip to Danau Girang yielded the information that researchers generally do not cover those distances in the forest and data delivery would be sporadic.

3.3.2 DigiMesh

Due to the poor range results of Wi-Fi, we created a second prototype of the network, using DigiMesh. DigiMesh is a proprietary wireless protocol, based on the 802.15.4 standard and designed for devices with limited power. Using the same frequency as Wi-Fi, DigiMesh has been reported to provide 7km of range, with a data rate of 250kbps.

In our prototype implementation, we are using Waspmotte sensor boards [68], a general purpose node that is capable of transmitting through various communication mediums. Our Waspmottes are provided with DigiMesh modules and a 2GB SD card to store sensed data.

When testing the range of the Waspmottes, we followed a similar method to that which is outlined in Section 3.3.1. One board is in a fixed location and running a C++ application to poll for nodes in the network. Once a node is found it sends a message to the node every 10 seconds. The second board is set to scan the network and receive packets as soon as a base node is found; this node is then moved to different locations.

The receiving node prints out variables related to the received packet, such as: Received Signal Strength Indicator (RSSI), source MAC address and packet ID. However, not all packets are received so the RSSI can display 0 if there are errors in reading or if packet collision occurs. We found this to affect the results and have just used the two nodes to identify the maximum distance they can be apart, while maintaining a stable connection.

Initial experiments were run in a moderately vegetated area in the UK which yielded 497m of range. Limitations with buildings preventing us from testing any further but the signal strength still proved to be strong.

The initial results for the range tests were positive and DigiMesh does seem to be a viable solution to account for the lack of range when using Wi-Fi. As the frequency is the same as 802.11g, thus licensing it for worldwide use, we expected similar results in Danau Girang.

Experiments were run in 2 areas of the rainforest around Danau Girang and the results yielded were not the same as we experienced in the UK; the thick vegetation had a significant impact on the range, reducing it by almost 50%.

In more open areas of the rainforest, we achieved 199m on average, more dense regions of the forest reduced this to 102m on average. While these results are not as high as we achieved in the UK, they are still suitable to use DigiMesh in the deployment of a WSN. The low range could be because we were using low-gain antennas and little configuration had been made on the DigiMesh radio.

3.3.3 Adapting and Optimising for Harsh Environments

All of the hardware that we used for experiments had their components exposed, as shown in Figure 3.3 and would have become compromised if moisture came in contact with them. To protect them from this, we used waterproof cases with a protective foam inside, known as Pelican cases, to keep the nodes watertight, but still allowing airflow to displace the heat generated, shown in Figure 3.5. External antenna can be fed through the lip of the case, using thin cable, ensuring that the range of the transmissions is not affected by the case.



Figure 3.5: Pelican Waterproof Case

After experiencing poor range results, we considered running antennae up through the trees to the canopy so there would have been no foliage to block signals and the humidity would have been lower. However, this was not possible because of conservation issues and the fact that animals could easily damage the antenna wire.

Typically, WSNs in urban areas can use mains power and more remote networks utilise solar, or other renewable energy sources. In an area of dense foliage and high canopy cover, solar power is not a viable energy source. While DigiMesh has a lower transmission rate than Wi-Fi, its power requirements are lower, which maximises network lifetime for WSNs that do not have access to constant power or a renewable energy source.

3.4 Software

In this section, we explain the image processing software we developed for the real time processing of images within the network, as well as the retrospective processing of images that have been taken during previous deployments.

3.4.1 Triton

Triton is a C++ program that makes use of the Open Computer Vision (OpenCV) library [21] for performing popular image processing functions [50]. Triton takes in a set of images, combines them and builds a Gaussian background model [128], a common background subtraction method commonly available in many computer vision packages. Using this model, the foreground is extracted from the image and this is scanned for objects. If an object is found, it is extracted to a new image and saved as a black and white template. This minimises the memory each template takes up but can still be used to assist with future classifications. An object identified from a set of images means that the set is then marked as interesting, i.e. it is believed to contain something. If no object is extracted, it is marked as empty. The template extracted can then be classified by a human and, in the case of our motivating scenario, those classified templates can be used with newly sensed images to determine a match and allow for classifications down to the species level.

Primarily, we developed this tool to detect the large number of empty images captured around Danau Girang, due to trees falling, movement of the Sun, fast animals or dirt on the lens. However, Triton proved to be quite effective when finding images of interest and it was modified to be used with our sensors. In this section, we explain how Triton was tested and compared with human observations. To ensure a valid comparison, we also compared both results to a random classifier.

During a typical three month deployment along the Kinabatangan River more than 40,000 images can be taken. A large proportion of these images can be classed as

false triggers; occurring when a camera is triggered by the motion detector but there is no object of interest. These can be caused by changes in the light, falling trees or animals moving too fast to remain in the shot in the delay between the motion detector triggering and the camera capturing an image.

We organised the images collected by the original camera that captured them and separating them into sets of three, because that is the number of images that the cameras were configured to take at each trigger. From this, we built a Gaussian background model [128] for each set and used the resulting model to detect objects in the foreground, classifying the detected foreground as the Region Of Interest (ROI), and extracting it.

Figure 3.6 shows a sequence of images taken by a Reconyx HC500 camera. The dynamic background and light levels make it very difficult to extract an object of interest. A background model is built from these images, which contains everything that is believed to be in the background. In this case, it should be the ground, the trees and the visible sky. The background is then subtracted from the image, leaving the foreground, and ROIs, larger than a user-defined threshold, are identified. The largest ROI is then extracted from the image and saved separately, shown in Figure 3.7.

From our visit to Danau Girang in 2011, we collected images from two different three month deployments, with the same cameras placed in different sites around Danau Girang for each deployment, giving us just over 70,000 images to test our approach. The images are sorted by the camera location and the date of collection. We process every 3 images as one set, building a background model of all three images.

We manually processed all of the images initially, marking images that are empty as false triggers. We then processed the images, using our application, and a resulting processed image is created from every set of 3 images. If nothing is detected in a set then no image is created and that set is logged as empty.

There are four classifications that can be made with images sets:

True positive: An ROI is extracted that contains the animal in the set.

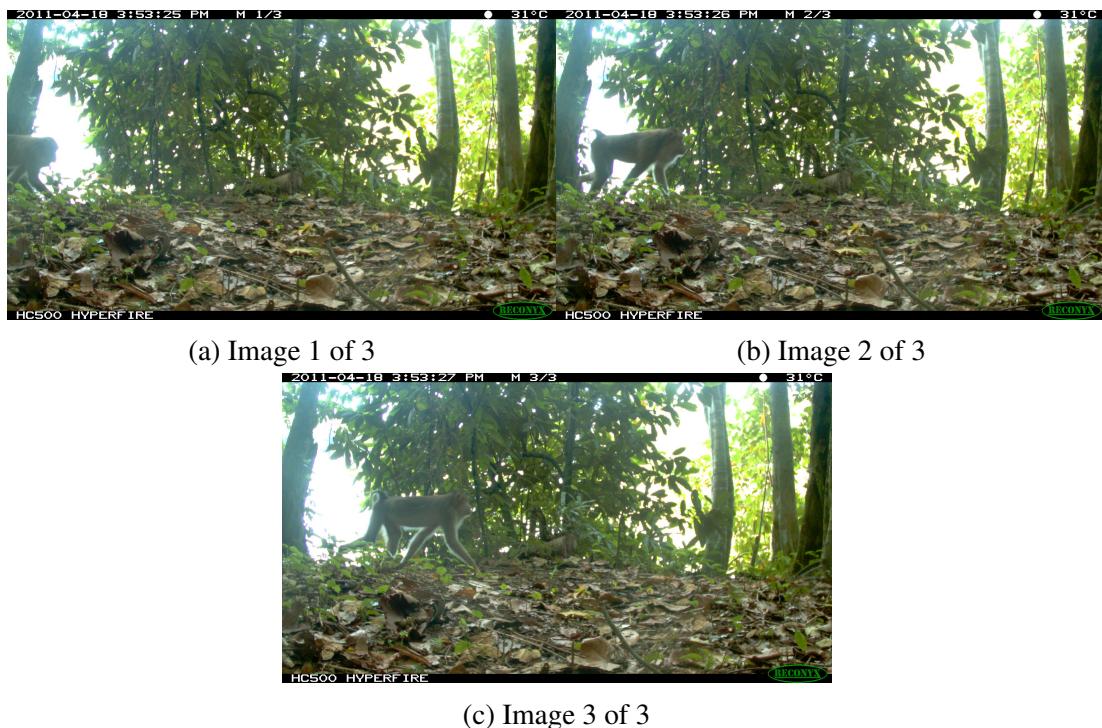


Figure 3.6: Sequence of 3 Images Captured on a Reconyx Trigger

False positive: An ROI is extracted that contains nothing of interest.

True negative: A camera is triggered with nothing of interest in the image and no ROI is extracted.

False negative: An image containing an animal has no ROI extracted.

The processed images are then compared with our manual findings. The accuracy of our application is calculated by the following equations:

$$Accuracy_{tp} = (Ex_{tp}/Ac_{tp}) * 100 \quad (3.1)$$

$$Accuracy_{tn} = (Ex_{tn}/Ac_{tn}) * 100 \quad (3.2)$$

Where Ex_{tp} is the number of true positive sets extracted, Ac_{tn} is the actual number of

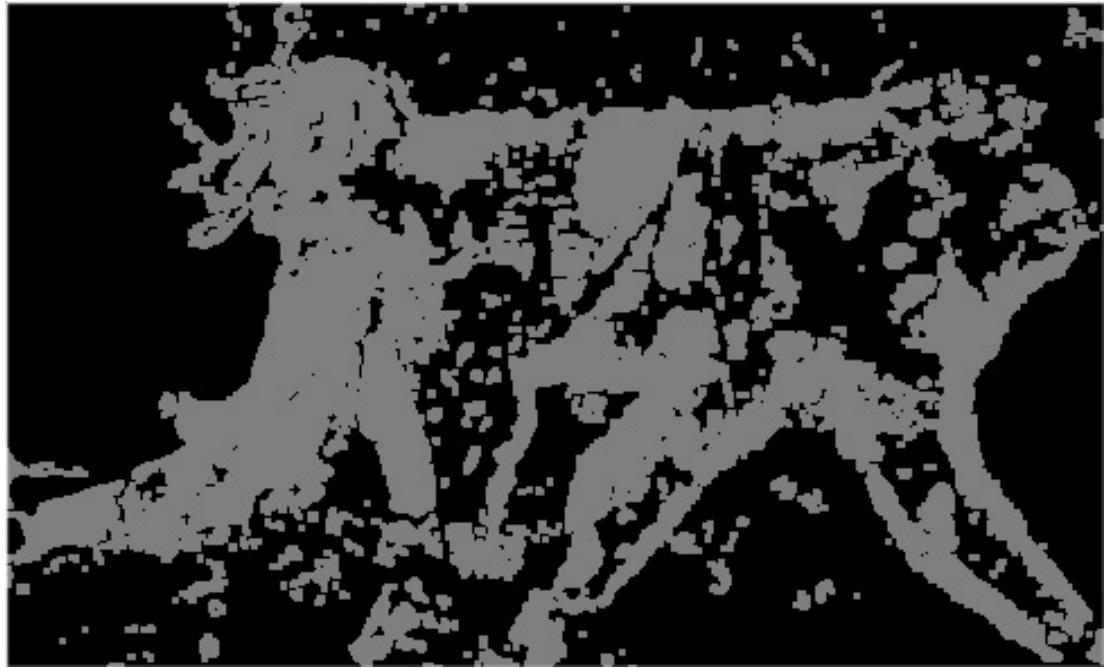


Figure 3.7: Resulting Processed Image from Original Set of 3

true positive sets, Ex_{tn} is the number of true negatives extracted and Ac_{tn} is the actual number of true negative sets.

Table 3.1 shows experiments run on images from a randomly selected camera deployed in Danau Girang, testing on 300 sets of 3 images. These images were manually processed and classified as interesting or empty, of the 300, 94 sets were identified as interesting. Using Triton, 77 of those interesting sets (true positives) were extracted, with another 17 false positives extracted. 201 true negatives were correctly identified and a further 5 were identified as false negatives. As Triton processed these sets, we tested the time it took to complete a set of images on a Pandaboard and we found the mean time to be 43 seconds.

Out of the 94 interesting sets, 77 were identified correctly, giving an accuracy of 82% for finding true positives (using Equation 3.1). Furthermore, of 206 empty sets, 201 were found, which gives a 98% accuracy at correctly identifying true negatives, using Equation 3.2.

True Positive	True Negative	False Positive	False Negative	Total Image Sets
77	201	5	17	300

Table 3.1: Classification Results

These preliminary results show that our method is effective at detecting images of interest (True Positives) and it appears that misclassification (False Positives) primarily came from black and white images taken at night, images where minimal movement of the animal has caused a trigger and images where an animal has caused a trigger but it has been too fast moving to be in the second two images.

After a longer deployment, we would be able to build up a more substantial background model to account for some of the animals being less dynamic in images and we expect this to decrease our error rate thus reducing the number of false negatives, although this is something that is yet to be implemented.

The results of Triton were compared to a random classifier, implemented in Python, that ran through a directory of images and generated a random number for each file. If the generated number was less than 0.25, then it was marked as true positive, between 0.25 and 0.49 marked it as false positive, between 0.5 and 0.74 was a true negative and between 0.75 and 1 was a false negative. Table 3.2 shows the outcome, with 72 *true positives* extracted, resulting in an accuracy of 80% (using Equation 3.1). 73.3 *true negatives* from a total of 206 were found, giving an accuracy of 36% (using Equation 3.2). Although this set does not compare to the number of images collected in a six month deployment, human classification is a time consuming process and to have 900 classifications is complex. Although these could be crowdsourced, the reliability becomes questionable, in part due to the specialist nature of the images. This does show that our approach is quite close to the accuracy of a human when finding images of interest, but is able to do so in a fraction of the time, processing hundreds of images every minute.

Triton, coupled with a set of human eyes at the heart of the network, should be an effective approach for prioritising images through the network, even when a classifica-

True Positive	True Negative	False Positive	False Negative	Total Image Sets
72	73.3	78.3	75.3	300

Table 3.2: Random Classifier Results

tion cannot be made through the use of local knowledge. More importantly, in an area where the environment is as dynamic as the Malaysian rainforest, 98% accuracy for detecting true negative images can be crucial, if only for saving network bandwidth.

Future Work

Triton is currently capable of processing a set of images and extracting the largest object of interest, if one is detected. The primary benefit of this is that it requires no training initially and works, with fairly good accuracy, from the time of deployment. However, it is possible that functionality could be extended by storing the extracted images and associating them with the actual content, i.e. animal name, future extracted images can then be matched to the templates, within a threshold, to assist with classifications within the network. Currently, the templates are extracted, stored and associated with their contents once classified by a human, but Triton does not use these. If Triton detected an interesting image, it could be envisaged that it could then search a folder of templates and use OpenCV to compare the images and provide a cursory species classification based on the closest match.

Extracted images are stored in black and white and are only a few kilobytes in size, this will be useful as we expect that a large number of templates would be required in order to accurately identify an object of interest. For example, in our motivating scenario, an animal could be any distance from the camera, its legs could be in different positions or its angle it faces towards the camera could be different, giving many possible images. This process could be optimised by only matching regions between extracted images, such as the head or body, but this would require further experimentation.

3.5 Conclusion

In this chapter, we have shown the current hardware choices available for more computationally capable sensors and detailed our experimental results on how rainforest environments impact wireless transmissions. Although wireless technologies such as Wi-Fi provide a high data rate, their range is limited and not suitable for sparsely located nodes covering a large area; especially in a humid environment. Commercial Wi-Fi solutions are becoming more widespread in open, public places, such as: shopping malls, airports and universities, but these places allow for multiple routers, constant power supply and a greater range as routers are placed as high as possible to maximise coverage. Newer wireless technologies, designed for long-range communication in sensor networks, are becoming increasingly more viable and, while they do have lower data rates and less robust protocols, they are more suitable for a resource constrained WSN that requires minimal power draw when transmitting. Using general-purpose hardware also means that there is no protection against water, humidity and animal or human intervention. We adapted existing waterproof cases to suit our needs and preliminary short term deployments showed that they are able to prevent moisture from entering the case. We have also developed an image processing application, Triton, that attempts to extract an object from a set of images that it believes to be the subject. Image sets that do not have any subject are classed as empty and should have no object extracted while those with a subject are interesting and should have an object extracted. Testing Triton on existing images at DGFC yielded 98% accuracy for detecting empty image sets and an 82% accuracy for finding interesting sets. A key feature of Triton is that it is not calibrated for a particular background or environment, so these results are especially promising.

Chapter 4

Wireless Sensor Network Architecture Design

In this chapter, we explain our proposed network architecture that uses local and global knowledge to make informed routing decisions and to classify sensed data within the network. Our approach, K-HAS, uses a three-tiered, hierarchical approach with each subsequent tier providing increased knowledge processing capabilities.

We aim to show that sensors capable of processing knowledge will provide a more efficient network and be able to prioritise the delivery of interesting sensed data. We also believe that human input is a valuable way for such a network to learn about classifications it has made correctly, and incorrectly, and use that knowledge to inform future classifications. To prove this, we have developed an architecture that uses different levels of knowledge processing throughout the network, the Knowledge-based Hierarchical Architecture for Sensing (K-HAS).

The rest of this chapter is structured as follows. Section 4.2 outlines the main aims of K-HAS and what it is capable of that typical sensor networks are not. Section 4.3 details the software and data standard choices we made for our implementation of K-HAS. Section 4.4 provides a sample walk-through of K-HAS, relating to our motivating scenario and Section 4.5 concludes the chapter.

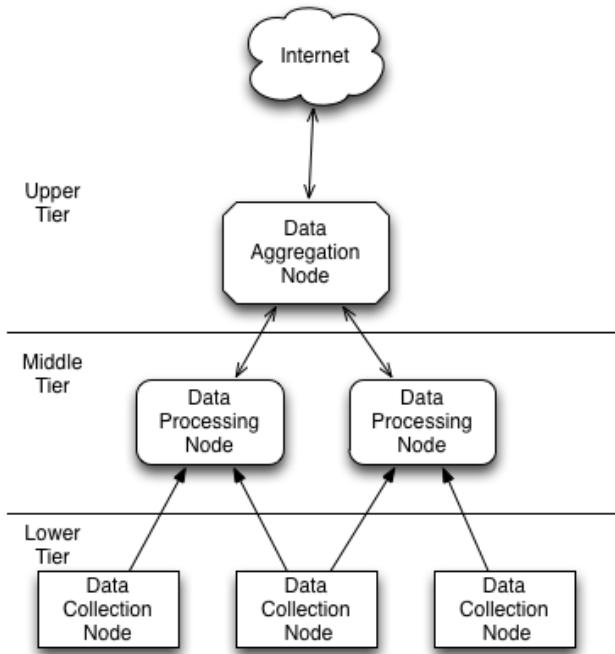


Figure 4.1: Overview of the K-HAS Architecture

4.1 The Local Knowledge Problem

While large in size, it is typical that a WSN would use low cost nodes with limited power, memory and computational capabilities; causing them to lack the ability to be aware of their surroundings or the data that they are sensing [15]. This means that, unless fixed by a routing protocol or a technician, data is delivered on a chronological basis and is then filtered at the base station, usually manually. Some WSNs store all of the data on the node and users of the network use a ‘pull’ model to query for data from nodes [99], but this requires some technical knowledge and, although it does increase the battery life of the nodes, it is a manual process again.

The environment that a WSN is deployed in is often a rich source of data to be sensed (such as inside a bird’s nest to sense temperature, humidity and movement), which often contains patterns that can be used to improve the performance of the network. For example, if a node knows that it has only been triggering between the hours of 6pm

and 5am for the past few weeks, it can then learn to enter a deep sleep outside of those hours or use that time to transmit data it has been storing while it assumes it will be inactive, based on previous days. Alternatively, this knowledge can be used to prioritise data throughout the network so that the most important data is received first, instead of the most recent. An example of this could be two camera nodes deployed facing the entry and exit of a building, tasked with looking for intruders between 5pm and 8am. If the camera facing the exit is triggered at 5:01pm and the camera on the entrance is triggered at 5:05pm, then the knowledge that the security guard leaves through the exit between 5:00pm and 5:08 pm will allow the entrance camera to prioritise its capture as more important, as it is an irregular occurrence.

This knowledge can be categorised as either *local* or *global*. Local knowledge is the knowledge of an area that has been gained through experience or experimentation [61]. For example, a native to the Amazon may know that three of the locations in which the nodes are to be deployed are flooded for two weeks of the year, rendering their readings useless for that time period and increasing their risk of failure. This is local knowledge, as it cannot be gained without experiencing the flooding in that area, or measuring local water levels.

While we believe that we are the first to use the concept of local knowledge within the wireless sensor network domain, the terms have been around for many years. In 1999, a book that referred to local knowledge as *indigenous knowledge* defined local knowledge as ‘systematic information that remains in the informal sector, usually unwritten and preserved in oral traditions rather than text’ [102].

Over the past twenty years, local knowledge has been used in various contexts, from researching lending and the credit market [104] to extracting local knowledge from natives to improve farming techniques [35]. This research, as well as work that will be covered later, showed us that there are two kinds of knowledge: global and local.

It was from agriculture research that we were able to refine our definition of local knowledge, [61] defines local knowledge as ‘knowledge that farmers have derived locally

through experience and experimentation'. They also say that indigenous knowledge is different in that it is culturally specific. From this definition, as well as our work with our motivating scenario, we were able to generalise the definition and expand upon it.

We now define local knowledge as *knowledge of an area, held by a domain expert, that has been gained through experience or experimentation*. The weather of a region is global knowledge because it can be found through a variety of media, whereas the saturation levels of the soil in a field would be local knowledge as it would require experimentation to gain that knowledge.

Using this definition, we believe that **utilising local knowledge within a sensor network can inform routing decisions to make better use of the bandwidth in resource-constrained WSNs by sending data that the node believes to be important first, rather than chronologically**. Patterns in the data, and knowledge of the environment surrounding a node, can allow a node to infer what data may be valuable, automate the processing of adding context to the sensed data, learn from previously sensed data and utilise knowledge of ongoing projects within the network to determine what data is thought to be of a higher priority.

To show this, we have developed a network architecture for WSNs that utilises knowledge from the data it senses, as well as its deployed environment. It is called the Knowledge-based Hierarchical Architecture for Sensing (K-HAS) and this thesis will show how K-HAS addresses the problem of delivering the most important data first and improving the overall efficiency of the network.

4.2 K-HAS

K-HAS has been designed as an architecture for WSNs that addresses the above problem by using knowledge to classify sensed data and adapt to changes in the structure of the network. By pushing knowledge bases out to the edge of the network, all nodes in

the network have some awareness of the data they are sensing, as well as how important it is, based on the current projects that the network is involved in. This is achieved by using rules with different levels of granularity based on the knowledge processing capabilities of that tier. Figure 4.1 shows a high level overview of K-HAS, showing the flow of data from each tier.

4.2.1 Data Collection

The data collection (DC) tier is very similar to sensor nodes commonly used in a WSN, using hardware that has similar, limited processing power and storage to nodes such as: the I-Mote and Wasp mote. These DC nodes are deployed at the edge of the network and tasked with sensing their environment, pre-processing the sensed data and using each other to relay data to the next tier.

DC nodes are capable of performing processing on data, such as the time it was recorded and its size, but their limited knowledge processing capabilities allow them to have an increased battery life and reduced size, making them suitable for a variety of deployments as they can be integrated easily with existing devices (i.e. cameras) and be easily disguised, or hidden.

Knowledge Base

Reduced knowledge processing capabilities and low memory restrict the knowledge that these nodes can hold and they are limited to a static knowledge base that is encoded at the time of deployment. DC nodes run an operating system designed for embedded devices, the size of the operating system is minimal when compared to an operating system used on modern desktop computers and provides much more limited functionality, such as sleep scheduling, basic file system access and the capabilities to run a variation of popular languages.

For example, TinyOS [66] is an operating system designed for sensor nodes and uses a dialect of the C language, called nesC, that is optimised for the memory constraints of most sensor nodes. The operating system exposes components that allow basic commands, such as reading from the sensor interface or sending a message. Commands are executed as a request to perform some action and events signal the completion of that action. Intensive tasks, such as data processing, can be scheduled to run at a later time, ensuring that the low-power node remains responsive and preventing any calls from blocking other events.

The Wasp mote nodes, explained in Section 3.2.3, do not run any operating system and any event handling, networking or memory management is performed solely by the single file application that is uploaded to the node through the boot-loader. Performing intensive tasks, such as local data processing, can drain the battery of the node, cause it to fill the available memory or even cause the node to go down. DC nodes do not have the safeguards in place that most operating systems and applications do, operational code must be lightweight, responsive and ensure that events do not run for longer than intended.

Some nodes within this class do not have the capability to process text files within the file system and, therefore, must have the knowledge added to their operational code. DC nodes only perform simple operations on the properties and content of the data that they sense, such as the time it was recorded, the location and its size. For more complex data, such as images and video, DC nodes do not possess the computational power required to process them and instead use the metadata associated. Unlike modern rule engines, these static rules do not use forward chaining and the outcome of one rule does not cause the rules to be fired again. Listing 4.1 shows an example of some of the rules in the knowledge base and highlights that the core structure of this file is simply a list of if statements that can be executed one after the other. Forward chaining is not used here because processing times need to be kept to a minimum and a chain of rules being fired could easily overwhelm the node. If any rules are fired that suggest

the observation may be interesting, it is prioritised through the network and potential classifications are encoded within the observation before it is sent on.

```

1 if month of observation is JUNE AND time of observation is between 17:00 and
   19:00 and active otter project is TRUE
2 add classification to observation('Potential Otter sighting')
3 prioritise observation as interesting
4 if temperature is 37 AND time of observation is between 01:00 and 05:00 and
   active leopard project is TRUE
5 add classification to observation('Potential Leopard sighting')
6 prioritise observation as interesting

```

Listing 4.1: Pseudocode DC Node Rules

When the data is recorded by the DC node, the knowledge base is fired and inferences are made about the contents of the data. Each DC node has a static knowledge base loaded onto it before it is deployed, which is based on the local knowledge of the area that it is to be deployed in. For example, a node deployed on the bank of a river would have a different knowledge base to a node deployed in the fields of a plantation.

Once a trigger has been processed, the data is packaged and then sent on to the Data Processing (DP) node.

4.2.2 Data Processing

DP nodes act as cluster heads of the network, serving a subset of all deployed DC nodes. When data is sensed, it is forwarded through all DC nodes to the DP node that is tasked with serving the originating DC node. These nodes have more knowledge-processing capabilities than a DC node and do not perform any direct sensing.

Due to the greater capabilities, DP nodes have a much shorter battery life than DC nodes and a network typically consists of fewer DP nodes. These capabilities, such as increased memory and higher processing power, allow DP nodes to run a rule management system (such as Drools [89]), that would not be possible on nodes that run on an OS for embedded systems, that is able to handle more complex rules and process more than just files, they can perform the same tasks as most modern computers, such as: image processing, audio processing or reading metadata from files that requires extra libraries. When a DP node receives data, it processes everything associated, this

includes metadata, the data itself and the inferences made by the DC node. If the DC node has inferred that the data is of a higher priority, then this data is processed first. This is done by prioritising the data at two stages: once it has been received and when it is about to be sent.

In our current design, DP nodes use two different radios, a Zigbee radio to allow long range communication from DC nodes and a Wi-Fi radio that provides short range communication that allows for higher data rates.

Knowledge Base

In our motivating scenario the network is image-based, this means that the DP node would perform image processing, as well as processing the image metadata. The increased knowledge processing capabilities allow DP nodes to run rules dynamically, learning from the sensed data and providing classifications that change based on changes in the environment. For example, if a DP node has not seen an elephant before, and it is not aware of the object in the image, then it will await a human classification. The node will then record the time period that it receives elephant pictures, i.e. June to July, and become more alert the following year. Similarly, the node will know not to look for pictures of nocturnal animals during the day. This local knowledge allows processing power to be saved and, thus, time; this ensures that the processing of sensed data is optimised as much as possible in order to reduce the time it spends in the network.

As previously mentioned, however, species level classification has not yet been implemented so DP nodes currently use image processing to determine if an image is interesting or empty and relies on human classifications at the endpoint to store the species.

The rule engine used should allow for rules to be dynamically inserted into the rule base, so that rules can be updated through network communication.

Upon receiving sensed data from a DC node, the rule base is fired on the meta data of each file received. If the rules determine that the data is of interest or, in the best case scenario, provides a classification, then the data is packaged and sent on to the Data Aggregation (DA) node.

4.2.3 Data Aggregation

Placed at the root of the network, these are nodes with the same knowledge-processing capabilities as DP nodes (although they typically have greater memory, processing power and a continuous power source) and would be accessible by users of the network. When DA nodes receive sensed data, it is unpacked and stored with a link to the node that the data originated from. Compared with a standard WSN, these nodes can be compared with a base station, or endpoint.

Any information added by the DP node is parsed and classifications are extracted. If a classification is found, it is stored and the DA node checks for any active projects that contain the classification. If a match is found then all users involved in the project are informed via their preferred method of communication. Using the motivating scenario as an example, the people involved with projects could be researchers and professors and they may be looking for images of leopards, requesting to be informed via Twitter.

All sensed data received, regardless of whether it has been classified, is accessible through a web interface hosted by each DA node. The interface shows all of the sensed data from each deployment, along with the associated classification. More importantly, it allows users to classify the data using a voting system. Users have roles which give them different privileges within the system. Normal users are able to vote and the majority vote is seen to be the current classification. While this is not a necessary feature of a DA node, it does allow the network to utilise the knowledge of human experts to inform future classifications.

However, privileged users are able to confirm a classification and prevent any further

votes. Once a classification has been confirmed, it is then sent back to the DP node it originated from. If the classification made by a user is different to the one inferred by the node, then it updates its knowledge base and acknowledges receipt.

This section of K-HAS is yet to be implemented but we believe that this system will be vital in the early stages of deployment as it will be used to build up a knowledge base. The more user classifications there are, the more accurate the network will be in the future. After a few months of classifications, K-HAS would then be able to use the knowledge base to make more informed classifications, requiring fewer classifications/confirmations from users. The exception to this is sensed data that cannot be matched to the knowledge base.

Data aggregation nodes should not be hampered by limited battery life that deployed nodes would experience as we expect them to be placed in a base station with power availability and access to the Internet. Therefore, DA nodes would typically be desktop computers with a constant power supply.

Knowledge Base

DA nodes do not typically experience the resource constraints that DC and DP nodes must compensate for. Because of this, they hold a global knowledge that contains a history of all observations made by all nodes, as well as the location and deployment times of all nodes in the network. When nodes are deployed, this is updated, by users, to show the location of the node and when it was placed.

Every classification sent from DP nodes is stored in a database and contains all information about an observation: date, time, originating node, route taken within the network, location, sensed data, classification of sensed data. They also hold information for all projects running, for example: a project to track elephant movements within the forest. When sensed data that could be related to the project is received, such as an image with an elephant classification, users associated with that project are informed

through email, or other means.

While DC do not store any of the observations they capture, and DP nodes only store part of the observation that can be used in future classifications, DA nodes store the complete observation made by every DC node, as well as any extra knowledge that is added by users upon receiving the sensed data. What is stored by the the DP node is dependent on the deployed purpose of the network. For example, in our motivating scenario, we store: the processed image, the classification, the node it originated from and the date and time it was captured.

As well as this, the functionality of DA nodes can be extended to provide administrative operations on the network, such as the recording of node locations, time of deployment and viewing all active nodes. This allows the DA node to monitor active nodes and alert users if a node has not sent any data in a while. The longer a K-HAS network runs, the more knowledge a DA node gains, both from users that classify observations and from changes in the sensed data. This knowledge is then relayed back to DP nodes, updating their knowledge bases as to what classifications were correct and which need updating for future observations.

Feedback Loop

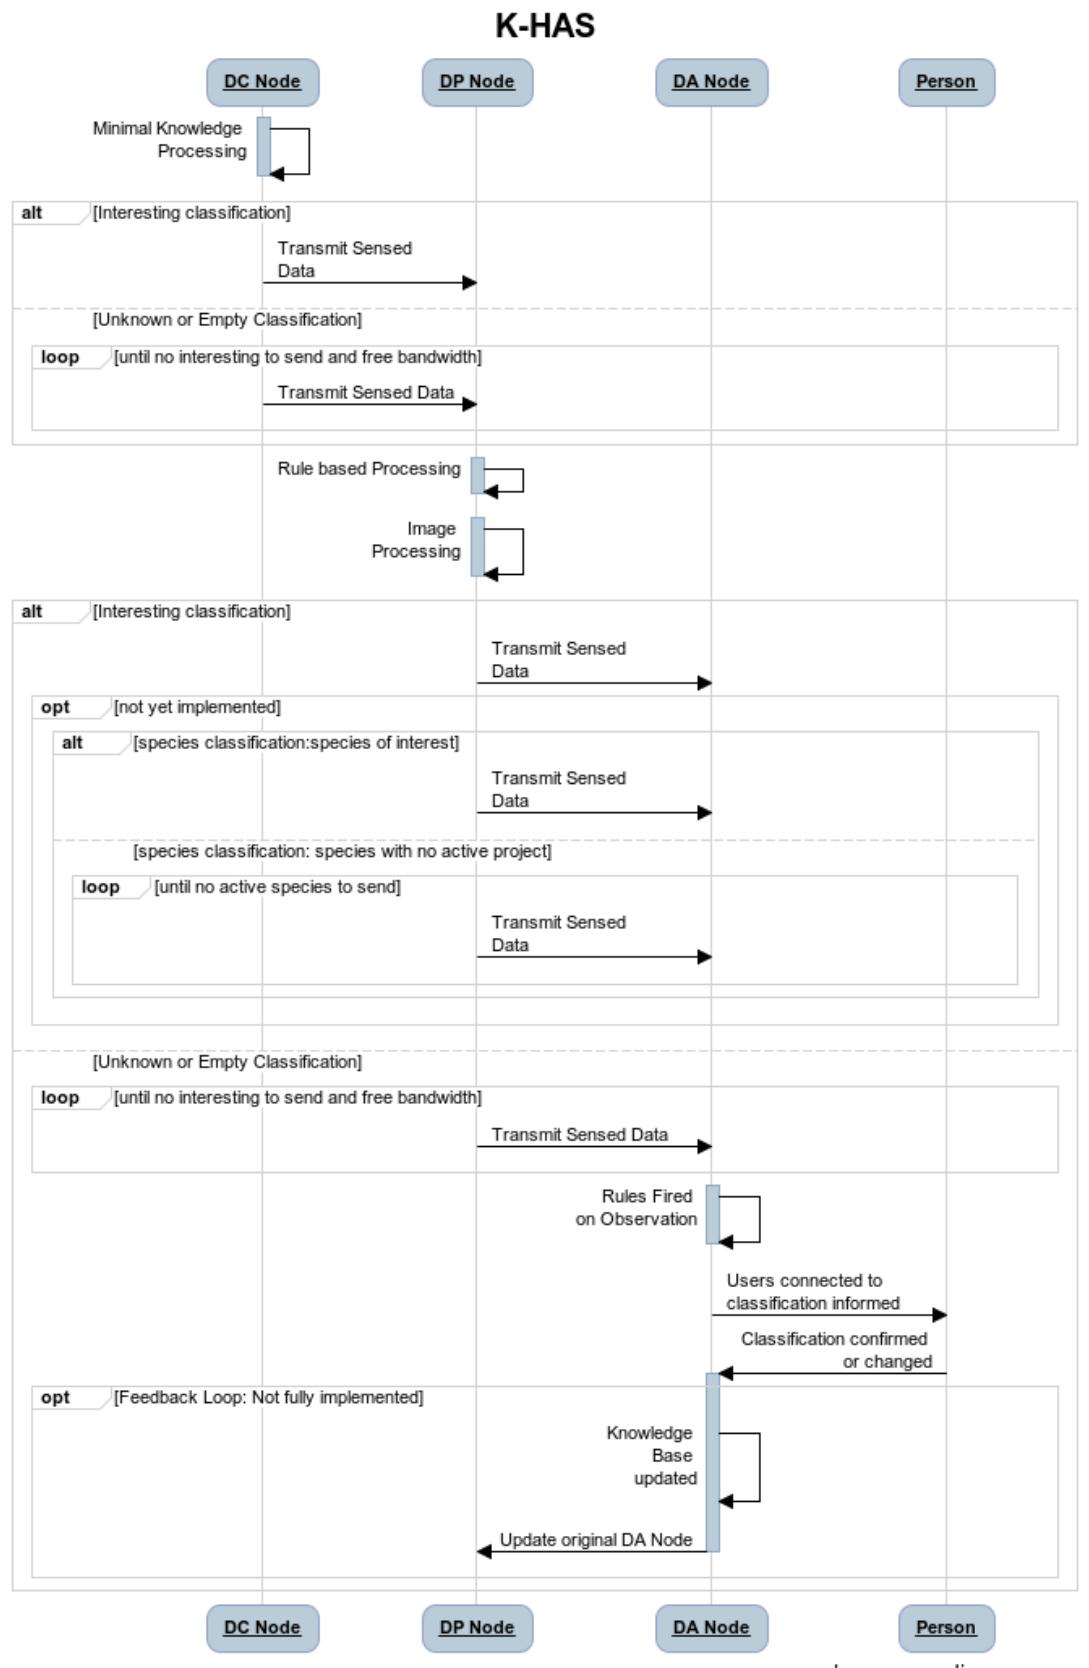
The feedback loop is a protocol within K-HAS, that uses human input, and other sources, to update DP nodes. When a DP node classifies sensed data, it stores some of that to assist with future classifications. What is stored depends on the type of sensed data. For example, images would mean that the DP node would store the resulting processed image, its classification and information about the time it was taken, the camera that took it and the location. When that is sent to the DA node, a human would then look at the image and mark the classification as correct, or modify it if it was not. Once that classification has been finalised, the DA node sends either a confirmation or a modification to the DP node that sent it, updating its knowledge base. This protocol seeks to reduce the number of incorrect classifications the longer the network is de-

ployed and allows the nodes to be dynamic and ‘learn’ throughout the lifetime of the network. It also allows nodes to adapt quickly to new data, if a DP node is unable to classify an image of an animal it does not yet have template images for, the knowledge of a human expert can provide those templates and the feedback loop will deliver that knowledge.

In the current design of K-HAS, the feedback loop has only been tested in a development environment and has not been fully implemented, we believe that the actual implementation of the feedback loop will allow the network to adapt dynamically to changes during its deployment.

4.2.4 Sequence Diagram

In figure 4.2, we have outlined the interactions that take place between nodes when an observation has been captured. Some parts of the architecture have not been fully implemented and those are highlighted throughout this chapter. This diagram provides a visual representation of how the tiers of K-HAS interact as well as how K-HAS prioritises sensed data and how K-HAS could utilise knowledge of professionals to improve the accuracy of its classifications.

**Figure 4.2: Sequence Diagram for K-HAS upon Capture of an Observation**

4.3 Technological Components

In this section, we describe the technologies used in our designs of K-HAS and how they integrate in order to use local knowledge based on their respective knowledge processing capabilities. The majority of components, both hardware and software, used in K-HAS are used so they are applicable for any WSN, but some choices have been made to remain in line with our motivating scenario and, thus, are more specifically suited for the capture of scientific observations.

4.3.1 Data Standard

To pass sensed data through the network, we first had to choose a standard format that would allow us to encode the sensed data, as well as enrich it with inferences made through processing. Darwin Core (DwC) is a body of standards with predefined terms that allows for the sharing of biodiversity occurrence information through the means of XML and CSV data files [119].

The Global Biodiversity Information Facility (GBIF)[5] indexes more than 500 million Darwin Core records published by organisations all over the web, allowing datasets that were previously siloed from the public to be accessed by both human and machine. The primary purpose of Darwin Core is to create a common language for sharing biodiversity data that is complementary to and reuses metadata standards from other domains wherever possible [119].

DwC Archives follow a star file structure, where a record can contain many occurrences, which is the recording of a species in nature or in a dataset. In an occurrence, there is an *event*, a recording of a species in space and time, enriched with other terms such as *identification* and *location*. DwC is the standard set of terms that can be used, while a Darwin Core Archive (DwC-A) provides the structure for data recorded using these terms. The core files in a DwC-A are:

1. EML.xml

2. Meta.xml

3. Data files

While DwC does not have the extensions available to OBOE, an extensible base ontology designed for ecological observations that is explained in detail in section 5.1.1, it is extremely concise for recording observations within the biological diversity domain and aims to be a standard reference for sharing these observations.

The record shown in figure 4.3 represents a DwC-A that conforms to the star schema. The ecological metadata language (EML) document contains all of the details about the project, such as who is involved, the institution code, contact details and the project(s) related to the observation. Listing 4.2 shows a fragment of the EML file and a complete DwC archive can be found in Appendix A.

```
1 <alternateIdentifier>e71fdalc-dcb9-4eae-81a9-183114978e44</alternateIdentifier>
2 <title>Images from Danau Girang during the PTY Project 2011-12</title>
3 <creator>
4   <individualName>
5     <givenName>Christopher</givenName>
6     <surName>Gwilliams</surName>
7   </individualName>
8   <organizationName>Cardiff University</organizationName>
9   <positionName>PhD</positionName>
10  <address>
11    <city>Cardiff</city>
12    <administrativeArea>Cardiff</administrativeArea>
13    <postalCode>CF24 3AA</postalCode>
14    <country>Wales</country>
15  </address>
16  <phone>(+44) 2920 123456</phone>
17  <electronicMailAddress>C.Gwilliams@cs.cf.ac.uk</electronicMailAddress>
18 <onlineUrl>christopher-gwilliams.com</onlineUrl>
19 </creator>
20 <pubDate>2012-07-26</pubDate>
```

Listing 4.2: Darwin Core Ecological Metadata File Fragment

The descriptor file is an XML document that contains the column headers in the attached files and the mappings of those headers to DwC terms, an example file is shown in Listing 4.3. This file shows that this is an archive containing the sighting of an individual, and that the sighting has been split into two files. The largest benefit of Darwin Core is its modularity. Extension files can be added to enrich the data for each

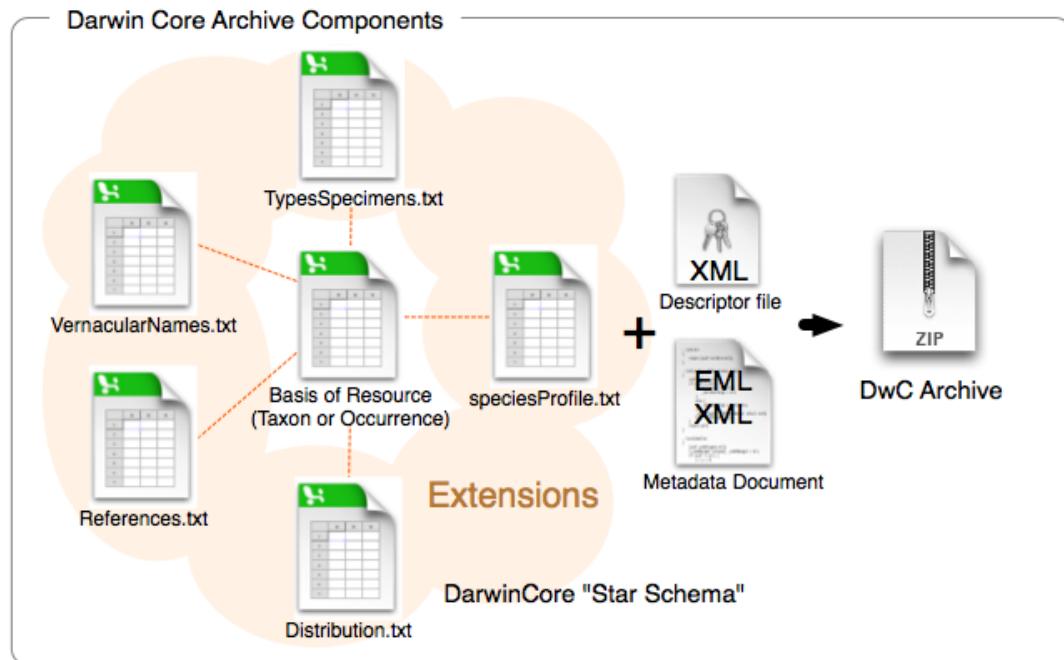


Figure 4.3: A Darwin Core Star Archive [119]

occurrence. In this example the extension file is named as *images.csv* and contains image-based evidence to support the observation.

```

1  <?xml version="1.0" encoding='utf-8'?>
2  <archive xmlns="http://rs.tdwg.org/dwc/text/" metadata="eml.xml">
3      <core encoding="UTF-8" linesTerminatedBy="\n" fieldsTerminatedBy="," fieldsEnclosedBy="" ignoreHeaderLines="1" rowType="http://rs.tdwg.org/dwc/terms/Occurrence">
4          <files>
5              <location>set.csv</location>
6          </files>
7          <id index="0"/>
8          <field index="0" term="http://rs.tdwg.org/dwc/terms/eventID"/>
9          <field index="1" term="http://rs.tdwg.org/dwc/terms/basisOfRecord"/>
10         <field index="2" term="http://rs.tdwg.org/dwc/terms/recordedBy"/>
11         <field index="3" term="http://rs.tdwg.org/dwc/terms/eventDate"/>
12         <field index="4" term="http://rs.tdwg.org/dwc/terms/locationID"/>
13         <field index="5" term="http://rs.tdwg.org/dwc/terms/scientificName"/>
14         <field index="6" term="http://rs.tdwg.org/dwc/terms/identifiedBy"/>
15         <field index="7" term="http://rs.tdwg.org/dwc/terms/dateIdentified"/>
16     </core>
17     <extension encoding="UTF-8" linesTerminatedBy="\n" fieldsTerminatedBy="," fieldsEnclosedBy="" ignoreHeaderLines="1" rowType="http://rs.gbif.org/terms/1.0/Image">
18         <files>
19             <location>images.csv</location>
20         </files>
21         <coreid index="0"/>
22         <field index="1" term="http://purl.org/dc/terms/identifier"/>
23     </extension>
24 </archive>
```

Listing 4.3: Darwin Core Descriptor File

Listing 4.4 shows the comma-separated central file (Basis of Resource) containing the

core details of the observation, i.e. the animal observed, and the column headings map to the descriptor file. Other linked files are typically linked by the unique ID of the observation, containing information that extends the observation and provides further context.

```
1 eventID, basisOfRecord, recordedBy, eventDate, locationID, scientificName,  
   identifiedBy, dateIdentified  
2 '1', 'MovingImage', '1', '2012-07-26 14:21', '1', 'Neofelis Nebulosa', '  
   Camera3', '2012-07-26 15:06'
```

Listing 4.4: Darwin Core Occurrence Data

All of these files are then archived and sent as a ZIP folder throughout the network. If the sensed data is media based, then the media is included as well. Software libraries to process DwC archives are included on both DP and DA nodes.

Darwin Core is suited to K-HAS because its use in ecological observations matches our motivating scenario and the archive can be easily created by a DC node, as it does not require any heavy processing and all of the files are commonly used formats.

EXIF

EXIF (Exchangeable Image Format) tags are written to images at the time of capture. Examples of these tags can be time, date or camera serial number. The capabilities of the camera do affect how detailed the EXIF is, for example, a camera with GPS capabilities will enrich the image with the location.

Wildlife cameras have more functionality than common digital cameras, with details like moon phase, temperature and/or GPS location. Some devices even include the saturation, brightness and hue of each image. These capabilities allow EXIF to be very detailed and can be used to find patterns in pictures that, when accompanied by local knowledge, assist with the classification of whether sensed data contains anything of interest. Figure 4.4 shows a snippet of the EXIF tags extracted from a Rekonyx camera.

EXIF — this group of metadata is encoded in 925 bytes (0.9k)

Exif Image Size	1,920 × 1,080
Exif Version	0220
Components Configuration	Y, Cb, Cr, -
Flashpix Version	0100
Color Space	sRGB
ISO	800
Exposure Time	1/30
Maker Note Reconyx	(713 bytes binary data)
Resolution	72 pixels/inch
Y Cb Cr Positioning	Co-sited

MakerNotes

Maker Note Version	0xf101
Firmware Version	3.0.1
Firmware Date	2010:05:18 4 years, 10 months, 25 days, 8 hours, 46 minutes, 7 seconds ago
Trigger Mode	Motion Detection
Sequence	2 of 3
Event Number	50
Date/Time Original	2012:04:16 05:48:17 2 years, 11 months, 26 days, 2 hours, 57 minutes, 50 seconds ago
Reconyx 0x0011	2
Moon Phase	Old Crescent
Ambient Temperature Fahrenheit	75 F
Ambient Temperature	24 C
Serial Number	H500BD09009453
Contrast	128
Brightness	0
Sharpness	32
Saturation	0
Infrared Illuminator	On
Motion Sensitivity	100
Battery Voltage	6.57 V

Figure 4.4: EXIF Tags from an Reconyx Camera

4.3.2 Middleware

The knowledge-processing capabilities of DA and DP nodes are the same and this is part of what makes K-HAS different from most other WSNs; both types of node run the sensor middleware, but each for different purposes. DA nodes use the middleware for administrating the network, receiving and archiving sensed data and allowing users to provide detailed classifications. DP nodes use it for the receiving, sending and controlling the flow of processing of sensed data before it is passed on.

Existing suitable middlewares have been detailed in Section 2.3.2 and our requirements for K-HAS were partially determined by the expertise of the users in our motivating scenario. Below is a list of our three core requirements:

Portability Heterogeneous WSNs utilise nodes with different architectures and capabilities, if middleware is to be used on the nodes it must be able to run on these varied devices.

Usability Users of K-HAS should not be expected to have knowledge of computer science or the underlying architecture, this network should be usable by almost anyone. The same must be said for the middleware as well.

Extensibility A closed-source middleware can be used, but it must then support all sensor nodes and data types, as well as receive regular updates. Open-source, or extensible, middleware can be used to add support for newer nodes.

GSN is a Java-based open-source middleware. New generic sensors can be added through XML files, while more complex sensors can be added through custom Java classes. GSN is covered in more detail in Section 2.3.2. Because GSN can run on any architecture that supports the Java Virtual Machine (JVM), it meets our portability requirements and the web interface to provide administrative functionality makes it usable by those without any domain knowledge. Finally, the ability to add new sensors

through XML means that it can be extended by almost any user of the network with very little guidance.

4.3.3 Knowledge Capture

GSN is packaged with a web interface that allows users to see all nodes deployed and view the latest sensed data received. The web interface is targeted towards users with domain expertise and has limited functionality focussed towards sensor administration. However, it does make GSN accessible by more than one computer, as well as a variety of different architectures. For example, the admin webpage could be accessed on the machine that runs GSN, or from a tablet computer connected to the same network. We used the same approach to develop a web-based tool that provides access to all sensed data, as well as a simple interface for performing tasks, such as uploading new rules or updating the location of nodes.

All sensed data is read from a MySQL database and users can view the metadata from each observation, such as location, date, time and temperature, as well as the data itself. From this, users are able to classify the data based on their role. Shown in the ontology in Chapter 5, K-HAS uses roles to control active projects and classifications; there are administrators and researchers. Researchers are involved in projects and receive notifications when relevant data has been received. They have access to the web interface and can vote on species classifications for sensed data. Administrators lead projects and can create/complete them, but they also have the ability to finalise classifications. K-HAS follows a knowledge hierarchy (Figure 4.5), with administrators at the top and DC nodes at the bottom. While there are more DC nodes in the network, their knowledge bases are more limited and classifications are trusted less than classifications made by DP nodes. Although there are fewer DP nodes, their knowledge bases are more detailed. Researchers and admins do not share a level on the knowledge hierarchy because we assume that administrators would be more experienced domain experts. For example, researchers could be students at Danau Girang, whereas

an admin would be a professor with more experience. When an administrator makes a classification, all prior classifications are ignored and the feedback loop protocol is used to update nodes.

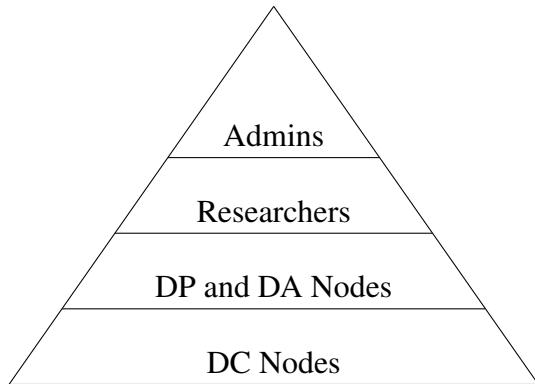


Figure 4.5: Knowledge Hierarchy for K-HAS

Figure 4.6 shows an observation where users can vote on the contents. An administrator can then confirm that classification and prevent further votes from being cast. This type of moderation means that it does not have to be specialists voting on sensed data that cannot be classified by DP nodes.

Viewing data for new observations is useful for gaining classifications and alerting members of a project, but viewing older sensed data allows patterns to be identified in order to create new rules. Figure 4.7 shows a map of all deployed nodes in the area surrounding the field centre in our motivating scenario. When users select a node, a table is populated with all of the classified observations that it has captured; this can then be used to extract patterns from the data and create rules. For example, the three observations of the Malay civet are only seen late at night, if further observations also showed this, then we could create a rule defining the active hours of the Malay civet and, potentially, list days that it is likely to pass. These rules can then be written and uploaded to the knowledge base.



Figure 4.6: Web Interface for Observations

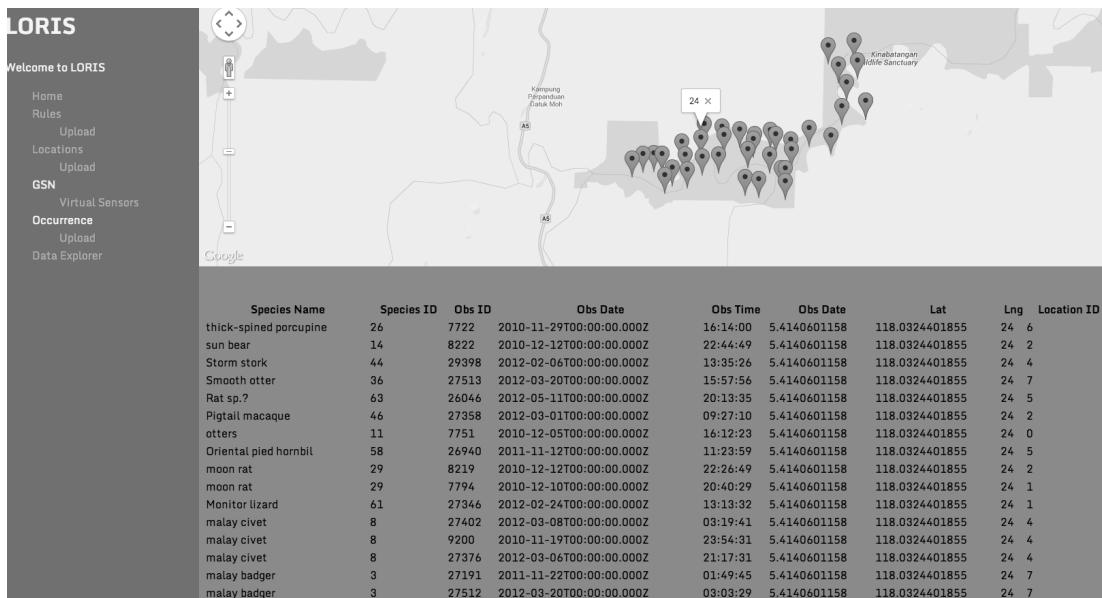


Figure 4.7: Web Interface for Classified Sensed Data

4.3.4 Knowledge Base

The Drools rule engine is a Java-based engine that uses forward chaining inference for the processing of rules [89], which means that rules are used to make meaningful

inferences about data. Unlike DC nodes, which have limited knowledge processing capabilities, Drools is able to chain rules together and a rule that may not have been triggered at the start of processing may be triggered later if another inference is made. For example, a rule that is specifically for small mammals may not be triggered until an inference has been made that the image may contain small mammals based on the time and location of the observation.

Drools is able to dynamically update its knowledge base, adding rules and firing them on observations that have already been loaded, as well as newer ones. This allows DP nodes to adapt to new rules and local knowledge whilst they are deployed. The use of *drl* files use a mixture of Drools and Java syntax to define rules, allowing them to modify, or create, Java objects. For example, a rule could be triggered on the receipt of sensed data and create a DwC object from the received data, process it and perform checks on the result that would trigger different rules based on that result. This is one of the main reasons we chose Drools, as it can work with GSN and DwC Java objects, as well as the ability to run on any architecture that supports Java. In order to add rules to the Drools system, knowledge of the Drools syntax and, ideally, Java is required.

The functionality of Drools is extensive and the engine is very powerful, however, it does require specialist knowledge to use and manipulate rules. Using a custom developed Drools web interface , detailed in Section 4.3.3, we created a simplified interface that uses a custom REST API for Drools, allowing users to create sessions, add rules, load data and fire rules, returning the output to the interface. Users can view, and load, existing *drl* files, shown in Figure 4.8. I

Once a file has been selected, users can view the rules and use the controls on the webpage to perform common operations. The *Load Darwin Core* button loads all DwC archives that are stored in the MySQL database into the current Drools session and the *Fire* button runs all loaded rules on the loaded data. If any of the rules trigger, then the output is presented to the user in the same page, allowing them to act on the results. For example, the location of observations could provide a narrowed down list of potential

```

File: Sample.drl
import com.encima.dwc.DarwinCore
import com.encima.dwc.Occurrence
import com.encima.dwc.Location
import com.encima.dwc.Species
import com.encima.dwc.Identification
import com.encima.utils.DBTools
import com.encima.utils.TimeTools
import java.util.Collection;
import java.util.Date;
rule "Time Frame"
when
    archive:DarwinCore(archive.getOcc().getEventDate() == "27-MAR-2012")
then
    System.out.println("Date is March");
end
rule "Otter River"
when
    $dwc:DarwinCore()
    DarwinCore(DBTools.getLocationDescription($dwc.getOcc().getLocationID()) contains "River")
    eval(TimeTools.checkTimePeriod($dwc.getOcc().getEventTime()) == "Afternoon")
then
    System.out.println("Otter");
    int id = $dwc.getOcc().getEventID();
    Date d = new Date();
    Identification ident = new Identification(id, 1, d, 1);
    $dwc.setId(ident);
    System.out.println($dwc.getId());
end
rule "Plantation Night"
when
    $dwc:DarwinCore()

```

Figure 4.8: Web Interface for Drools Operations

classifications, allowing an administrator to remove votes that do not match the list.

Users can write new rules based on patterns gleaned from observations received during a K-HAS deployment, which can be found using the web interface and querying the database. These patterns can be encoded as rules into a *drl* file and uploaded to have an immediate effect on the active knowledge base in the network.

Currently, these implementations have been developed for our motivational scenario, but much of these tools are general enough to be repurposed in order to apply to a variety of different WSNs. The Drools API can be used for any kind of sensed data and the web interface would require minor changes to be extensible.

4.3.5 Routing Protocol

The routing protocol we have selected for K-HAS is not fixed for every deployment but, from our research, we modified the commonly used Minimum Cost Forwarding

Algorithm (MCFA), outlined in Section 2.2.1, as it allows for changes in the topology of the network and does not require every node to have a global view of the network.

On deployment of all nodes, a DA node sends out a packet containing the number zero, representing the number of hops to the DA node. Nodes in range receive the packet, store it along with the identifier of the originating node and then send it once it has been incremented by one. The next nodes receive that packet and do the same until the edge of the network is reached. If a node has a number stored that is higher than the one it receives then it is replaced and sent on until the nodes at the edge of the network are reached.

When a node wants to send data, it queries neighbouring nodes and sends to the node with the lowest hop count to the root. We modified this to run in accordance with our tiered architecture as we expect the topology to remain the same for much of the deployment.

In MCFA, nodes do not store path information and messages are broadcast to all nodes when they are sensed.

Firstly, our protocol runs in two modes: configuration and running. During the configuration mode, we use MCFA. Whereas MCFA does not store path information, we use the method described above to store both the hop count and the nearest neighbour. This process is carried out until all nodes of the network have a hop count (and neighbour) and a final pass is made by all nodes to find, and store, their neighbour with the lowest count. The main difference is that, unlike MCFA, not all nodes process all packets. If a packet originates from a DA node, then it is only stored by DP nodes but DP node packets are stored by both DC and DP nodes. DC nodes store it and send all data through their nearest DP node and DP nodes store other DP neighbours to delegate processing to, should a situation arise where they have too much data to process and cause a bottleneck. Configuration mode can then be run at a set interval throughout the deployment of the network, or initiated manually. Figure 4.9 shows how DC nodes process routing discovery packets from other nodes.

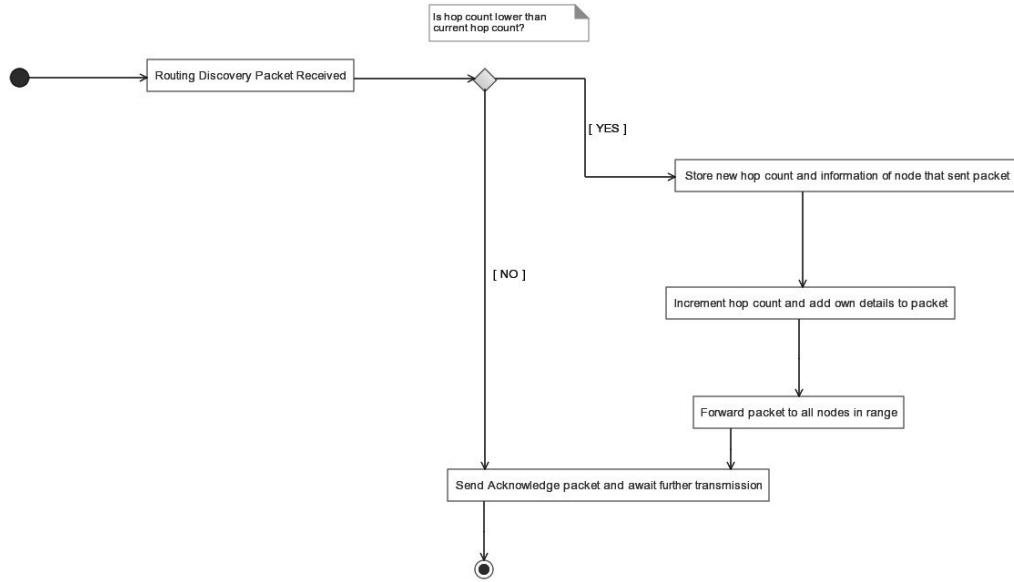


Figure 4.9: Activity Diagram for DC Nodes during the Discovery Phase

Whilst in running mode, nodes do not query for the neighbour with the lowest hop count, or broadcast the sensed data to all nodes in range, they send to the node stored as their nearest neighbour. If that fails, then a query is sent out to find other available nodes in range and then sent to the one with the lowest hop count. If the nearest neighbour node is unavailable for more than three attempts, then it broadcasts a request to run the configuration mode again.

4.4 Walkthrough

In this section we will explain the steps involved in the capture, and processing, of an observation when using the K-HAS architecture. Each tier is responsible for performing different actions upon the observation to ensure it is received by the DA node with an inference as to what it may contain.

Not all of the features described in this section have been implemented within K-HAS and some features have only been tested on a small set of observations. The image

classification down to species level is a concept that we have not implemented but our proposed future work on template matching is a commonly used practice within image processing [67]. The feedback loop, described in Section 4.2.3, has also not been implemented.

4.4.1 Scenario

This walk-through will use our motivating scenario and the type of sensed data will be images of animals in the Malaysian rainforest. In this example, we have a collection of wildlife cameras, with nodes attached to them, deployed in the forest. Projects for the rare clouded leopard and sun bear are currently active at Danau Girang. The clouded leopard is a nocturnal carnivore that uses existing paths and hill trails to travel through the rainforest and the sun bear is the smallest bear in the world and sightings are rare. It is also nocturnal and claw marks can be seen on trees that they have climbed. All of this information has been encoded onto the DP nodes and DC nodes know that images taken at night will be of a higher priority, as well as to prioritise further images at night from DC nodes that are deployed on ridges or existing trails.

In order to explain the K-HAS architecture, we need to show the planned topology for the network in Danau Girang, which is based on the positions of the cameras in 2010. Figure 4.10 shows a section of the proposed topology around the field centre, with the rest of the network spreading out along the river on both sides. The triangle icon shows the location of the DA node (at the field centre), with DP nodes (square icons) placed near the DA node because of the poor range of Wi-Fi (Section 3.3.1). Circle icons represent the DC nodes and they link to the nodes with the fewest hops to a DA node, as long as that route includes a DP node.

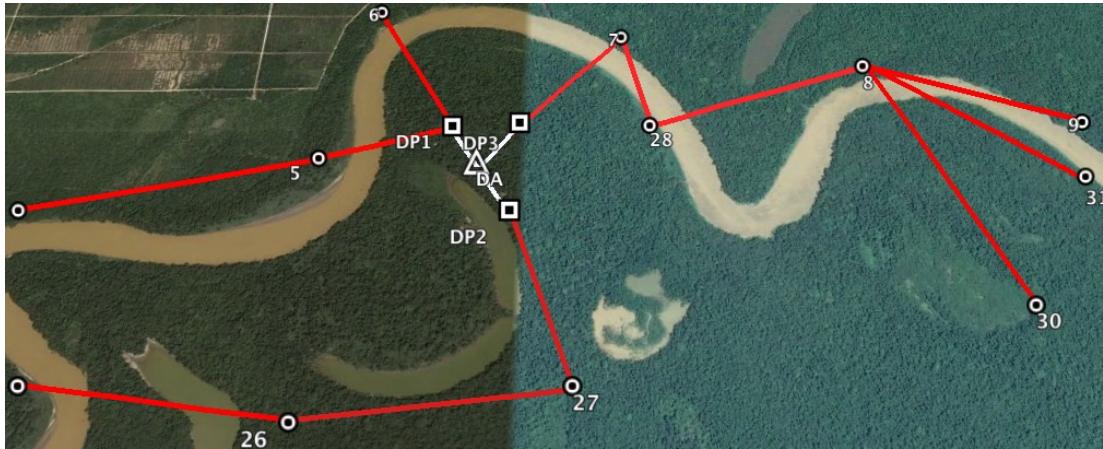


Figure 4.10: Section of Proposed Topology for K-HAS in Danau Girang

4.4.2 Data Collection

A DC node is deployed along a ridge in the rainforest and consists of a wildlife camera with a wireless node attached. At 0200, the infra-red sensor detects movement and the camera triggers a set of 3 images to be captured. The DC node creates the DwC archive for the observation. Terms that describe the observation, such as time, date, species identified and location, is added to the meta.xml file and links to CSV files that contain the data for each term. Any field that can be completed, such as time, location and date, is added to the set.csv file. A separate CSV file is created that holds the filename of each image that was taken. The image is shown in Figure 4.11.

The DC node runs its rules on the metadata of the images and infers that the image may contain a clouded leopard, this is because the image was taken in the early hours of the morning and the camera is deployed on a ridge. DC nodes run simple, non-chaining rules based on the file metadata and details about the node, such as location, but these rules are fixed from the time of deployment and are not updated until the network is redeployed. When bandwidth is restricted, DC nodes use a queue to prioritise both their own observations and those received from other nodes. Observations marked as interesting are moved to the front of the queue and any others are sent afterwards.

The inference is included in the archive and is compressed. The node then sends it



Figure 4.11: Clouded leopard Image Capture

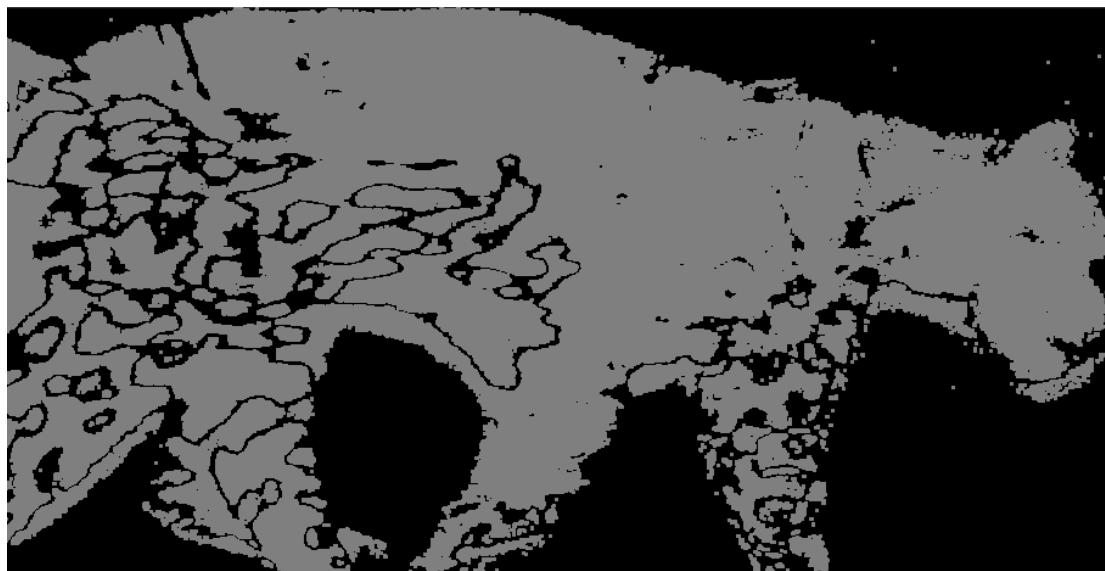


Figure 4.12: Processed Image of clouded leopard

through every DC node between the originating node and the DP node assigned. To achieve the long range communications in the forest, Digimesh is used; the low transfer rate does mean that an archive can take several minutes to send but it allows for a range of up to 1km.

4.4.3 Data Processing

The DP node receives the observation and it is unzipped and processed by the Darwin Core library. If the data has been preprocessed by a DC node and marked as interesting, then the processing is prioritised, otherwise it is added to the queue. In this case, the DC node believes the observation is interesting and may contain a clouded leopard, so it would be processed before other observations that may be queued and not marked as interesting. The images are read from the filenames provided in the CSV and processed using two methods. The EXIF tags in the image are extracted and the images themselves are processed using the Open Computer Vision (OpenCV) library. A unique feature of the DP node is that it uses two radios to allow links to both DA and DC nodes. DC nodes send archives using Digimesh, to achieve long range communication, and DA nodes use Wi-Fi, to provide a faster transfer rate than Digimesh and a more standard method that allows other devices to connect, such as mobile phones or laptops.

In this example, the knowledge base on the DP node is aware that clouded leopards and sun bears are nocturnal, but clouded leopards have previously only been seen when the temperature is between 30 to 35°C and only when the moon is not full. However, data on sun bear is not as complete and the knowledge base only shows that the bear is nocturnal and can be seen at any time of night in any area of the rainforest. The DP node uses EXIF tags from the images and identifies that the moon phase is not full and that the temperature is 32°C, from this it determines that the image could contain either animal and it cannot make a final conclusion.

Image Processing

Our Triton program, described in Section 3.4.1 is run on the set of three images. These images are converted to black and white and combined to build a background model for the complete set. The detected background is then removed and the final image is

then searched for objects, where objects in the foreground will be shown with white pixels. The largest object is then found in the image and extracted to create a template, shown in Figure 4.12.

Processed images of previously sensed images are stored on the DP node and associated with the confirmed classification, confirmed by a human or a node. Although the memory available on a DP node is typically around 32GB, this could easily fill in a matter of months if 3 full HD images were stored for every observation. Storing a single black and white template that contains a portion of the image is much more efficient and can still easily be associated with the classification made. The extracted image is then compared with the existing images, using the knowledge base to prioritise templates for comparison. In this example, nocturnal animals are prioritised and especially nocturnal animals with active projects associated. If the DP node has received an observation from the same DC node recently, then it will check for a classification on that and check for a match there first.

As explained in Section 3.4.1, species classification has not been implemented in K-HAS and templates are currently only associated with human classifications but this walkthrough describes how the process would be carried out if classification were to be implemented. These findings are written into the set.csv file of the DwC archive (Listing 4.4), using the identifier of the DP node as the ‘person’ that identified the image and the scientific name for the clouded leopard as the species identified in the image. The archive is then zipped and sent on to the DA node, sending observations of interest first and delaying observations that have been found to contain nothing of interest.

This observation is the first trigger from the DC node in the past few hours, so there are no recent classifications. However, processing of the metadata showed that the image was taken at night, so the DP node would use its knowledge base to match the images to templates of nocturnal animals first. Triton could then find a match to an existing template of a clouded leopard and completes its classification.

Classification

The metadata processing of the image shows that it could be any nocturnal animal that is known to come out when the moon is not full and the temperature is 32°C. This is not a complete classification but the image processing has found a match.

4.4.4 Data Aggregation

Upon receiving a DwC archive, it is unarchived and processed by Darwin Core libraries, called by the middleware running on the node. The resulting archive is then inserted into a database and the files themselves are stored in a directory that maps to the DC node that captured the original observation. At the field centre, three users of the system have subscribed to updates for observations of clouded leopards.

As the archive is processed by the library, the species is extracted and this triggers a rule to notify the subscribers. The rule then queries the database and finds their preferred method of communication. In this case, one is a lecturer and wants to be emailed while the two remaining are students and want to be notified via Twitter. An email is generated that contains the time, location and content of the observation, with the images attached, and sent on to the lecturer. The students are sent a short tweet that tells them a clouded leopard has been spotted and a link to the middle image in the sequence is provided; the middle image is used because local knowledge has shown that it is the most likely to contain the full subject in the image. In order to maintain privacy, a ‘direct message’ can be sent on Twitter so this message is not public.

The middleware on the node supports a web interface to allow users to perform administrative functions on the network, such as deploy a new node, on top of this there is a custom made website that shows all observations for every DC node. This allows users to log on and classify the images. In this case, the lecturer receives the email notification, reviews the attached images and clicks on the link to access the website to inform the DP node that the classification was correct. Due to the administrator

position the lecturer has on the system, he is able to stop any users voting on the image and to simply confirm the classification.

When a user classifies the observation, they see that a clouded leopard has been spotted in the same area on the same day for the past 5 weeks and they create a rule (in Drools syntax) to automatically classify images from this camera that have a similar time (within an hour) and have an object extracted from them by the image processing. The user can then upload the rules through the same web interface and it will instantly become active on the system. In the current implementation, rules can only be added by humans and the Drools API we have implemented then updates the rule base. The web interface allows users to study patterns in existing sensed data and perform queries on the database, from this they can identify rules and upload them.

If there was no classification, then users would be able to vote on the contents and use the classification with the highest vote, or the classification made by an administrator. Once a classification has been made, it is stored in the database and written to the archive. This triggers the DA node to send that classification on to the DP node that sent the original archive. In this case, the DA node informed the DP node that it has been confirmed as a clouded leopard and the DP node then stores the extracted image in the directory of clouded leopard templates, to assist with future classifications. This updated template causes the rule base of the DP node to be updated so similar data processed in the future would be correctly identified. The longer the network is deployed, the more knowledge DP nodes could gain and the more accurate their classifications can be. For example, a change in season could cause a new, previously unknown animal to migrate to the rainforest. With human assistance, the animal can be identified and determined whether it is of interest. This knowledge can then be stored and sent on to DP nodes to prioritise and classify correctly the next time that it is captured. The feedback loop protocol is currently only a design and has been minimally implemented as a proof of concept, we believe it can improve the accuracy of classification in a WSN whilst nodes are still deployed but it would require testing

in the future.

4.5 Conclusion

In this chapter we have explained the architecture we have developed to allow knowledge to be encoded and utilised within a wireless sensor network. Using tiers of nodes, with varying levels of knowledge-processing capabilities, we can process observations within the network and deliver data that has, where possible, already been classified. Working as more of a subscribe-push method, users do not have to check a DA node for new data, instead it is sent to them if it has been found to be part of a project they are subscribed to. If not, then the data is accessible to all users of the network through a web interface.

Using GSN as the application middleware allows sensors to be added, modified and maintained by those without technical knowledge and ensure future interoperability. However, it is not a standardised approach and using an architecture that implements the standard OGC SWE [96] would allow for interoperability through the use of standards not just for sensor networks but the Web as well.

One of the key features of K-HAS is that it is not a static deployment. The knowledge that the network holds at the time of deployment will rarely be the same as the knowledge held after a few months. Humans enrich the existing knowledge base and the nodes are able to make inferences about the data they are sensing, improving their classifications the longer they are deployed. When developing this architecture, GSN was the most robust architecture of those that we researched and tested, the support for many databases, administrative interface, native support for many widely used sensors meant that it was a better choice than a middleware that adopted the OGC SWE standards, especially as we were not interacting with SWE systems in our motivating scenario. However, while GSN is stable and mature, its large codebase does mean that there are dated features, such as using SOAP instead of REST and an unintuitive web

interface that does not utilise web sockets. A middleware with similar automation on receipt of sensor data could be used instead of GSN that did follow the standards set out by the OGC. We believe that this should require few changes to the core K-HAS architecture as it currently stands.

In Chapter 6, we explain how we implemented a variation of K-HAS in our motivating scenario and Chapter 7 shows our evaluation of the K-HAS architecture, but Chapter 5 will first outline the development of an ontology to support the architecture described here.

Chapter 5

An Ontology for Knowledge-Based Wireless Sensor Networks

In this chapter, we explain the ontology we have proposed to formally define the components within K-HAS, the structure of the sensed data, the users involved with the network, as well as the format of the sensed data that is passed through the network.

This ontology is for those wanting to deploy a WSN that uses knowledge to classify the ecological observations recorded, or even to select parts of the ontology that meet their requirements and implement those. Terms have been used to map to widely-used ontologies in the domains of ecological observation, sensors and people. However, our proposed architecture (Chapter 4) defines new terms that do not exist in current ontologies, thus we extend the alignment of these existing ontologies to create an ontology that encompasses the ontologies as well as defining new terms to match our architecture.

5.1 Background

K-HAS was developed in order to provide a generic architecture for wireless sensor networks to utilise the local knowledge contained within their environment to process sensed data and, therefore, make more efficient use of the network bandwidth by prioritising sensed data that is deemed to be more valuable. We have defined local

knowledge as knowledge of an area that has been gained through experience, or experimentation.

Before we were able to implement K-HAS, we needed a high-level model that defined relationships between each tier of the architecture, as well as the data standard used to transport sensed data from DC nodes to an endpoint. Developing an ontology for K-HAS means that we can do this, as well as provide a computer-readable model for all of the classes and components used by K-HAS.

Making it computer readable has allowed us to reuse classes in the development of software for each tier. For example, we were able to develop a common Darwin Core java library that is used in this ontology and in our GSN middleware to unzip and process received archives.

During the development of the K-HAS ontology, we researched existing ontologies that were commonly used in the domains that K-HAS covered. These included scientific observations and sensor hardware.

Looking into these existing ontologies, we found that there had been many surveys on representing sensors in the semantic web: [30], [59], outlines the existing work. These surveys clearly highlighted that these ontologies had been split into two branches; observation-centric and sensor-centric.

Observation-centric ontologies, such as OBOE [72] and O&M [96], focus on the data that is sensed, and its content; whereas sensor-centric ontologies, such as SensorML[96], detail the components that make up a sensor and the operations they perform to turn sensed data into an output.

K-HAS has been developed to be used with any sensors and is not specific to wildlife cameras, therefore we also looked into sensor-based ontologies that concentrate on the hardware and the individual capabilities of each device within the network.

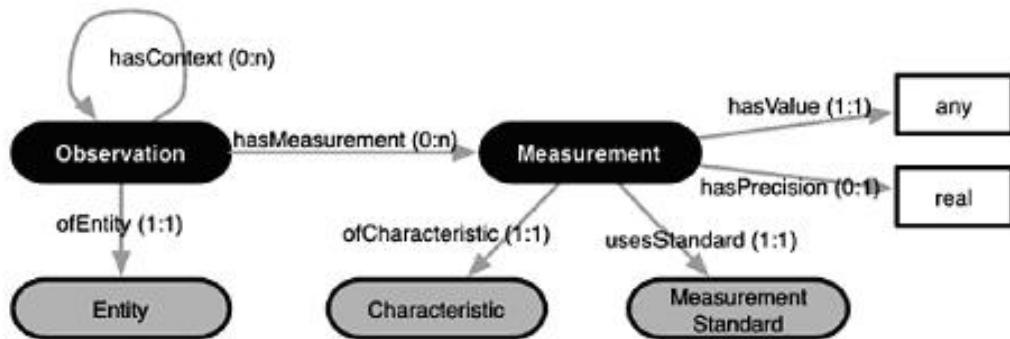


Figure 5.1: The Core of an OBOE Observation [71]

5.1.1 Observation-Centric Ontologies

OBOE

The Extensible Observation Ontology (known in reverse as OBOE) is a popular suite of ontologies used to represent scientific observations [72]. Initially starting as base ontology for ecological observations, it has now grown into a suite of extensions that make it suited for chemistry, bioinformatics, anatomy and others. OBOE is represented by OWL-DL [79] and allows the characteristics of a generic scientific observation to be linked to domain specific characteristics.

OBOE focuses on the concept of an *observation*, which is made up of an entity, a measurement and a characteristic [71]. An example of an observation could be a researcher observing an animal (the entity) and recording the gender (the measurement) as male (the characteristic). A single observation could then consist of multiple observations within it, such as gender, location, species and the number of species observed.

Figure 5.1 shows the basic structure of a core OBOE observation, outlining the five key classes that are linked by seven properties. While this is the core structure of OBOE, domain specific extensions have been implemented by utilising ‘extension points’ that are part of OBOE’s core. OBOE follows the O&M Standard (below) very closely, providing extensions to the core classes that allows more information about each to be

encoded, adding context and enhancing the value of an observation.

The primary benefits of OBOE are that it is generic enough to cover almost all types of scientific observation and domain extensions allow for more specific details to be stored.

O&M

The Open Geospatial Consortium (OGC) Observations and Measurements Standard aims to provide a framework suitable for recording any observation made by a sensor, regardless of the domain [96].

A key feature of O&M is that *observation* and *measurement* are not just classes. They also denote an action. An observation is an action that causes a result, yielding a value and a measurement is a set of operations that provide some result(s).

O&M provides a conceptual model, as well as XML encoding for observations and measurements. The listing 5.1 shows an observation of a vehicle in a given time and place. Similar to the encoding of a measurement with the standard and protocol used in OBOE, O&M provides support for the recording of the procedure used to gain the measurement for the observation.

```

1 <?xml version="1.0" encoding="windows-1250"?>
2 <om:GeometryObservation gml:id="geom1610"
3 xmlns:om="http://www.opengis.net/om/1.0"
4 xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
5 xmlns:xlink="http://www.w3.org/1999/xlink"
6 xmlns:gml="http://www.opengis.net/gml"
7 xsi:schemaLocation=".Specialization_override.xsd">
8   <om:samplingTime>
9     <gml:TimeInstant>
10    <gml:timePosition>2009-09-16T17:22:25.00</gml:timePosition>
11  </gml:TimeInstant>
12 </om:samplingTime>
13 <om:procedure xlink:href="urn:ogc:object:procedure:ifgi:GPS"/>
14 <om:observedProperty xlink:href="urn:ogc:def:phenomenon:OGC:Shape"/>
15 <om:featureOfInterest xlink:href="urn:ogc:object:feature:vehicle"/>
16 <om:result>
17   <gml:Point srsName="urn:ogc:crs:epsg:4326">
18     <gml:pos>40.7833 -73.9667</gml:pos>
19   </gml:Point>
20 </om:result>
21 </om:GeometryObservation>
```

Listing 5.1: An Observation of a Vehicle encoded in O&M

The listing shows that an observation centres around a *feature of interest* that can be a physical object and the measurement is the detection of the vehicle at the recorded location. Figure 5.2 shows a basic diagram of the ontology. While the event of a feature is linked, it is clear that the main focus is on the observation and the measurement associated with it.

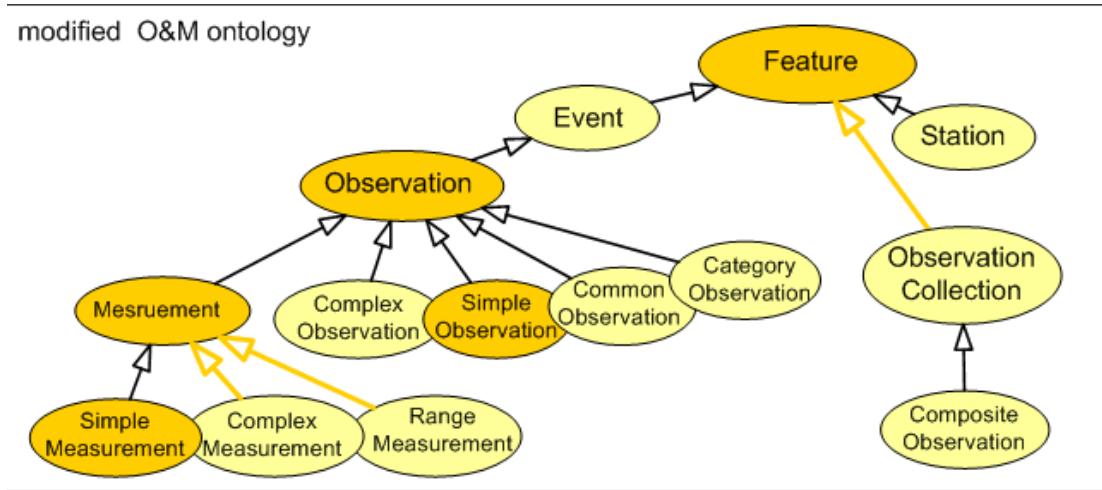


Figure 5.2: The OGC O&M Ontology [88]

The structure of an observation within O&M is focussed on the action and the result, this makes it suited for sensor networks across many domains that perform a wide range of observations. Because the users of the O&M standard are spread across many domains, each with their own terms and definitions, the creation of an ontology for the standard aimed to remain as generic as other observation-centric ontologies.

Darwin Core SW

In order to represent DwC occurrences, covered in Section 4.3.1, in an ontological format, work has been done to represent Darwin Core terms, as an ontology, in OWL. Darwin-Semantic Web (Darwin-SW) [19] is the project that aims to do this and many of the core terms associated with an occurrence have already been formalised and Figure 5.3 shows the entity relationship model of these terms.

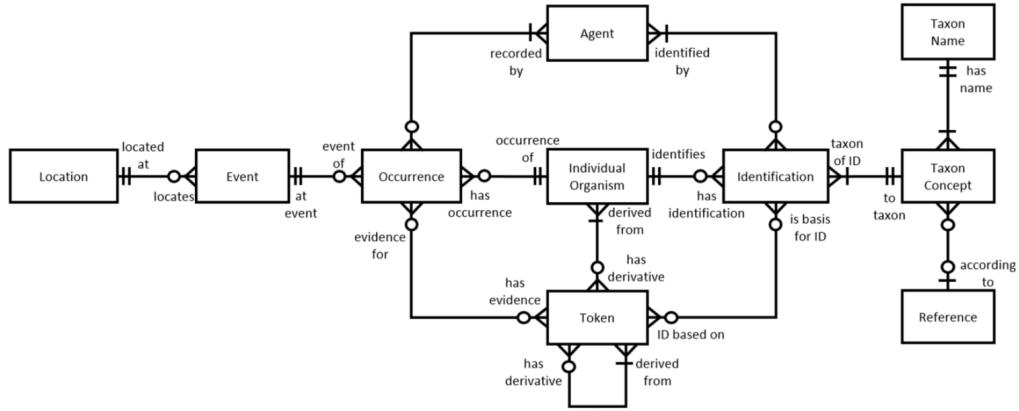


Figure 5.3: Darwin-SW Entity-Relationship Model [19]

While Darwin-SW does not represent all of the classes within the DwC namespace, it contains the core classes required to record an occurrence, to a more-detailed level than OBOE. The downside of this is that the specificity of the terms limits DwC to ecological observations, making it far less generic than OBOE and other alternatives. There are positive aspects that, for those that need to record ecological observations, DwC allows for a greater level of detail and can be mapped to OBOE with ease.

Listing 5.2 shows a fragment of a DwC occurrence of a living specimen which, in this case, is a tree [18]. Darwin-SW, and Darwin Core, terms are used to describe the identification of the tree, such as the by whom it was identified, the date and the taxon. The full source can be seen in Appendix B and includes links to existing related occurrences, other images of the individual and relationships to other resources.

```

1 <dsw:hasIdentification>
2   <rdf:Description rdf:about="http://bioimages.vanderbilt.edu/vanderbilt
3     /12-126#2002-04-10baskauf">
4     <dcterms:description>Determination of Ginkgo biloba L. sensu Flora of
      North America (1993) for the individual http://bioimages.
      vanderbilt.edu/vanderbilt/12-126</dcterms:description>
6     <rdf:type rdf:resource="http://rs.tdwg.org/dwc/terms/Identification" />
7
8   <!-- In lieu of stable external identifiers for taxon concepts, Im
      defining some onsite -->
9   <dsw:toTaxon rdf:resource="http://bioimages.vanderbilt.edu/
10    taxonConcepts#183269-fna1993" />
11
12 <dwc:identifiedBy>Steven J. Baskauf</dwc:identifiedBy>
13 <dsw:idBy rdf:resource="http://bioimages.vanderbilt.edu/contact/
14    baskauf" />
15 <dwc:dateIdentified>2002-04-10</dwc:dateIdentified>
16
17 <!-- Relationship of the identification to other resources -->
18 <dsw:idBasedOn rdf:resource="http://bioimages.vanderbilt.edu/baskauf
19    /10554"/>
20
21 </rdf:Description>
22 </dsw:hasIdentification>

```

Listing 5.2: Darwin-SW Representation of an Identification

5.1.2 Sensor-Centric Ontologies

Sensor-centric ontologies are more focussed on the structure of the sensor, the network and the sensing processes involved.

SensorML

The Sensor Markup Language (SensorML) Standard has been developed by the OGC, and complements the O&M Standard, to enable the discovery and tasking of internet-connected sensors [96].

SensorML provides an XML schema for describing a sensor, its capabilities and the processes available. At the core, SensorML comprises of:

1. Component - A physical process that transforms information from one form to another.
2. System - Model of a group of components.
3. Process Model - Atomic processing block used within a Process Chain.

4. Process Chain - Composite block of Process Models.
5. Process Method - Definition of the behaviour of a Process Model.
6. Detector - Atomic part of a Measurement System.
7. System - Array of components, relates a Process Chain to the real world.
8. Measurement System - Specific type of System, mainly consisting of sampling devices and detectors.
9. Sensor - Specific type of System that represents a complete Sensor.

These definitions outline the core concepts of SensorML, a *system* that performs one (or more) process(es) and is comprised of a group of *components* [98]. A SensorML document allow for a general, formal specification of a *sensor* and its capabilities. The document describes a *component*, outlining what data it reads in and the output once it has been processed. Several of these *components* can be used to create a *system* and the primary goal of SensorML is to describe the process of how an observation came to be, focussing on the technical features of the node.

5.1.3 Combined Sensor and Observation Ontologies

In this section, we review ontologies that represent the hardware of the sensors within a network, as well as the observations they record. This allows them to model the sensing capabilities of each sensor, the measurements they capture and the data they record.

SSN

The Semantic Sensor Network (SSN) Ontology is the most fitting ontology for our requirements, as it covers systems processes and observations. Developed by the W3C

after an extensive review of existing ontologies [31], the SSN ontology is designed to allow focus on a variety of perspectives, such as the sensor within the network or the data that has been observed.

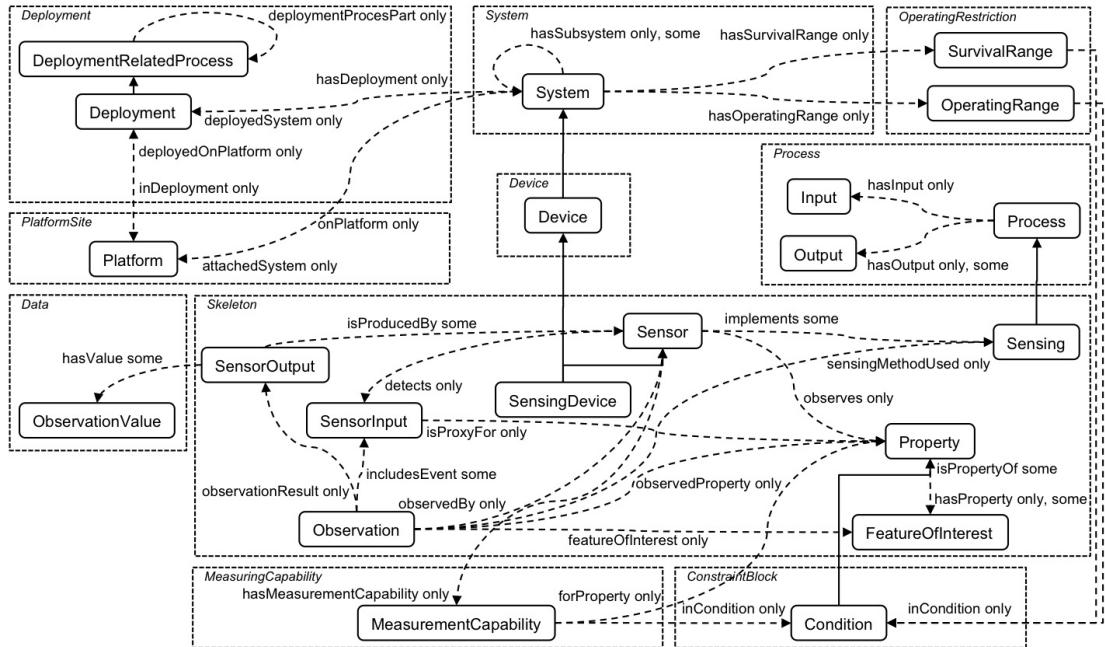


Figure 5.4: SSN Ontology

Figure 5.4 shows the model for the SSN ontology, displaying the relationships that connect each class, as well as their associated properties. It also highlights the modular approach that has been taken, separating the system from the process and the observation.

The observation pattern of SSN is centred around ‘Stimuli-Sensor-Observation’, which can be simply described by an event causing a sensor to trigger and create an encompassing observation to store details about the event, as well as the device that recorded the event. While SSN is focussed on sensors, capturing the measurement capabilities of each sensing device that makes up a system and specifics on its lifespan, the development of the ontology arose from the review of sensor-centric ontologies as well, recording data about observations in a structure very similar to O&M.

SUMO Extension

The Suggested Upper Merged Ontology (SUMO) is a general purpose ontology [83] that provides general-purpose terms and is intended to be extended for domain specific ontologies. In [37], SUMO has been extended to link sensor hardware and sensor data ontologies in order to assist with searching and evaluating distributed and heterogeneous sensor networks. The work combines the Sensor Hierarchy Ontology (SHO), the Sensor Data Ontology (SDO) and the ability to ‘plug in’ extension ontologies.

The SHO describes the hardware of a sensor node, as well as its accuracy, transmission medium and data processing capabilities. The SDO, however, describes the sensing properties of a device beyond the hardware and the context of the sensor can be enriched with information about spatial and/or temporal observations. Similar to GSN (Section 2.3.2), SUMO uses the notion of virtual sensors where a group of sensors can be described together to provide abstract measurements. In [37], the example of a humidity sensor, temperature sensor and wind speed sensor being collectively described as *weather sensors*. The SHO ontology is another extension of SUMO, from [37], that represents the hardware of a WSN, including the node itself, data transmission units, data processing units and individual sensors. The data model describes a sensor with metadata describing features such as measurement type, transmission range and physical properties.

5.1.4 Findings

There are many ontologies that are suited for sensor-based scientific observations and ontologies that allow for the description of sensor hardware cover the hardware capabilities of sensors as a whole and sensor systems comprised of multiple sensors in great detail. SSN and SUMO were the closest that we found for this, but SSN was still very hardware focused and went into greater detail of the hardware capabilities of each measuring devices than we needed.

While the SHO and SDO extensions of SUMO are classed as both an observation-centric and sensor-centric ontologies, they are separate ontologies that model both sensor hardware and sensed data and, while the development of these ontologies did follow a similar approach to the one described in this chapter (combining ontologies), the core concepts of SHO and SDO are more focussed on retrieval of sensed data through queries constructed using both hardware details and properties of the sensed data. However, an important feature is the Extension Plugin Ontologies (EPO), support for domain specific plugins to integrate with SUMO and integrate with the SHO and SDO and this would also be a useful feature for K-HAS.

5.2 Method

We found that SSN, SensorML and Darwin Core satisfied many the hardware and much of the sensed data subsections of K-HAS completely. We have used these existing ontologies to develop an aligning ontology that connects ontologies across multiple domains to support our proposal of K-HAS, extending the concepts in existing ontologies with K-HAS specific concepts.

The SSN ontology is a modular ontology created by combining concepts from existing, commonly used ontologies and allows for domain specific concepts to be imported. Some of the main uses cases for the SSN ontology are provenance and data discovery, which are also key within K-HAS. However, the tiered structure of K-HAS did not map directly to SSN and it proved difficult to represent the flow of knowledge through a network, as humans can also perform similar operations on observations and the observations are enriched as they pass through the network. SSN describes the measurement capability, operating range and hardware specifics of a sensor but, while SSN does also model the process that a sensor may undertake to create an output, K-HAS sees sensors as nodes with computational power; not only able to sense their environment but perform classifications utilising on their knowledge bases. We found

it difficult to model the processing power of each sensor without modifying both the *Skeleton* and *Process* modules of the ontology. Also, while SSN does contain *ObservationValue* and *FeatureOfInterest*, shown in Figure 5.4, the Darwin Core standard that we adopted was more detailed with scientific observations. Using the *Data* module of SSN, we could have extended it with the Darwin Core ontology but K-HAS overlaps *ObservationValue* and *FeatureOfInterest* to become more focussed on the classification than the observation itself.

However, while we did not use the ontology directly, many of the concepts can be mapped directly (shown later in Table 5.4) and it would be possible to use import modules from SSN and extend them with the concepts specific to K-HAS, SSN does not describe domain concepts (such as time or location) as this is expected to be handled by imports from more specific ontologies. K-HAS allows humans to act as sensors, which would require slight modifications to SSN, as well as extending with the ability to provide more detailed classifications for observations.

Before we began development, we researched existing methods that have been used to develop current ontologies. From this, we found three well-documented methodologies; the methodology used to develop the Toronto Virtual Enterprise (TOVE) ontology, the Methontology and the methodology used to develop the Enterprise Model ontology.

The methodology by Uschold and King [112], used to develop the Enterprise ontology, is a four step method that was most suited to the processes we expected to follow. The steps are not covered in great detail in the original paper but, research since provides greater detail for each step.

5.2.1 Identify Purpose

The aim of this step is to identify why the ontology is being built and what requirements it is supposed to fulfil. This includes considerations, such as the audience, the intended purpose and the specificity of the ontology.

Other methodologies have used more structured methods to informally identify the requirements of the ontology. Gruninger and Fox [46] propose competency questions; these are questions that are used to identify the problems that the developed ontology is developed to solve. These questions can act as a benchmark, showing that an ontology is sufficient if it can solve the questions raised here.

5.2.2 Build the Ontology

Building the ontology can be broken down further into 3 smaller steps: capture, code and integrate existing ontologies.

Capture

The capture stage involves defining the concepts and terms that the ontology will model, as well as how they map to the real world. This can be done in one of three ways:

1. Top Down - Starting with the core concepts, create the more specialised classes until you have identified all subclasses.
2. Bottom Up - This process begins with the more specialised classes and grouping them into the more general classes towards the top of the hierarchy.
3. Middle Out - A combination of the two, this process involves specialising, and generalising, the classes identified in the middle layer of the hierarchy.

As for the capture of the knowledge used to identify the classes needed, this is an area that has been documented, but primarily provides recommendations. [39] suggests interviews with domain experts, iteratively brainstorming with a group actively involved in the development of the ontology. This stage has often been referred to as the knowledge-acquisition stage.

Code

This step involves coding the terms identified in the previous step into a formal language, such as the Web Ontology Language (OWL) [79].

Integrate Existing Ontologies

This step can be carried out at the same time as the step above, so that the ontology is developed with existing ontologies in mind. This also allows overlap to be identified, and incorporated, early. Being aware of commonly-used ontologies within the domain, before development begins, is a more logical approach and does avoid the need to create new terms for existing concepts.

Existing ontologies can be integrated by importing them and linking the existing terms to the terms identified in the ontology that is being developed. However, there is also the method of developing the ontology completely, so that it is consistent without the need to rely on existing ontologies and creating ‘sameAs’ relationships between the terms identified and only the required terms in the existing ontologies.

In [60], it is recommended that, when importing external ontologies, the whole ontology is not used. Rather, one should aim to extract only the required fragments from the external ontology that are relevant to the concepts in the developed ontology.

5.2.3 Evaluation

For the evaluation, Uschold and King adopt the definition of [40]: ‘to make a technical judgement of the ontologies, their associated software environment, and documentation with respect to a frame of reference... The frame of reference may be requirements specifications, competency questions, and/or the real world’.

There are some well documented methods for evaluating ontologies in the literature,

[111] proposes using the competency questions, used in the first step, to ensure that they can all be fully answered by the finished ontology.

5.2.4 Documentation

Although this step is listed at the end of the development cycle, it seems more fitting to document all major aspects of the ontology as it is being developed. Documentation of the ontology should include: any assumptions, all concepts introduced, all ontologies that have been incorporated (by whatever means) and any primitives used for the definitions of concepts.

5.3 Results

This section details our results when using the Uschold and King methodology, accompanied by more recent research for particular steps, to develop the K-HAS ontology.

5.3.1 Identify Purpose

The need to create the K-HAS ontology was partly due to the reasons outlined by Gruninger and Fox [46], we had identified a problem as we were developing a sensing architecture that utilised local knowledge: how could we formally represent the flow of knowledge and sensed data throughout a wireless sensor network?

To determine the scope that K-HAS should cover, we identified a set of competency questions that represent what we expect K-HAS to cover within the domain. In order to present this, we used an approach similar to that of [27], which can be seen in Table 5.1.

These questions allow us to identify the core concepts that the ontology needs to represent, as well as serving as a tool to evaluate the completed ontology.

Competency Questions
Find all Occurrences of an Individual
Find all Occurrences of an Individual at a Location
Find all Occurrences within a specified Date and Time range
Find all Sensors that have recorded an Occurrence of an Individual
Find all Locations of an Individual
Find the storage location of a Project
Find all Projects containing an Individual
Find all People involved in a Project
Find all the Evidence that supports an Occurrence
Find all the Types of evidence that supports an Occurrence

Table 5.1: Competency Questions

5.3.2 Build the Ontology - Capture

This part of the development cycle is the identification of the concepts and their implementation. The first step is capture the knowledge that will be used to identify the core concepts within the ontology.

To capture the knowledge for the ontology, we used a combinatorial approach of those outlined in the literature. We interviewed domain experts, as well as overseeing a basic implementation of a system, and we iteratively brainstormed the concepts throughout the development cycle.

Over the course of eighteen months, which involved 2 field visits and several brainstorming meetings, we identified the core concepts that the ontology would need to contain, as well as the properties that would link them. From this, it would seem that we followed the *Top-Down* approach, but the first visit with domain experts also allowed us to identify some of the more specialist classes early on in the development cycle. Thus, it seems that in practice we followed a more *Middle-Out* approach.

Table 5.2 outlines the core concepts we have identified for K-HAS to be complete, as well as the definitions we have used for the architecture. When integrating existing ontologies, the definitions of similar terms would need to match with our definitions or

we would not deem them the same as the K-HAS concepts.

Concept	Subclass Of	Definition
Occurrence Identification	-	A text-based recording of the content of an Occurrence
Evidence	-	Media to support the Identification, such as a photo or recording.
Location	-	The location of the Occurrence.
Date	-	The date of the Occurrence.
Time	-	The time of the Occurrence.
Project	-	Project(s) that can contain many Occurrences
Node	-	A device with sensing capabilities.
Data Collection (DC) Node	Node	Node with limited knowledge-processing capabilities charged with sensing a feature (or features) of its environment.
Data Processing (DP) Node	Node	Node with knowledge-processing capabilities charged with serving a subset of DC Nodes and processing their sensed data.
Data Aggregation (DA) Node	Node	Node with knowledge-processing capabilities that stores all knowledge and sensed data for the whole network.
Person Administrator	Person	Person in charge of a project (or projects).
Worker	Person	Person involved with a project.

Table 5.2: K-HAS Concepts

Mapping these concepts, a diagram of the base K-HAS ontology is shown in Figure 5.5. The next step is to create the ontology and map the concepts identified to existing

ontologies within the same domain space(s).

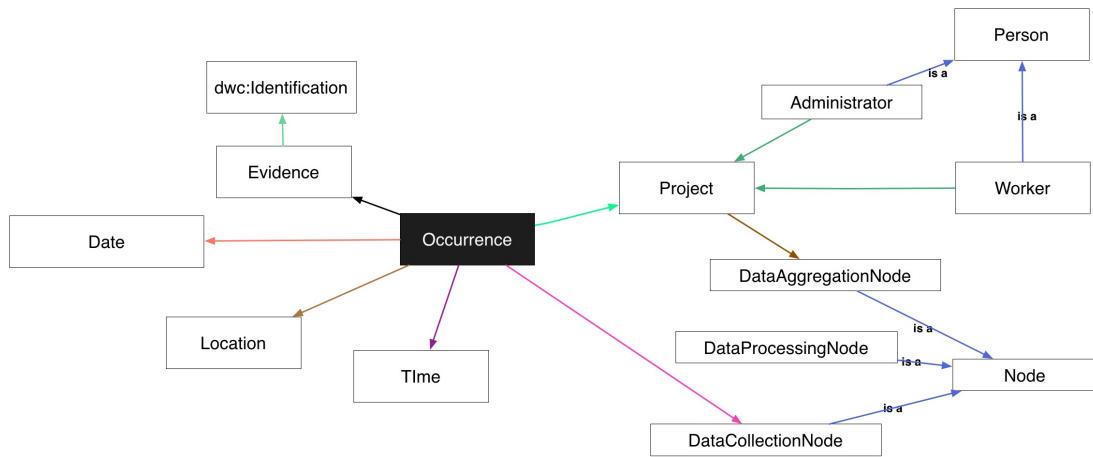


Figure 5.5: K-HAS Base Ontology

5.3.3 Build the Ontology - Code

Once the concepts had been identified (and agreed upon), the ontology must be coded. We used Protege 4.2.2 [84] and implemented K-HAS in the Web Ontology Language (OWL), creating a core file that could be expanded; should we need to import existing ontologies.

5.3.4 Build the Ontology - Integrate Existing Ontologies

Our proposal for K-HAS is focused on scientific observations. Because of this, the focal point of our existing ontology research is on ontologies that are centred around scientific observations. This allows us to create a more generic ontology that is still able to capture all of the semantic details associated with a wildlife observation.

As explained in Section 5.1, our research of existing ontologies covered two categories: observation-centric and sensor-centric. We identified ontologies, within multiple domains, that satisfied many of our requirements for K-HAS, but not all. Using the core

concepts outlined in Section 5.3.2, we created a minimal ontology and used the results of our research to map K-HAS concepts to those that had already been identified.

During this process we found that we had determined concepts that mapped to existing ontologies, but may not share identical structures, or even the same name. For example, the OBOE ontology describes the concept of an occurrence, which is very similar to our occurrence and the occurrence in the Darwin Core-SW ontology. When we found these mappings, we used *sameAs* relationship to form a link that allows data stored according to these existing ontologies to map directly to K-HAS. However, the structure of an OBOE observation is, as outlined in Section 5.1.1, more generic for all types of scientific observation and depicts the observation of an entity, containing a measurement of a particular characteristic. Whereas the structure of a Darwin Core observation is more limited to scientific observations of taxa which allows it to have more predefined terms, such as location, species name and the evidence for the recorded individual. Because of this, it was difficult to create a structure that encapsulated both OBOE and Darwin Core due to the generality of an OBOE observation compared to the more specific structure of DwC.

Research showed that Darwin Core maps to K-HAS' requirements more completely, as well as the fact that there has been work to represent Darwin Core Observations in OBOE [103], K-HAS occurrences map directly to Darwin Core and elements that are similar to OBOE have been linked by the *sameAs* relationship. This does mean that full OBOE observations cannot map directly, but K-HAS, and Darwin Core, occurrences can be converted, if necessary. However, using the terms in our ontology we can create a subset of an observation that does not include all of the terms defined in OBOE but can still map to all three ontologies.

For the sensor hardware of K-HAS, we found that the SensorML ontology maps directly to our concepts and we also realised that we did not need to recreate concepts that may already exist in popular ontologies outside of the domain spaces we researched. For example, the Friend of a Friend (FOAF) ontology [23] is an ontology designed to

create machine-readable pages that describe people, so it is more logical that K-HAS reuses existing terms from popular ontologies to allow pre-existing data to be mapped with ease.

5.3.5 Extend the Ontology

Whilst researching widely-used ontologies, we became aware that some classes identified for K-HAS were not satisfied by what is currently available, these are prefixed with khas and shown in Figure 5.6. The final step for creating the ontology was to add these concepts to the linked ontologies, we call these *extension concepts*, concepts unique to K-HAS which do not exist in any other ontology. These terms were linked to the unique layered architecture of K-HAS and we defined them within a new K-HAS namespace, changing our ontology from an alignment of existing ontologies to an extension of these.

The final ontology is shown in Figure 5.6 (with the complete code in Appendix C) and shows how each concept maps to existing ontologies. The figure shows that there were only five extension concepts unique to K-HAS that can be subclassed from the *person* concept in the FOAF ontology and *node* in the SensorML ontology.

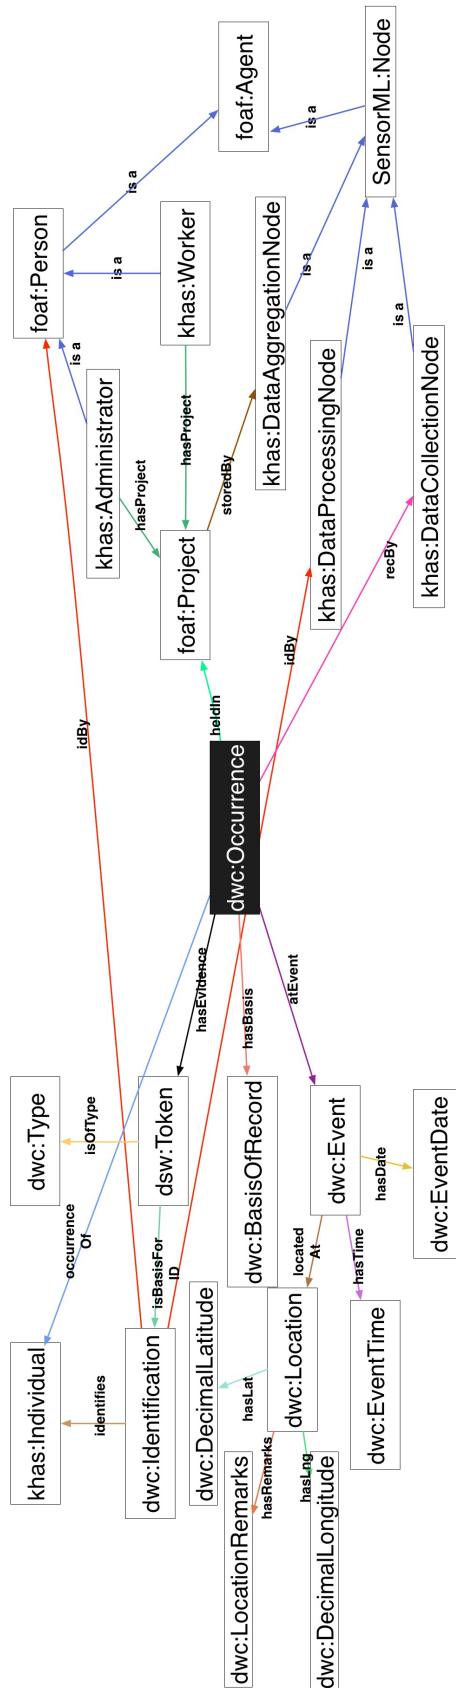


Figure 5.6: K-HAS Ontology

5.4 Evaluation

Uschold and King's ontology development method does not explain how to evaluate the created ontology and much of the literature describes a number of methods that can be employed, [22] outlines a number of different evaluation techniques and separates them into categories. These categories are listed below:

1. Comparing the ontology to a ‘golden standard’.
2. Using the ontology in an application and evaluating the results.
3. Comparing the ontology with a collection of documents from the domain to be covered.
4. Manual evaluation by humans to test if the ontology meets a set of predefined criteria.

The last category mentioned is similar to the method in Section 5.2.3 that uses the competency questions to determine the effectiveness of the ontology. Because there is no agreed method, we have evaluated K-HAS by testing it in the application it was developed in (Protege), mapping it to existing documents within the intended domain and ensuring that it can satisfy all of the competency questions we outlined in Table 5.1.

5.4.1 Using the Ontology in an Application

We developed the K-HAS ontology in Protege, a Java based ontology editor, that also provides the functionality to reason over the ontology and check for inconsistencies. Using the Hermit reasoner that is built into Protege, the ontology was determined to be logically consistent.

To confirm the validation provided by Protege, we also used a web-based ontology validation tool, called the OntOlogy Pitfall Scanner (OOPS), that scans an ontology for common errors, such as defining incorrect inverse relationships or using recursive definitions, that can occur during the development phase [87]. The results of this tool showed that no pitfalls had been detected.

We have a customised build of K-HAS running that receives data from a variety of sources and stores it in a MySQL database that allows users to classify through a web site. The schema of the database maps to our ontology, scientific observations are received, unzipped and stored in the relevant fields in the database. Using real sensed data, we have successfully stored over two hundred occurrences and mapped existing DwC occurrence to K-HAS.

5.4.2 Competency Questions

In Table 5.1 we have identified some basic competency questions that we expected K-HAS to be able to satisfy with the concepts we had identified, Table 5.3 shows how the original questions are satisfied by K-HAS. We have mapped the terms we originally used to the terms used in K-HAS (and the linked ontologies) and shown the concepts involved with each question, as well as the relationships linking them.

Competency Questions	Concepts	Relationships
Find all Occurrences of an Individual	Occurrence; Individual	Occurrence occurrenceOf Individual
Find all Occurrences of an Individual at a Location	Occurrence, Token, Individual, Location, Identification	Occurrence hasEvidence Token
Find all Occurrence within a specified Date and Time range	Occurrence; Event; Date; Time	Token isBasisForID Identification Identification identifies Individual Occurrence atEvent Event
Find all Data-Collection Nodes that have recorded an Occurrence of an Individual	DCNode; Individual; Occurrence	Event hasTime Time Event hasDate Date Occurrence recordedBy DCNode
Find all Locations of an Individual	Occurrence; Individual; Event; Location	Occurrence occurrenceOf Individual
Find the storage location of a Project	Project, Data Aggregation Node	Occurrence occurrenceOf Individual
Find all Projects containing an Individual	Individual, Project; Occurrence	Occurrence hasEvent Event
Find all Persons involved in a Project	Person; Administrator; Worker; Project	Event locatedAt Location
Find all the Evidence that supports an Occurrence	Occurrence; Token	Project storedBy DataAggregationNode Individual occurredIn Occurrence
Find all the Types of evidence that supports an Occurrence	Occurrence; Token; Type	Occurrence heldIn Project

Table 5.3: Competency Questions

The table shows that all of the competency questions outlined in Section 5.3.1 have been satisfied by the final ontology and each concept used maps quite easily with little to no change.

5.4.3 Comparing the Ontology with a Collection of Documents

Within the domain of scientific observations, Darwin Core is a popular choice for observations of wildlife and plants, it was because of this that we chose to use many of the pre-existing concepts from Darwin Core in K-HAS.

In the data driven approach suggested by [28], a corpus of documents, related to the domain that the ontology covers, and the keyword content is matched with the terms used in the ontology. Because Darwin Core observations are structured into archives of files, explained in Section 5.1.1, we evaluated the K-HAS ontology by combining the approach of comparing a corpus of documents related to the domain with human evaluation of ensuring an ontology met a set of requirements to create a method that ensured existing scientific sensed observations could be mapped to K-HAS with little to no modification.

As explained in Section 4.3.1, an archive of Darwin Core files consists of a minimum of 3 files: a metadata file that contains information about the creator of the archive and the project it relates to (eml.xml), a metadata file that describes all of the files that contain the occurrence and the fields within them (meta.xml) and a csv file that contains the data relating to the occurrence itself. Within the archive, the core files that need to be mapped to K-HAS are the meta.xml and eml.xml files as this allows us to store what and who.

We used files from a DwC archive made available online to extract the terms associated with an occurrence and we mapped these terms to K-HAS concepts, the results can be seen in Table 5.4. This table also shows how K-HAS terms could also be mapped to the SSN ontology.

Darwin Core Term	K-HAS Concept	SSN Term
occurrenceID	Occurrence	Observation
basisOfRecord	Basis of Record	-
recordedBy	Person/Node	Observed By
associatedMedia	Token	Observation Value
eventDate	Event Date	-
eventTime	Event Time	-
locationId	Location	Region
scientificName	Individual	Feature of Interest
identifiedBy	Person	-
dateIdentified	Identification	-

Table 5.4: Evaluation of K-HAS against Darwin Core Occurrence Terms

Each term within a Darwin Core observation maps to K-HAS with only a few changes to the terms, while the definitions do not change. As K-HAS is an alignment ontology, there are extra concepts that do not map to Darwin Core, but can encapsulate each occurrence. For example, a project in K-HAS can contain many occurrences. As previously mentioned, [103] shows how a DwC archive can be represented in OBOE, which means that it would not be too complex to map an OBOE observation in K-HAS.

5.5 Conclusion

In this chapter we have presented an ontology for the K-HAS architecture and described our methodology for development, as well as showing that existing ontologies are not complete enough to cover our requirements for K-HAS. The K-HAS ontology we present is an alignment of ontologies that are spread across multiple domains and provides a complete solution for a sensor network that deals with scientific observations. We have also extended this alignment to include concepts that model the unique features of K-HAS, but this ontology should be suitable for any sensor network that deals with observations.

While the main benefits of this ontology is that it provides a high-level overview of our network architecture and allows those that want to implement K-HAS to re-use all, or

some, parts of the ontology, it is also used in the network itself to allow sensed data to be mapped to data stores, define how classifications should be made and determine if certain users have permissions to edit an observation. Using the ontology, we can ensure that all nodes use the same data standard and use a pre-defined user model that can be easily implemented in variations of K-HAS. On top of this, the competency questions, outlined in Table 5.3 also form the basis of queries that can be carried out on sensed data, nodes, users and projects within K-HAS, linking them all through a common model.

Our evaluation has shown that the ontology is logically consistent and that existing scientific observations, in other formats, can be mapped across to K-HAS with few changes. The competency questions we identified during the design phase of the ontology can all be satisfied by the resulting ontology and our K-HAS network currently stores data with a schema that follows the design of this ontology.

In the future, work could be done to make the Darwin Core terms more modular and users will then be able to ‘plug in’ their occurrence structure of choice. Currently, this ontology has been developed with the scientific observations of our motivating scenario in mind. The K-HAS, FOAF and SensorML modules can all be reused in almost any WSN but not all networks will be tasked with recording individuals. In these cases, DwC (and some K-HAS) terms would need to be replaced with terms suited to the task of the network, such as flood monitoring.

The tiered architecture of K-HAS means that observations are classified at each step, with DC nodes applying their limited knowledge at the time of capture, DP nodes processing the data in-depth and applying knowledge from previously sensed data and humans either confirming or modifying these classifications. Currently, this only updates the observation as it passes through the network but, using this ontology, we would like to implement a provenance model that supports the classification of an observation at each stage of the network and edits made by humans once it has reached a DA node. In [80], provenance is defined as referring to “the sources of information,

such as entities and processes, involved in producing or delivering an artifact” and the W3C Open Provenance Model (OPM) [80] has been developed to “support a digital representation of provenance of any “thing”, whether produced by computer systems or not.” In future versions of this ontology, and our architecture, we would like to implement this model to support a complete history of edits made to an observation and show a detailed flow of knowledge as it passes through the network.

Chapter 6

Deployment in the Field

In this chapter, we detail the design and deployment of a modified K-HAS in the Malaysian rainforest. We also highlight the issues we experienced deploying the architecture in a humid, dense rainforest for use by those without domain knowledge of WSNs.

Experiments carried out in the rainforest had already shown that the range of wireless communications could be reduced by up to 80% and we expected that the lifetime of nodes in such conditions would be affected, based on experiments explained in Section 3.3.

Yearly visits were made to the Danau Girang Field Centre (DGFC) to test different nodes, gather knowledge and trial iterations of K-HAS. The first visit showed us just how much the rainforest affected communications, but also allowed us to gather knowledge from researchers and cameras that had previously been deployed. Subsequent visits were then used to test our own nodes and software, based on the knowledge we had gained from the first visit, explained in Section 3.3. We developed K-HAS with a view to deploying it in Malaysia, however, it soon became clear that our design was beyond what was currently available, as well as the time constraints.

Developing a camera with both wireless and processing capabilities, as well as being waterproof, proved to be an extremely difficult task, while wireless wildlife cameras were already commercially available with range much higher than what our experiments had yielded. This could be because of their test environments being places such

as American forest with less dense foliage, lower canopies and much lower humidity and using wireless technologies that are not licensed for worldwide use. For our own node designs, we attempted to use various node types connected to a camera, such as the Raspberry Pi and Wasp mote nodes; both yielded problems with mounting an SD card that was readable by both camera and node. At the time, we were unable to create a camera combined with a node that could be left untouched for months at a time in a rainforest.

However, we still needed to implement a network to evaluate our approach, thus we developed a modification to the K-HAS that utilised the latest commercially available hardware to provide an architecture that provides similar capabilities. Because these changes were made due to issues with deployment in Danau Girang, we chose to name it LORIS, after a famous animal in Malaysia. A Slow Loris is a small, nocturnal primate commonly found in South and South-east Asia but in our context it is an acronym that stands for Local-knowledge Ontology-based Remote-sensing Informatics System.

LORIS has been developed specifically for the Malaysian rainforest, our motivating scenario and, as such, much of the hardware and software has been implemented to reflect this.

The rest of this chapter is structured as follows. Section 6.1 highlights the changes we had to make from K-HAS and explains the hardware used at each tier. Section 6.2 explains how we planned and deployed the network. Section 6.3 details the results from the deployment in Malaysia and Section 6.5 concludes the chapter with a summary of our results.

6.1 Deployment Architecture

The aim of LORIS was to keep as much of the architecture of K-HAS as possible and reuse the ontology without the need for modification; serving as a form of validation. In order to do this, we identified the tiers of the network that could be deployed in the

rainforest, given the hardware and wireless constraints, and the areas that needed to be addressed. Regular power and cheap computers with good knowledge-processing capabilities meant that DA nodes were simple to deploy and, with the growth of powerful microcomputers like the Raspberry Pi, DP nodes were also in abundance. However, finding a reliable camera that could integrate with a node capable of processing observations that was also reliable enough to withstand the humid rainforest proved to be difficult.

6.1.1 Data Collection

Experiments were run yearly in Malaysia to test the performance of variations of hardware. As covered in Section 3.3.1, the first year involved testing the range of Wi-Fi with an IGEP board. The limited range of 30m meant that, despite the high transfer rate, it was not possible. Wi-Fi was not designed for use in resource-constrained sensors, so the power consumption of the radio meant that it limited the lifetime of the node.

The following year, we used the Digimesh protocol, explained in Section 3.3.2, which provided a longer range and was developed for use in sensor networks. While the range was suitable for the rainforest, it was not as much as we had anticipated. The other difficulty proved in mounting an SD card on two devices at the same time, the existing wildlife cameras and our Wasp mote nodes, without causing corruption from read/write clashes.

From these experiments, we began to look into commercial alternatives that combined both the node and the camera. The most usable solution we found was the Raspberry Pi coupled with a camera attachment [114], but it was not ready for use in the Malaysian rainforest or to be used for an extended unsupervised period. Existing wildlife manufacturers were then looked into and we found that wireless cameras had been manufactured, but they had no local processing capabilities and had not been used

much in research. Reconyx had created a cellular camera that, in Malaysia, could send images via SMS [95], but cellular coverage on the rainforest floor is not common and the cost of the camera was high. Based on these findings, we used Buckeye X7R cameras [25], shown in figure 6.1. Using the same Digimesh protocol as the Waspmove, they had a tested range of 1 mile and their casing had been developed to withstand harsh environments. The 1 mile range was tested in a more open woodland in the US but the higher quality components used, as well as their modifiable antenna supported our theory that we would achieve better range with their cameras.



Figure 6.1: Buckeye X7R and Battery Pack

We expected to achieve around 800m of range and these proved to be suitable DC node replacements, despite the fact that we had to forgo local processing or the storage of any local knowledge.

6.1.2 Data Processing and Aggregation

Because our DC nodes were wildlife cameras that did not have any knowledge processing capabilities, and because our DC node sites were near enough to the field centre, we combined the DP and DA nodes into one desktop stored at DGFC, that contained both a Digimesh radio and an Internet connection. The benefit of using commercial grade hardware was the software that accompanied it; remote management and configuration software allowed users to modify the settings on each camera, as well as handling the retrieval of images from every node deployed. EXIF tags written to images could be modified, as well as how many images were captured for each observation. The software had been created to be used by those without any specialist knowledge and was easily used by researchers in the field centre.

Each observation was saved into a directory that matched the ID of the camera it originated from and this meant that the existing software used in K-HAS could be used without modification. Sensor middleware, such as GSN, is running on the same machine and virtual sensors listen for changes in each directory; where each virtual sensor represents a deployed camera.

When new observations are detected, Drools runs on the images to start metadata and EXIF processing. Some rules had already been extracted from the processing of 120,000 images collected from our yearly visits to DGFC; other rules had to be added by users of the WSN. When an observation had been processed, rules were run once again to determine if a match could be made to existing projects and, if so, who should be notified.

The web interface for classifying and viewing all observations was accessible by those within DGFC, as well as uploading rule files and new locations. Using this interface proved to be easier for those without a WSN background and the GSN interface had a steeper learning curve. Figure 6.2 shows the basic metadata of an observation, outlining where it was captured and when; as well as the images themselves.



Figure 6.2: LORIS Web Interface

6.2 Deployment

In June of 2013, a three week visit was made to Danau Gurang, with three Buckeye cameras and the software required to deploy LORIS. The network was first tested within the field centre and one of the Buckeye cameras had been broken during transport, giving us only two cameras. Due to the protected nature of the forest, we were unable to nail any cables to trees to use the high gain antenna and it was not possible to use the cables without first securing them, because animals tampering with cameras was a common occurrence.

For the initial week of the deployment, we wanted to test the robustness of the network before focussing on the data. This meant placing the cameras in locations where they would be triggered often and in an area where they would be affected by rainfall and humidity. Figure 6.3 shows the locations of the cameras during both of the weeks that it was deployed. Buckeye 1 was deployed for the first week along the main path leading to the field centre from the river, this experienced the most traffic due to the number of

researchers present at the time.

The cameras were initially configured to send a single image for each trigger, as opposed to the three that is used by the current Reconyx cameras at DGFC, because we wanted to test range and data rate for a single transmission. This was fine for images of humans but we found that triggers caused by animals could miss and have no content. We therefore changed the number of images to two so that network traffic was still minimised, especially as images could only be sent when power was available at the field centre.

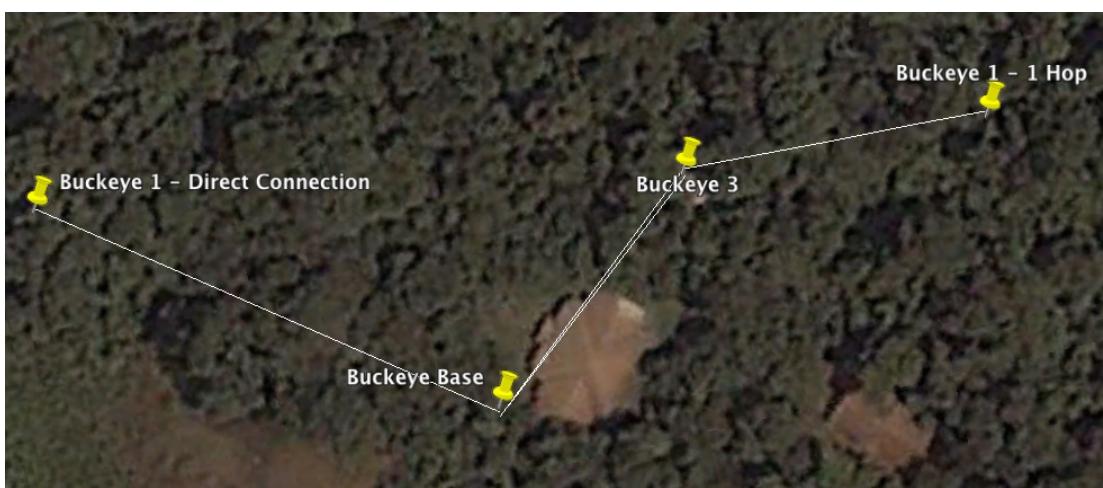


Figure 6.3: Camera Locations Around Danau Girang

6.2.1 Direct Connection

For six days, we tested two Buckeye cameras that were near enough to the field centre to maintain an active connection to the base station. This meant that no hopping was involved and cameras were not reliant on each other to transmit images. Buckeye 1 was placed on a main path that guaranteed human foot traffic and Buckeye 3 was placed on a trail in the forest that, while experiencing minimal foot traffic, was expected to yield small mammals and birds.

For the duration of the six days, the cameras did not report a drop in battery levels and

a total of 1076 images were sent. Buckeye 1 was in a more open environment, despite being further away, and average speeds of 4KB/s were achieved.

Transmissions from Buckeye 3 were significantly hampered by the dense forest that it was surrounded by, as well as the fact that a low-gain antenna was used. We also speculate that the field centre itself acted as a barrier to the signal, as the base station was located on the opposite end. Because of this, we experienced an average speed of 1.2KB/s, taking around three minutes to receive an image. This is consistent with experiments conducted in previous years where dense, humid forest has led to a decreased range of, up to, 78%. Lower frequencies do experience a longer range, but the data rate is significantly impacted.

6.2.2 1-Hop Network

For a further five days, Buckeye 1 was moved further into the forest and Buckeye 3 was set up as a routing camera that forwarded images onto the base station, 63 images were taken during this period.

The location of the moved camera is also shown in Figure 6.3. As was expected, the traffic of the network reduced significantly when the cameras were moved from the main path, with almost all 63 pictures being of animals. Due to animals moving a lot faster than humans, the number of pictures taken per trigger was increased to two, this would ensure network traffic was not overwhelming while increasing the chance of capturing the subject.

Surprisingly, the speed of transmission from Buckeye 1 was faster than Buckeye 3, despite being routed through it, with an average speeds of 1.8KB/s.

6.3 Results

Deploying a network for two weeks is not a robust experiment and does not show that LORIS can withstand months of running without human intervention. However, it does serve as a proof of concept that, using commercial hardware, a modified K-HAS network can be implemented and utilise the knowledge of its environment.

The pictures of humans from the first six days are not of use for the current motivating scenario, focussed on animals in the rainforest corridor, but one could see this network also being tasked to warn the authorities of hunters in restricted areas of the rainforest. However, the images we have of animals, while few are animals of interest due to the proximity of the cameras to humans, were processed and used to create templates for future observations.



Figure 6.4: First Image of Lizard by Buckeye 3

More interestingly, we were able to infer further rules from the short deployment that

helped us to narrow down the choice of animals to match to based on the time of day and location. For example, Figure 6.4 shows a lizard crossing the path of Buckeye 3 in the middle of the day and Figure 6.5 shows the lizard walking past in early afternoon on a different day. This allows us to encode a window of time in which the lizard is more likely to appear onto the static knowledge base of DC nodes or, in the case of LORIS, onto the DA/DP node. Images an hour later were captured by the camera, showing the lizard walking in the opposite direction. While this was not every day, and the deployment time was too short to establish a pattern, it does help us to infer that a lizard passing by in early afternoon is likely to walk the same path approximately an hour later.



Figure 6.5: Second Image of Lizard by Buckeye 3

6.4 Rules

This section explains the rules used in LORIS and K-HAS in more detail, to highlight how classifications are made as well as how nodes in the network are monitored. Any examples made here relate to our motivating scenario. The rules on each node are a vital way of encoding local knowledge and a practical use to show how local knowledge, from users or previously sensed data, can be used classify newly sensed data.

As highlighted in Section 4.3.4, K-HAS uses the Drools rule engine, a Java based rule management system that uses forward chaining based inference. Forward chaining starts with some data, in our case a set of images, and uses inference rules to extract more data. We use this method to make meaningful inferences from properties of the sensed data. For example, an image taken at 8pm could be run through Drools and an inference rule will tell the system to start by looking for nocturnal animals.

6.4.1 Data Processing and Data Aggregation

The increased knowledge-processing capabilities allow DP nodes to run a complete implementation of Drools, this means that forward chaining can be used and the knowledge base is dynamic, i.e. it can be updated whilst the network is deployed.

Rules are written in *drl* files that are loaded into a knowledge session. Sessions, as well as the firing of rules, are handled by the Java code. The benefit of using Drools with a Java based middleware, such as GSN, allows the simple integration of the two technologies, allowing a sensor class in GSN to insert an archive into the session and run the rules. When sensed data is received from a DC node, it is unarchived and Drools is run on the DwC files describing the original observation, the metadata, the originating DC node and the knowledge base of the DP node.

Listing 6.1 shows a psuedocode example of some more simple rules that could be contained within the knowledge base, a full listing is in Appendix D. In a rules file,

existing local knowledge is used with the properties of a new observation to infer the contents and then update the DwC archive, or return suggestions as to what the contents may be. For example, the sample rule uses the location ID of the camera to check if it is near a plantation and uses local knowledge of the area and global knowledge of the time to infer that the image may contain a sun bear or Malay civet, based on the knowledge that they have been known to forage in the plantation at night when there are no humans. This rule maps to the terms defined in the ontology (Chapter 5) where the individual is a sun bear or a Malay civet, the evidence is an image and the occurrence is the DwC archive created by the triggering of a DC node.

```
1 rule "Plantation night"
2   when New_Sensed_Data
3     if location of node is Plantation and time of day is night
4       write to file set.CSV
5         dateIdentified = \today
6         identifiedBy = node ID
7         Print "Likely a sun bear or Malay civet"
```

Listing 6.1: Pseudocode of a Drools Rules File

These rules are very simple and only scrape the surface of the local knowledge that can be utilised, as well as how DwC archives can be effectively used to carry the data of these inferences, but it does show that Drools is a powerful rule engine and that these rules can be extended further to make full use of the knowledge base.

DA nodes use the same rule engine as DP nodes but rules are used for the post-processing of Darwin Core archives. For example, when a classified archive is received from a DP node, Drools is run to check if the classification matches an existing project. If it does, then it searches for those that are subscribed to the project and notifies them via their selected medium. Drools can also be used to monitor the nodes deployed in the network and check their status. Rules are run to check for archives from DC nodes and, if a node has not sent an archive in a set number of days, then administrators are informed that it may have run out of battery.

More importantly, rules are vital for the exchange of knowledge between nodes in the network. If a user classifies an observation on a DA node, then rules are fired that

identify the DP node that sent it and update the knowledge base, on the original DP, so that the image template it uses in future classification matches the one identified by the user.

6.4.2 Existing Data

Throughout our visits to Danau Girang, we collected local knowledge through field work and interviews with the researchers based at Danau Girang. As well as this, we collected species classifications made by the researchers for images that had been taken during camera deployment. Using these data, we have extracted patterns in animal movements in order to create some basic rules for the DC and DP nodes. In this section, we will describe the data we have and the methods used to extract rules.

Semi Structured Interviews

On our final visit to Danau Girang, we designed a set of 14 semi structured interview questions that we asked to ten researchers based at the field centre, with the majority working on different projects. We used these questions to gain insight into common trends between projects and patterns recorded about the animal subjects. A subset of the questions have been listed below:

- What species are you looking at?
- Do you have specific sites that you look into?
- What specifics do you know about the target species?
- Do others use your research? If so, how?

These questions have been designed to be as open as possible to allow the interviewees to provide as much detail as they wish, or even transition to a different topic. The

interviews were recorded and later transcribed to text. We then used a software package called Dedoose [33], which is an online tool that assists with performing manual qualitative analysis on text based documents. Uploading our documents to the service allowed us to manually review the responses to each question and highlight, or ‘tag’, excerpts of interest and link to other excerpts that had either been mentioned by another interviewee; or were related to other comments. Each excerpt was tagged with the subject it was discussing, for example, a comment on the accessibility of an area of the forest during certain times of the year were tagged as *local knowledge*. The excerpts were combined and exported into a CSV file that detailed who the interviewee was, the content of the excerpt and the tag. Appendix E shows a complete interview transcript with a researcher about the clouded leopard and Appendix F provides an extract of the same interview, once it has been annotated within Dedoose. The extract included shows an example of questions that were spawned from asking about specific sites and exploring further based on the answers given and show highlighted texts that have been tagged and coloured, by Dedoose, based on the tag. The extract can be seen in Appendix E from lines 22 to 35.

These face to face sessions gave us insight into patterns and observations that researchers had learnt during their months, or years, at Danau Girang; something we would have otherwise been unable to learn in our short periods at the field centre. An example of this is learning that the clouded leopard is not nocturnal around Danau Girang, while it is in the rest of the world. It is common knowledge that clouded leopards are nocturnal animals but they have been seen at all hours of the day around DG. Researchers believe that this could be due to the fact that the rainforest is secondary, growing back after heavy logging 40 years ago, or human impact from palm oil plantations that border the edges of the rainforest or the availability of prey. Whilst the local knowledge that clouded leopards can be seen throughout the day cannot be directly encoded as a rule, it does mean that the animal can be removed from rules that may filter out nocturnal animals as possible classifications for images captured during the day.

While we were not able to construct a full set of rules from these findings, the results have provided support for patterns extracted from existing data that has been classified by researchers, which we discuss further in the next section.

Existing Data

When images are manually collected at Danau Girang, researchers process each image set and they have recently been recording the classifications. However, these classifications are only being recorded for currently ongoing projects and other images are ignored. During one of our visits, we collected a CSV file of 2650 confirmed species classifications made by researchers. The data had not been cleaned and multiple users had used different names for the same species, we cleaned the data manually and matched it with the images in our database. We then used the classifications for each species to determine any patterns by analysing the details recorded for each observation, which are listed below:

- Time
- Date
- Location
- Temperature
- Moonphase
- Classification

SQL queries provided these details for each species and we recorded patterns that could be identified. For example, if a Samba Deer was only spotted between the hours of 7pm and 4am then we could write a rule that the species is nocturnal. More detailed rules can be created if there are patterns across multiple details. A species that is only spotted in two of the twenty sites between the hours of 7pm and 4am means that nodes

with local knowledge can make accurate classifications without the need for image processing. However, these rules are based on large amounts of data collected from an active, or previous, deployment. DC nodes in their early deployment stages would not have the data to make such detailed rules.

Stored procedures were created to join the results from multiple tables and manually explore patterns in the data for the individual species. However, many species either did not have enough observations or the observations they did have had so much variation that there was no discernible pattern. If this work were taken further then we would like to have our entire dataset classified by professionals and Danau Girang is currently undergoing this process for all images taken since the camera trapping project began.

```

1 if TIME > 1500 and TEMP <= 26:
2   if TIME < 0000:
3     if MOONPHASE <= 2:
4       CLASSIFICATION = GOAT (25\% chance)
5     if MOONPHASE > 2:
6       if DATE in JAN:
7         CLASSIFICATION = HUNTER (2.6\% chance)

```

Listing 6.2: Example Rule created from Existing Data

Using the Weka package [54], we initially constructed a J48 decision tree (an open source Java implementation of the C4.5 algorithm [90]) but found that the accuracy was only 27% and the rules extracted from the model related to individual times that had only seen a single observation. From this, we then used a decision table [64] within Weka that created a model yielding an accuracy of 53%. We used the resulting model to extract a collection of 281 rules that could be run on a DC node. Figure 6.2 shows a rule that was created from the output. This rule checks the time of capture for the observation, the temperature and the moonphase; which has been converted into a numeric value. If the temperature is fewer than 26 degrees and the time is between 3pm and midnight, then there is a 25% chance of the species classification being a goat. The if-statements are executed in order and the classification that matches the properties of the observation, and has the highest percentage chance, is forwarded to a DP node.

With Weka, we experienced the same problem as with SQL, we only had 2700 classifications made by people with domain knowledge, so any models generated from Weka

were limited not just to those species in those classifications, but to those species with a sufficient number of observations. Of those 2700 classifications, there were only 45 different species and some had fewer than 10 sightings. For example, the set we used primarily contains goats and the resulting model provided 84% accuracy when given an observation of a goat. However, the Malay badger had very few sightings in the set and the model only yielded 17% accuracy. With future work, we would like to add new classifications to this model and test it extensively with new, unclassified data, as well as add more variables on top of moonphase, temperature, date and time but ensuring they are variables that would be available to data collection nodes.

Because the features used to generate the rules are available in every observation, and do not require any external information, DC nodes are able to process the series of if-statements quickly. This method of knowledge-processing comes at the cost of accuracy, when compared to using existing data, image processing and/or a dynamic knowledge base, but the speed and simplicity of these rules mean that they can be used by almost any node, regardless of computational capability.

6.5 Conclusion

K-HAS was developed, with the motivating scenario in mind, as an ideal, general-purpose network architecture that is able to utilise the knowledge of its environment. However, experiments to support the design and deployment of such a network within the timeframe of this PhD showed that the current technology was not ready to support a network that would run without human intervention for extended periods. LORIS is our more specialised, cut down approach that can be implemented using commercial hardware that is readily available, as well as software that we have open-sourced.

Our deployment in Malaysia has shown that local knowledge can assist with classifications and can be gained from even a few images collected within the first few days of deployment. Ideally, we would like to leave this network running for an extended

period to infer rules from the sensed data and test the reliability of the network in harsh conditions. From both the visit and the deployment, we have offered evidence for the existence of local knowledge and provided a method that can be used to extract it from field experts.

The rules extracted from the interviews and classified data we collected while deploying LORIS can be used as a static knowledge base for DC nodes in the K-HAS architecture, or as the basic starting knowledge base for the DA node in LORIS. As more classifications are made to new sensed data, this same process can be used to create new rules for previously unrecorded classifications, as well as refine existing rules.

The testing of these rules have been drawn from a small set of existing data collected from Danau Girang, although we believe the results are promising enough to show that, in principle, local knowledge could be used to make accurate classifications on sensed data and a larger dataset could generate more specific, more accurate rules in order to do so.

LORIS has shown that local knowledge can be used to enrich sensed data, and automate classification, when it has been received at the field centre, but we maintain that K-HAS' aim of pushing local knowledge right out to the edge of the network makes a network more effective and allows for a more timely delivery of important sensed data. The main drawback of LORIS is that it relies on the chronological delivery mechanism used by the commercial cameras, whereas K-HAS can improve this mechanism by sending the data that it infers to be of a higher importance, rather than what was simply captured first.

Chapter 7

Simulation Experiments and Evaluation of the Architecture

In this chapter, we present the results of simulation experiments developed to evaluate and explore choices offered by our K-HAS architecture. Using nodes with knowledge processing capabilities to deliver interesting data quicker than a standard WSN, we have developed a simulation for our proposed architecture, K-HAS, as well as variations on the knowledge processing capabilities of the nodes at each tier.

The simulations were developed to determine whether K-HAS is the best mix of data processing and data collection nodes that splits between minimal knowledge processing to prioritise from the edge and high knowledge processing towards the centre to provide more accurate prioritisation, thus minimising the transmission time of interesting data. This was done by using a network structure, that matched our motivating scenario, and changing the knowledge processing capabilities on each node at every tier; ranging from no knowledge on every node to the maximum knowledge processing capabilities across the network. We aim to show that the more knowledge processing capabilities that are pushed out towards the edge of the network, the more effective the network becomes at prioritising data that it believes to be interesting and delaying what it believes to be empty.

This chapter is structured as follows. Section 7.1 describes the tool used to develop the network. Section 7.2 outlines the scenarios used in our experiments and Section

7.3 describes the parameters used to configure each run and what was observed. Finally, Section 7.4 explains the results from multiple runs and Section 7.5 concludes our findings and highlights areas that require further experimentation.

7.1 Simulation Environment

Using RePast Simphony [29], an agent-based network simulation tool developed in Java, we created a network to emulate K-HAS. RePast is an agent-based modelling system that allows for agents to be created and placed on a grid. Ticks denote a period of time and simulations can run for a fixed number of ticks, or until stopped. Ticks can also be used to schedule events, such as searching for neighbours, by calling methods that last for a set number of ticks, or begin at a particular tick. For example, a camera sensor node may be tasked with taking a picture every three hundred seconds. When the simulation reaches three hundred ticks (or six hundred, nine hundred, twelve hundred and so on), a scheduled event is run to simulate the camera’s capture of an image and transmitting it to an endpoint.

RePast was chosen because we did not require the low level network configuration provided by other tools, such as NS2 [78] or the MultiAgent Simulation Environment (MASON) [70], but we did need to modify and record the behaviour between nodes as they capture and process sensed data. RePast’s event scheduling allows for nodes to be modelled as agents and the dynamic configuration allowed us to modify the simulations during run time. The aim of these simulations was to visualise how different knowledge processing capabilities can affect the prioritisation and transmission time of observations, as well as the accuracy of their classifications. RePast allowed us to utilise existing Java code we were using in our K-HAS middleware and develop a base simulation that could be configured easily with XML files.

Simulations are shown in a GUI that allows the network to be visualised and configured, with results exported to other applications, such as MATLAB and R.

7.2 Simulation Scenarios

While the aim of these simulations was to show the effectiveness of K-HAS over the current solution, we also wanted to determine if it was the best tradeoff between factors. We believe that the solution to reduce transmission time of interesting images and reduce the number of undelivered observations would be to attach nodes with high knowledge processing capabilities to all cameras in the network, however we experienced shorter battery life with more powerful nodes and that could mean that replacements would be made as often as the current manual solution, detailed in Section 3.1.

Throughout this chapter, we will be referring to nodes tasked with different purposes as the three definitions listed here. K-HAS is a tiered network architecture, with a hierarchy similar to that which is described in Section 2.2.3, below we have described the nodes at each tier :

- Sensing Node (SN): A node that has been tasked with capturing sensed data and forwarding from other sensing nodes to a processing node.
- Processing Node (PN): A node that acts as the cluster head for sensing nodes; it is tasked with receiving sensed data from sensing nodes, processing it and forwarding to the central node.
- Central Node (CN): A node with similar functionality to a typical base station, tasked with storing all sensed data and providing an interface to users.

At the processing and sensing tier, the degree of knowledge processing capabilities can range from the levels outlined below:

- No Knowledge (NK): The node possesses no knowledge processing capabilities.
- Minimal knowledge (MK): The node possesses basic knowledge processing capabilities and contains a static rule base (Section 4.2.1).

- High Knowledge (HK): The node possesses high knowledge processing capabilities and is able to process data, metadata and use a dynamic rule base (Section 4.2.2).

The scenarios we have tested show various combinations of knowledge within the network, as well as the combination of knowledge processing at the centre, as a baseline. We can list all the scenarios in a triple that highlights the knowledge level at each tier. The leftmost descriptor is the knowledge level for sensing nodes, the middle describes processing nodes and the rightmost is the knowledge level for the central node. The list below describes each scenario and are ordered to highlight the progression of pushing knowledge from the central point of the network out to the edge:

- NK-NK-MK: Sensing and processing nodes have no knowledge processing capabilities with minimal processing capabilities on the central node.
- NK-NK-HK: Sensing and processing nodes have no knowledge processing capabilities with high processing capabilities on the central node.
- NK-MK-NK: Sensing nodes possess no knowledge processing and processing nodes have minimal knowledge.
- MK-NK-NK: Sensing nodes possess minimal knowledge processing capabilities.
- MK-HK-NK (K-HAS): Sensing nodes have minimal knowledge and processing nodes possess high knowledge. This scenario matches the K-HAS architecture we proposed in Chapter 4
- HK-NK-NK: Sensing nodes have high processing capabilities.

The higher knowledge processing capabilities of HK nodes allow them to classify observations with a greater accuracy, but their battery life is much shorter than MK nodes

due to their increased power needs (Section 3.2). In contrast, MK nodes can run for a longer period without requiring battery replacement. MK nodes, however, are unable to classify observations to the same level as HK nodes. While HK nodes can classify an observation as interesting or empty, MK nodes can only assume an observation is interesting using the image’s metadata as they lack the capabilities to reliably determine whether an observation is empty or not. The scenarios we have implemented cover combinations of HK, MK and NK, in a twenty five node network. We use twenty five nodes because DGFC had between twenty and twenty two active cameras during our first and second visits and we knew that the first implementation of the network would require a single endpoint. From this, we chose to use four processing nodes so that each could handle an equal number of sensing nodes, if they were all within range, and a single central node. Our experiments in Danau Girang were restricted to a smaller number of nodes, due to cost, but we expected to deploy a sensing node onto all of the active cameras. The hierarchical nature of the network is shown in Figure 7.1, with green nodes representing sensing nodes, yellow nodes representing processing nodes and the red node representing the central node. Blue arrows depict slower Zigbee links between sensing and processing nodes and the yellow arrows represent faster Wi-Fi links between processing and central nodes.

Before implementing, we designed the agents required based on the nodes described in the ontology we proposed in Chapter 5. Using that, we created a hierarchy of nodes inheriting common properties from a node object. As previously mentioned, we had metrics on range and transmission times from previous experiments and the deployment of LORIS. We used these to create properties for each transmission medium that could be used by each node object. Table 7.1 shows how the K-HAS terms, introduced in Chapter 4, and LORIS terms, detailed in Chapter 6, map to the node types described above. The nodes used in our simulations map directly to the ontology.

Using a Java library we developed for DwC archives during the implementation of LORIS (Section 6.1), we were able to implement DwC archives as the data standard in

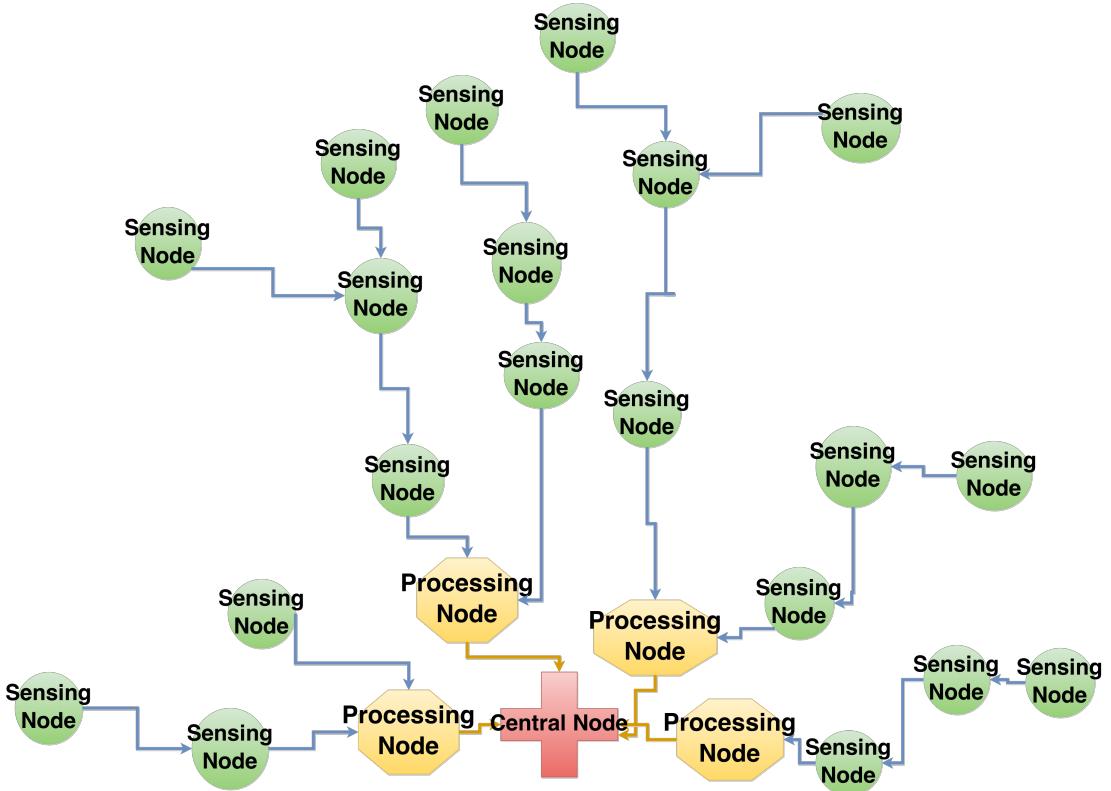


Figure 7.1: Simulation Example in RePast Simulator

Network Type	Original Term	Maps To
K-HAS	DC Node	Sensing Node + MK
	DP Node	Processing node + HK
LORIS	DA Node	Central Node + HK
	Buckeye	Sensing Node + NK
	DA and DP	Routing/Central Node + HK

Table 7.1: Mapping of K-HAS to Simulation Terminology

our simulations, this allowed us to model the prioritisation of sensed data, as well as the classification of observations, in the same way that a K-HAS deployment would.

7.3 Network Setup

7.3.1 Parameters

The parameters used in these simulations were either gained through experiments performed at Danau Girang, gleaned from the data or extracted from technical specifications. Repast allows these parameters to be set in a configuration file and can be manipulated once loaded into the simulation. An example of this file can be seen in Appendix H. The table below outlines the parameters we have used and some that require further explanation are covered in this section.

Parameter	Description	Value(s)	Constant?	Observed?
SN Count	The number of sensing nodes in the network	20	Yes	No
PN Count	The number of processing nodes in the network	4	Yes	No
CN Count	The number of central nodes in the network	1	Yes	No
Height	The height of the available space	1200	Yes	No
Width	The width of the available space	1200	Yes	No
Capture Chance	The chance of an observation being captured each second in normal circumstances	0.000857703189	Yes	No
Interesting Chance	The chance of an observation containing an item of interest	0.207	Yes	No
Transmission Rate	The transmission rate of sensing nodes in the network	Ideal or Variable	Yes	No
SN Transmission Rate/kbps (Zigbee)	When <i>Transmission Rate</i> is set to <i>ideal</i> , this is set as the maximum. <i>Variable</i> sets this randomly between 20 and the maximum	20-250	No	No
SN Transmission Range/metres (Zigbee)	The transmission range of Zigbee	30	Yes	No
PN Transmission Rate/kbps (Wi-Fi)	Always set as the maximum due to the proximity to the central node (kbps)	54000	Yes	No
PN Transmission Range/metres (Wi-Fi)	The transmission range of Wi-Fi (metres)	200	Yes	No
Bandwidth	Determines the available bandwidth in the network, if ‘saturated’ then the chance of capture is increased ten fold	Normal or Saturated	Yes	No
HK Processing Time/secs	Time to process an observation with high knowledge processing capabilities (Central Node or Processing/Sensing Node)	30 (CN) or 43 (PN/SN)	Yes	No
MK Processing Time/secs	Time to process an observation with minimal knowledge processing capabilities	5	Yes	No
Observation Size/KB	Size of the generated observation, the combination of 3 randomly generated image sizes	1362-2598	No	No
Transmission Time/secs	The time an image takes to pass through the network, from point of capture to being received by the central node	-	No	Yes

Table 7.2: Parameters Used to Configure Repast Simulations

Transmission Rate

Each scenario, listed in Section 7.2, can be run with either a *variable* or *ideal* transmission rate. This means that the transmission rate of a sensing node can be fixed at the maximum rate at the start of each run or a randomly generated number (between 20 and the maximum) is set at the start of each run. The maximum value was taken from the Zigbee technical specifications.

Bandwidth, Capture Chance and Interesting Chance

As with transmission rate, each scenario can be run with the bandwidth under *normal* or *saturated* load. When the bandwidth is normal, the chance of capture is the value shown in Table 7.3.1, but saturated increases that chance ten fold. A random number is generated on the sensing node every second and compared with the capture chance, generating an observation if the value is less than the chance.

Using the existing data collected from Danau Girang, we calculated how often a camera triggers in a six month deployment, as well as how often the observation contained interesting content.

To calculate the count of interesting images, we processed every directory of images to extract the largest object in the foreground, using our Triton program. Once processed, we iterated through every directory, counted the total number of images and the total number of extracted images; an extracted image is a black and white image containing the largest object that has been found in the observation and are only created when the processing believes that the observation is interesting. This gave us a 20.7% chance of an image being interesting, across every camera.

The chance of a camera being triggered each second was calculated by the total number of observations (13,399) divided by the number of seconds in six months (15,552,000). This gives a chance of 0.000861561 of a camera trigger in any given second.

Transmission Range

From our experiments at Danau Girang, see Section 3.3, we tested the range of Wi-Fi and Zigbee in a humid rainforest, using those values we set the transmission range as the maximum that we achieved during our experiments.

Observation Size

Using the observations we extracted from Danau Girang, we listed the size of all 120,000 files and extracted the minimum and maximum size from the set, when an observation is generated in RePast, it consists of three images and each image is a randomly generated size between the maximum and minimum bounds outlined in the table.

Processing Times

While the processing times are fixed, central nodes can perform higher knowledge processing slightly faster and they are able to process up to 4 observations at once, whereas processing and sensing nodes can only process one at a time.

Run Count and Duration

Each scenario is configured to run for a period of three months, this is solely due to time constraints due the run time of each simulation. Tests showed that three months was a long enough period to model a saturated network and the capture chance is kept at the six month chance because the network would be deployed for six month periods and these runs serve to show a snippet of their deployment, rather than cover the whole period.

Twenty five runs were completed of each scenario, as well as for the possible combinations of transmission rate and bandwidth. For example, a network simulating NK-NK-

NK would have 25 runs for *ideal* transmission rate and *normal* saturation, 25 for *ideal* transmission rate and *saturated* bandwidth and so on.

7.3.2 Routing Protocol

The routing protocol used needs to be dynamic in order to adapt to nodes being added and removed during deployment, while minimising traffic in a resource constrained network. In our approach, we use a modification of the Minimum Cost Forwarding Algorithm (MCFA), described in Section 4.3.5. A cost is assigned to each node, based on how far they are from the central node, with neighbouring nodes choosing to connect to the node with the lowest cost. However, in normal implementations of MCFA, all nodes are of the same type and simply need to connect to a central node. This modified MCFA is used in all scenarios.

In our K-HAS architecture, sensing nodes cannot connect directly to a central node because processing would not take place. Because of this, we used the same routing method across all scenarios. Our implementation of MCFA works with a discovery phase and a transmission phase. The discovery phase is a scheduled event, taking place at the start of deployment but it can be run throughout deployment to react to nodes being added or removed.

Discovery

Discovery begins at each central node, scanning nodes in range for processing nodes and sending a broadcast packet, with a cost of 0, to inform them that they are within range of a central node. Links between Central and processing nodes use W-Fi in all of our scenarios. Once received, processing nodes increment the count and forward the packet to any processing nodes within range of them, where we use the range of Zigbee. We found that this method overloaded the processing nodes and all sensing

nodes within range would connect to the first processing node they receive the broadcast from. We then implemented a method, called *load balancing* [48], which uses the sensing nodes connected to a processing node to calculate whether it should offload new nodes to a neighbouring processing node.

The maximum connections a processing node can have is determined by the total number of sensing nodes in the network divided by the total number of processing nodes, which is held in the knowledge base of the central node. Once a processing node has the maximum number of connections allowed, it starts to offload to a neighbouring processing node that is also in range of the sensing node requesting a connection. If there are no neighbouring nodes then the processing node exceeds the maximum number of connections allowed, to save sensing nodes being left with nowhere to send their data.

If the sensing node that receives the broadcast does not have an existing route to a central node, or the cost of the current route is higher than the received route, it adds an edge to the processing node, increments the count and forwards it to all nodes in range. This process continues until the broadcast reaches the edge of the network. Nodes do not have global knowledge of the route to the central node, only of their neighbour with the lowest cost. This can be seen in Figure 4.9 in Section 4.3.5.

This phase can be repeated throughout the course of the deployment, simply by scheduling it as an event to occur every n ticks. However, the simulation currently only uses the discovery phase at the beginning of the deployment.

A sensing node will always be at least 2 hops as all observations must go through a processing node. However, as seen in Figure 7.1, there could be more hops depending on the node's position within the network. All hops between sensing and processing nodes use the slower Zigbee transmission medium, to allow for a longer range, while the final link between the processing and central node uses a faster Wi-Fi connection.

Transmission

Once the discovery phase has been completed, providing nodes are within range of the central node, the transmission phase begins where only DwC archives are then sent across the network. Observations are captured based on the mode of the simulation and sent to the lowest cost neighbour.

In order to manage transmissions, nodes have a *SendState* object that contains the next archive to send, the time it will take to send it and whether it is currently sending. This is used to determine what operations to perform. Once an archive has been sent, it is deleted from the SendState and the sending flag is set to false. A new archive is then added and sent when the opportunity arises.

When a processing node receives the archive, it begins processing. processing nodes use the SendState as well, but they only add an archive once it has been processed and they then select the oldest archive that has been classified as interesting, providing an archive is not already waiting to be sent. The archive stores information about the route it takes, recording every hop, as well as the time it took from capture to central node.

Scheduled sending events run every thousand ticks, which is configurable, to check the sending state of the node and send any archives in the SendState. The node then waits for the number of ticks that it will take in order to transmit the archive.

Once the simulation is completed, either manually or through a defined number of ticks, the archives in each node are iterated over and written to a CSV file, with details such as the path it took, total transmission time and time of capture.

7.3.3 Observation Transmission

When a node has no knowledge, it simply sends the observation that was first added to its queue. With MK, it finds the first *interesting* observation and sends that, if one is not in the queue then an observation marked as *unknown*. HK performs much the same

as MK except it instead looks for *empty* observation if an *interesting* cannot be found. If a node has any form of knowledge processing power, it will not send an observation unless it has first been processed. Nodes can only send an observation one at a time and must also wait for the receiving node to be free.

An observation is marked as delivered once it has been received by the central node, if there is only processing at the central node, then it is only marked as delivered once the central node has processed it.

7.3.4 Processing

Repast has built in functionality that allows for methods to be scheduled to start at a certain tick and repeat at set intervals. The way a scheduled method works is that it is run at the time specified and further scheduled methods can be prevented from running until it has completed. Much of the node's functionality is handled in this way. A processing method is set to run every 10 ticks (seconds) to replicate the functionality of checking a file system for new observations. The earliest (unprocessed) observation is found and a processing method is scheduled, depending on the knowledge level of the node. The function for DA nodes is the same, except they are able to process 4 observations simultaneously. This is shown in listing 7.1 and listing 7.2 shows how classifications are made.

Observations are passed to the processing methods and those methods run for the specified durations. High knowledge processing nodes are able to determine an empty observation with 97% accuracy and detect an object of interest with 82% accuracy. A random number is generated and compared with these values, to determine how the observation is classified. Once the observation has been classified, the node marks itself as not processing to allow a new observation to be classified.

```

1 kL = getKnowledgeLevel()
2 if(kL == NoKnowledge OR isProcessing) {
3     return
4 }else{
5     obs = getEarliestObservation()
6     if(obs.identification IS NULL OR (obs.identification NOT NULL AND obs.
7         identifiedBy.getKnowledgeLevel() < kL)) {
8         if (kL == MinimalKnowledge) {
9             schedule method for minimal processing to run as soon as possible for 5
10            ticks
11            setProcessing(true)
12        } else if (kL == HighKnowledge) {
13            schedule method for high processing to run as soon as possible for 43
14            ticks
15            setProcessing(true)
16        }
17    }
18 }
```

Listing 7.1: Processing Function for DA and DP Nodes

```

1 void highKnowledge(obs) {
2     rand = generateRandomNumber()
3     interestingAccuracy = 0.82
4     emptyAccuracy = 0.97
5     if(obs.actualContent == "INTERESTING") {
6         if R < interestingAccuracy
7             obs.id = "INTERESTING"
8         else
9             obs.id = "UNKNOWN"
10    } else{
11        if R < emptyAccuracy
12            obs.id = "EMPTY"
13        else
14            obs.id = "UNKNOWN"
15    }
16    if(this == centralNode)
17        obs.processingTime += 30
18    else
19        obs.processingTime += 43
20    setProcessing(false)
21 }
22 void lowKnowledge(obs) {
23     rand = generateRandomNumber()
24     interestingAccuracy = 0.1
25     if(obs.actualContent == "INTERESTING") {
26         if R < interestingAccuracy
27             obs.id = "INTERESTING"
28         else
29             obs.id = "UNKNOWN"
30     } else{
31         obs.id = "UNKNOWN"
32     }
33     obs.processingTime += 5
34     setProcessing(false)
35 }
```

Listing 7.2: Processing Functions for Nodes

7.4 Results

In this section, we show plots for each scenario and the variations on bandwidth and transmission rate. The resulting CSV files from each run were processed with the R programming language [91] and summary statistics were output to find the mean transmission time for all observations, interesting and empty, as well as those that were delivered to the end point during the runtime of the simulation.

Table 7.4 shows the mean number of observations delivered for all runs of each scenario and their variations. We can see that almost all observations are delivered when the network is not saturated and the transmission rate is ideal, despite Zigbee's low transmission rate. An observation is marked as undelivered if it has not yet reached the central node or, in the case of central processing scenarios, until it has been received and processed by the central node. As knowledge is pushed out into the network, the processing time becomes more of a factor (as processing nodes cannot process multiple observations at once) on top of the transmission time and we can see that this reduces the number of received observations. MK-HK-NK is the lowest here and we believe that is because it is the only scenario that processes an observation twice (with MK and then HK), but we discuss this in greater detail as we analyse the results further. As we discussed previously, processing on the central node means that an observation is only marked as delivered when it has been processed. However, the short processing times (when compared to long transmission times) and the fact that a DA node can process up to four observations in parallel means that this is not really a bottleneck.

Mode	Bandwidth	Transmission Rate	All	Delivered	Undelivered	% Delivered
NK-NK-MK	normal	ideal	44408	44255	153	99.7
NK-NK-MK	normal	variable	44539	35131	9407	79
NK-NK-MK	saturated	ideal	444714	129224	315490	29
NK-NK-MK	saturated	variable	444569	66454	378115	15
NK-NK-HK	normal	ideal	44481	44478	3	100
NK-NK-HK	normal	variable	44512	32804	11708	74
NK-NK-HK	saturated	ideal	444899	134552	310347	30
NK-NK-HK	saturated	variable	444760	57105	387655	13
NK-MK-NK	normal	ideal	44501	40614	3887	91
NK-MK-NK	normal	variable	44429	27981	16448	63
NK-MK-NK	saturated	ideal	444858	106656	338202	24
NK-NK-NK	saturated	variable	444649	50643	394006	11
MK-NK-NK	normal	ideal	44484	44206	278	99
MK-NK-NK	normal	variable	44533	26307	18226	59
MK-NK-NK	saturated	ideal	444594	92904	351690	21
MK-NK-NK	saturated	variable	444989	51225	393764	12
MK-HK-NK	normal	ideal	44513	40927	3586	92
MK-HK-NK	normal	variable	44455	23638	20817	53
MK-HK-NK	saturated	ideal	443211	79864	363347	18
MK-NK-NK	saturated	variable	444397	46454	397943	10
HK-NK-NK	normal	ideal	44423	43378	1045	98
HK-NK-NK	normal	variable	44437	29295	15142	66
HK-NK-NK	saturated	ideal	445257	97790	347467	22
HK-NK-NK	saturated	variable	444817	46061	398756	10

Table 7.3: All Observations Captured and Delivered

7.4.1 Normal Network

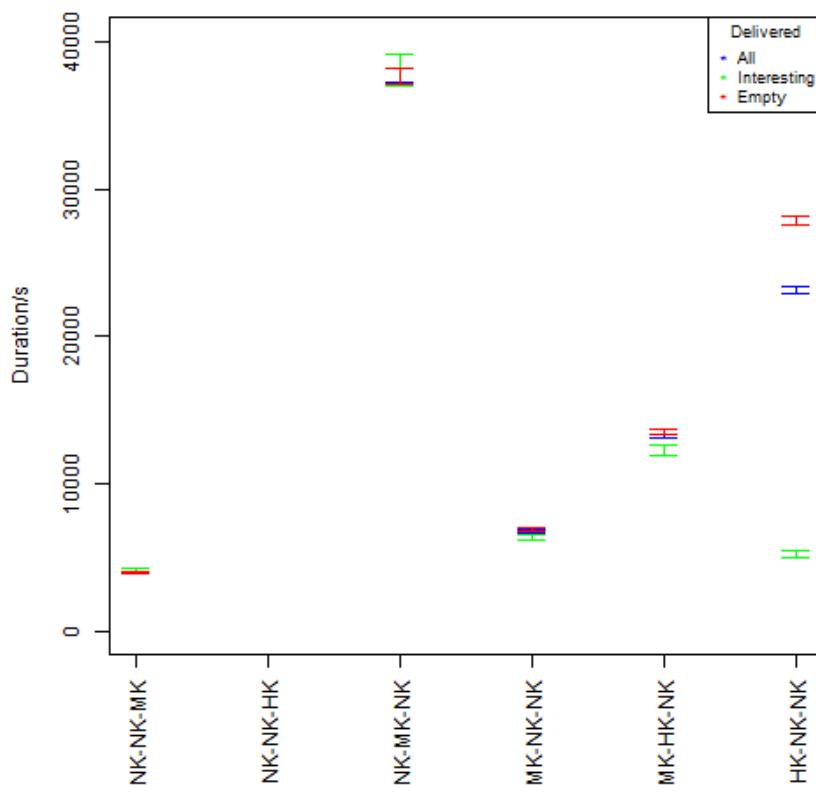
When the network is not saturated and Zigbee is able to transmit observations at its maximum transmission rate, we can see that the transmission time varies largely across all scenarios (Figure 7.2a). This plot is shown with error bars outlining the mean with a 95% confidence interval, this was done because the standard deviation showed such variation that the error bars ranged from a few hundred seconds to tens of thousands for each scenario. NK-NK-MK shows the lowest transmission time for all observations but it is only able to accurately identify 10% of interesting observations and cannot identify empty observations. NK-NK-HK is not visible in the figure because it has such a low mean transmission time (206 seconds), this appears to be an anomaly that is not apparent in other scenarios, but is consistent across all runs of normal bandwidth and ideal transmission rates. MK-HK-NK is higher than even MK-MK-MK because of its increased processing time, although MK-HK-NK is able to more accurately determine truly interesting observations. Whilst HK-NK-NK has a high mean transmission time for all observations, it is clear that the ability of this scenario to accurately prioritise from the edge significantly reduces the transmission time for interesting observations.

Figure 7.3a shows that all scenarios delivered the majority of the observations captured. MK-HK-NK delivered around 90% but this is likely due to the processing time discussed previously. The reduced number for NK-MK-NK was unexpected but we believe it is because the DP nodes become a bottleneck for processing as they can only process a single observation at a time. processing nodes are only able to process a single observation at a time and, when they are the only nodes with processing capabilities, the four nodes take on all observations from every sensing node. In this scenario, it could be better to pass on some of the processing to the central node since the link between them is a much faster Wi-Fi link, effectively making it NK-MK-MK. Because the network was not saturated, the delivered observations are consistent with the configuration parameters outlined in Table 7.3.1. When the transmission rate varies, we see a greater decrease in the number of delivered observations, show in Figure 7.3b.

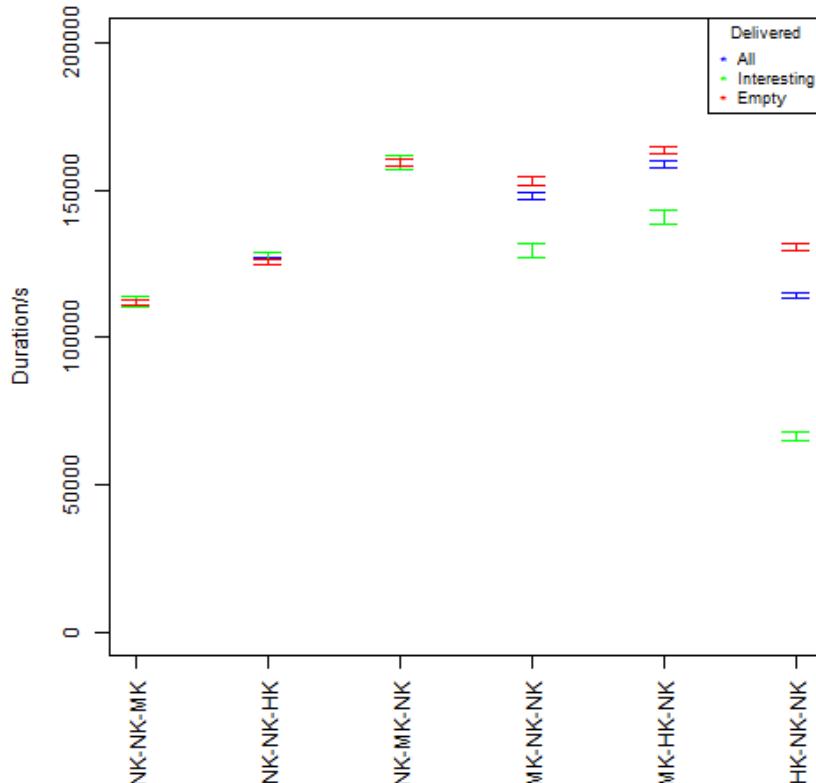
With the exception of HK-NK-NK, as knowledge processing capabilities are pushed out into the network, the fewer observations are delivered. This is consistent with Figure 7.2b as the transmission times also increase. HK-NK-NK is able to accurately prioritise from the edge, and every sensing node has processing capabilities, so it does fit that the mean transmission time would be reduced with this scenario, also allowing more observations to be delivered.

With a variable transmission rate, nodes could transmit anywhere from 20kbps to the maximum 250kbps that Zigbee provides, Figure 7.3b shows how this affects the number of delivered observations, reducing the number of delivered observations as knowledge is pushed out. MK-HK-NK delivers the fewest, most likely because observations are processed by both sensing and processing nodes. HK-NK-NK, on the other hand, deviates from the trend and delivers a much higher proportion, this is again most likely due to the fact that, even though higher knowledge processing takes longer, the capabilities are on all sensing nodes so there is less chance of a bottleneck situation occurring.

Figures 7.4a and 7.4b both show the same split of interesting vs empty observations because almost all observations captured for each scenario are delivered when the network is not saturated, so the percentage split is equivalent to the chances of an interesting observation captured shown in Table 7.3.1.



(a) Ideal Transmission Rate



(b) Variable Transmission Rate

Figure 7.2: Mean Transmission Times for Delivered Observations for a Normal Network.

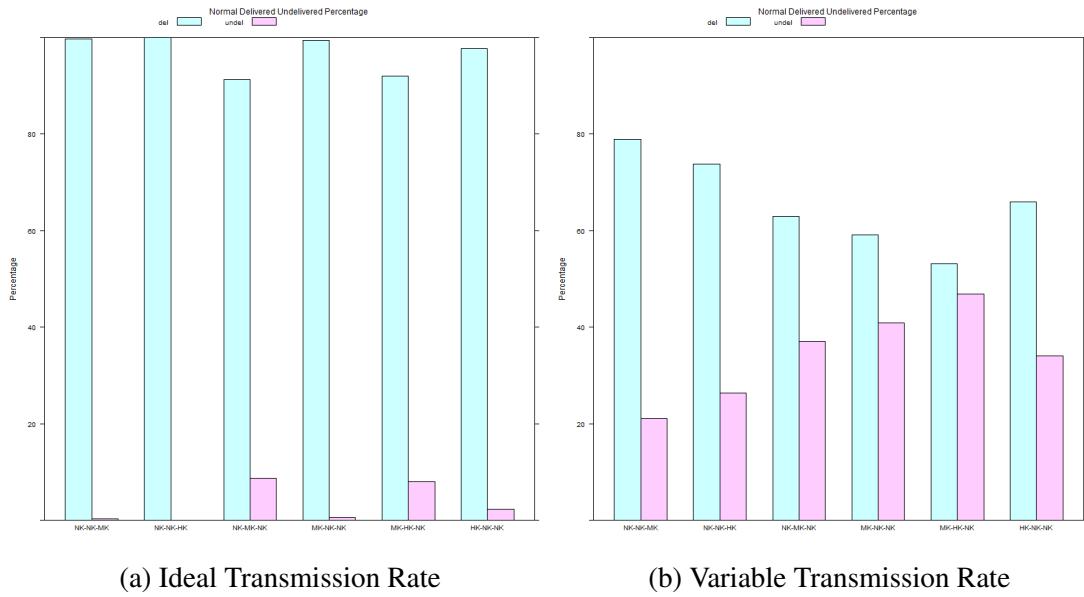


Figure 7.3: Delivered vs Undelivered Observations for a Normal Network

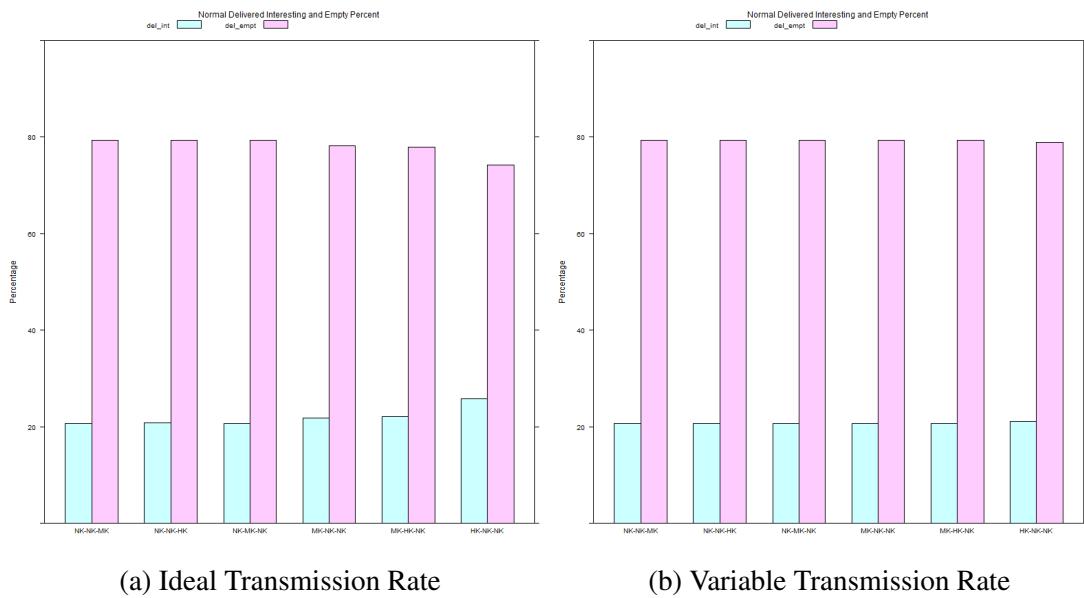


Figure 7.4: Interesting vs Empty Delivered Observations for a Normal Network

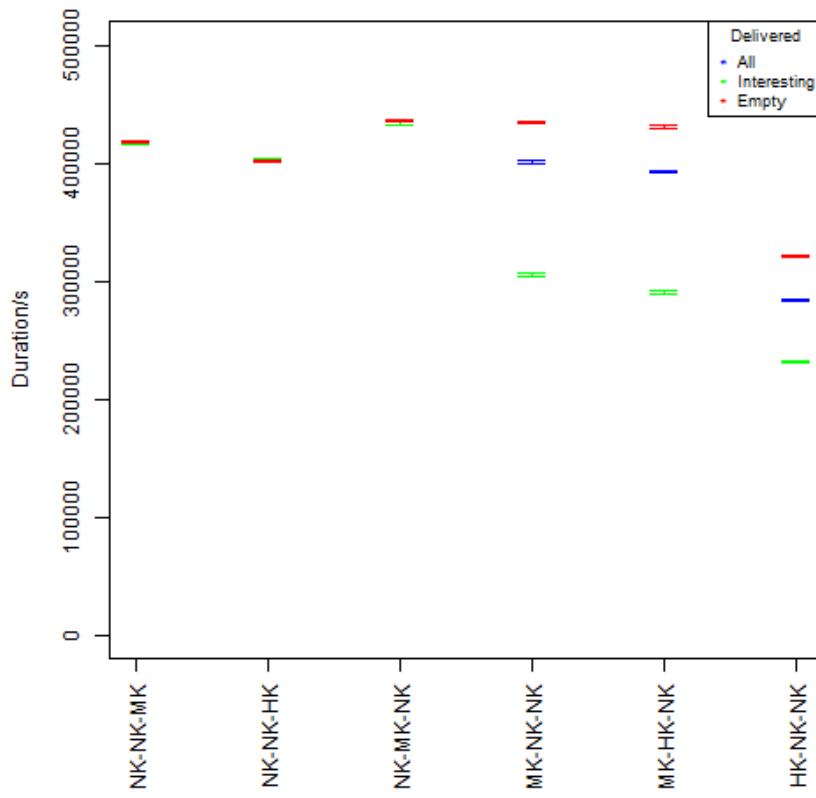
7.4.2 Saturated Network

The transmission time for both ideal and variable transmission rates is greatly increased when the network is saturated and we expected these scenarios to show how push-

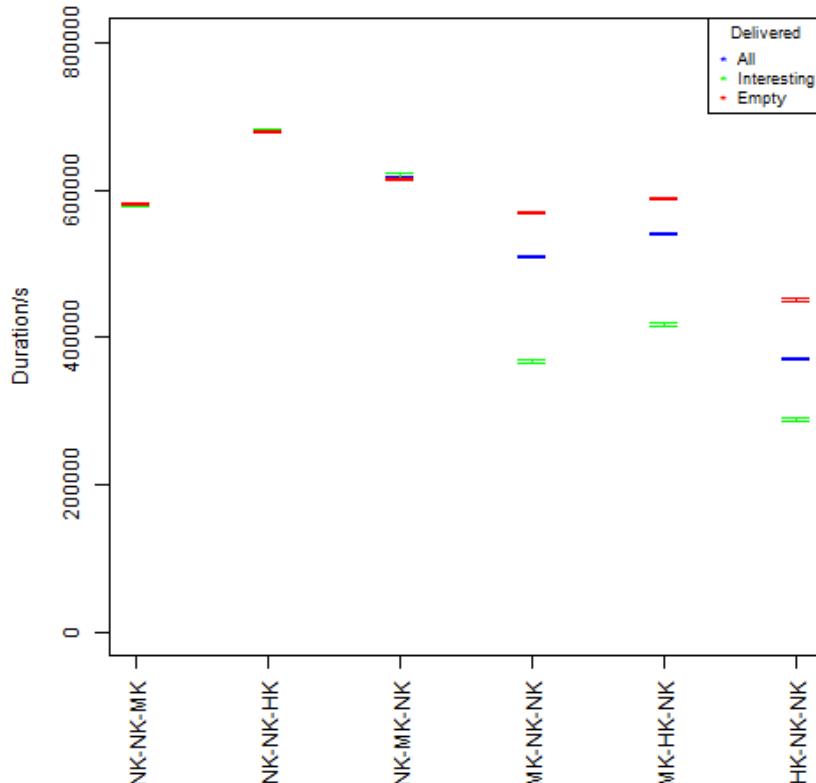
ing knowledge further into the network provides more interesting observations faster, prioritising when the network is aware that it cannot deliver every observation it has captured.

Despite variable transmission rates, Figures 7.5a and 7.5b show a similar trend. Scenarios of central processing gave consistently high transmission times for all observations, as is NK-MK-NK because nodes cannot prioritise data from the edge. HK-NK-NK is much faster when delivering interesting observations, which would also reduce the mean transmission time for all observations but we did not expect to see a reduction in the time taken for empty observations when compared with other scenarios. This is most likely because the processing time of HK-NK-NK means that fewer observations are delivered in each run (Figure 7.6a) but could also be because of empty observations misclassified as interesting being prioritised through the network as it is the only scenario that can mark an observation as empty from the edge. However, the higher accuracy of HK-NK-NK should mean that fewer misclassifications are made. As we expected, Figure 7.6b shows that, when the transmission rate is variable, fewer observations are delivered as they take longer to transmit.

Figures 7.7a and 7.7b show the number of interesting observations compared to the number of empty delivered when the network is saturated. Both show that, as knowledge is pushed further out into the network, the percentage of interesting observations delivered to the central node increases when there is not enough bandwidth available to deliver every observation captured.



(a) Ideal Transmission Rate



(b) Variable Transmission Rate

Figure 7.5: Mean Transmission Times for Delivered Observations for a Saturated Network.

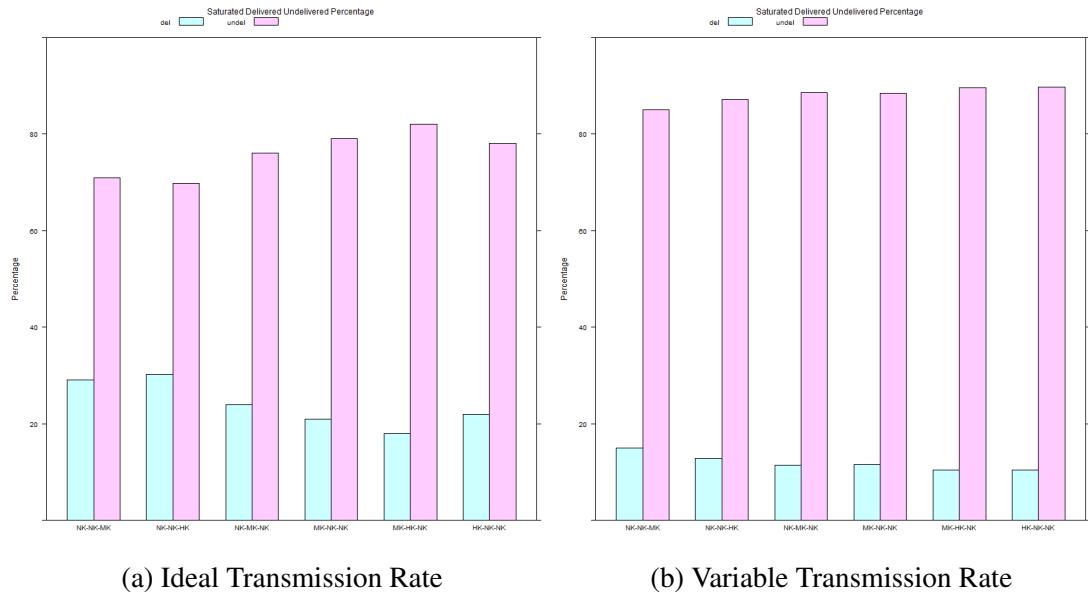


Figure 7.6: Delivered vs Undelivered Observations for a Normal Network

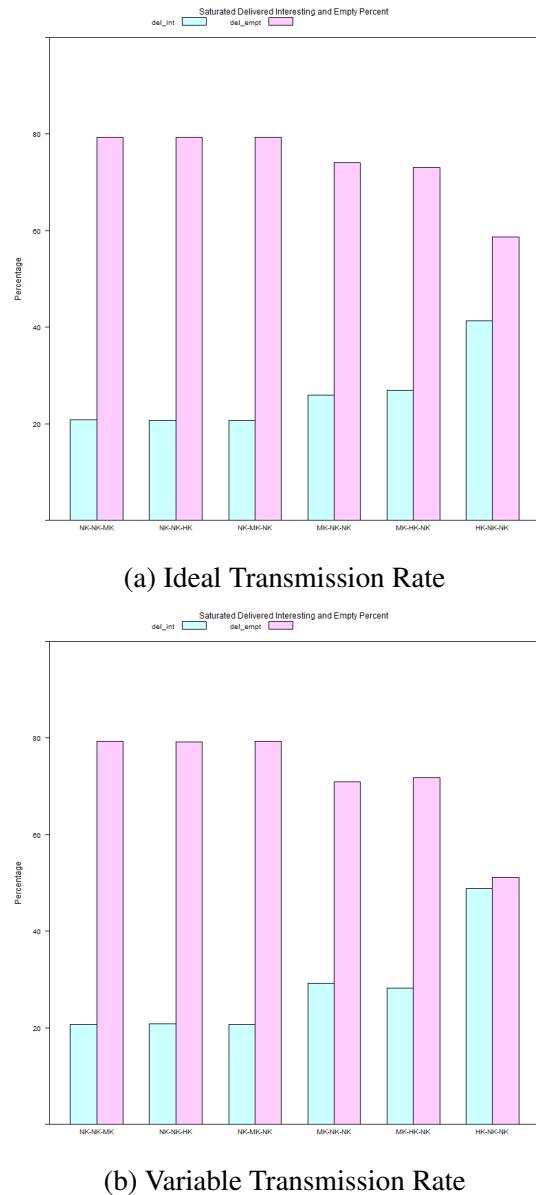


Figure 7.7: Interesting vs Empty Delivered Observations for a Normal Network

7.5 Conclusion

The general trend that is clear from all the figures is that pushing knowledge further out into the network increases the mean transmission time of all observations, when compared to central, but it does reduce the mean transmission time of the delivery of

interesting observations, especially when the network is saturated.

These simulations first show that saturating the network simply by increasing the chance of capture ten fold can have a huge impact on the transmission time of observations and really tests how each scenario prioritises data within the network. While processing each observation can take time depending on whether the node possesses minimal or high knowledge processing capabilities, it does not seem to be the bottleneck here. With nodes only having a single radio, sending in parallel is not possible and waiting for a neighbouring node to be available can delay an observation massively.

Pushing knowledge further out into the network does reduce the time taken for an interesting observation to reach the central node. When the network is saturated, it is clear to see how the prioritisation can impact what data is delivered and that higher knowledge processing capabilities allows for interesting data to be delivered faster and with a higher priority. The downside to MK-HK-NK is that this processing needs to occur twice: once at the data collection node and again on the data processing node and it seems that this adds a delay to observations. However, in many situations, it may not be possible to put high knowledge processing nodes throughout the network as they are more expensive.

The figures for both scenarios show that each scenario has benefits that depend on the purpose of the deployed sensor network. If a network was deployed for quick delivery of all sensed data to the DA node then central processing is the best scenario to choose. However, saturation will slow delivery and there is no prioritisation of data. If you wanted quick delivery of the most interesting data first then this depends on the accuracy you require. Using MK-NK-NK will give you quick delivery of interesting data, but the accuracy will be low and much of the data received could be incorrectly classified. Deploying MK-HK-NK does slow the delivery of data, due to the processing on both levels, but the network is able to use higher knowledge processing capabilities to improve the accuracy of classifications, as well as determine whether an observation is empty or not through the DP nodes. HK-NK-NK is the best option here as it allows

for high accuracy and prioritisation from the edge of the network.

Chapter 8

Conclusion and Future Work

In this thesis, we aimed to show that utilising the local knowledge of an environment in a WSN improves the efficiency of the network by giving it the ability to prioritise sensed data based on the results of in-network processing. We believe that pushing knowledge out farther towards the edges of the network improves the overall performance. To show this, we have developed a three-tier WSN architecture that uses knowledge-processing capabilities to process sensed data as it is forwarded through the network. Most current WSN implementations deliver data chronologically, or store it on the node to be retrieved by queries. We believe that this knowledge can be used to infer how valuable sensed data is and prioritise that through the network, delivering the most interesting data first. However, resources are still limited in WSNs and our architecture had to utilise these resources effectively, such as battery life and bandwidth. Using Data Collection (DC) nodes at the edge of the network, they capture observations and use their limited knowledge-processing capabilities to enrich the sensed data before sending it on. Data Processing (DP) nodes use more powerful knowledge-processing features to attempt to classify observations and prioritise the sending of them to Data Aggregation (DA) nodes. DA nodes make the data available to users and use their classifications, and input, to dynamically update its knowledge base.

We used a collaboration with a research centre in Malaysia with the aim to implement K-HAS in the Malaysian rainforest, in an area that had experienced logging and contained a diverse range of rare wildlife. Using K-HAS, we wanted to deploy a network that would prioritise images of rare wildlife and only send images of common wildlife

when bandwidth was available. Over the course of 3 visits, we gathered knowledge from the area, researchers and locals to build a knowledge base and create rules that could be used to classify data. We also collected images taken from their current, manual solution to infer patterns and use existing classifications for new sensed data.

Current sensing technology means that K-HAS is not ready to be implemented as the architecture dictates. DC nodes are not yet readily available, and this is especially true for image capturing sensor nodes. Because of this, we modified the architecture to use commercial hardware with fewer capabilities that would allow us to actually deploy a sensor network in the Malaysian rainforest that does use local knowledge. We call this architecture LORIS, Local-knowledge Ontology-based Remote-Sensing Informatics System, and this solution combines the DA and DP node to hold all of the knowledge-processing capabilities.

LORIS has shown that even using local knowledge at the base station of a WSN means that sensed data is processed and organised within minutes of being received. This method also means that users are alerted to data that they have subscribed automatically.

Our simulations of K-HAS show that the bandwidth in a WSN is used more effectively when knowledge is pushed out towards the edge of the network, allowing nodes to perform knowledge-processing to make inferences about the contents of the data and prioritise, or delay, its delivery appropriately. We have no reason to think that this would not work in practice.

8.1 Summary of Contributions

In this section, we summarise the contributions detailed in this thesis, focussing on our deployment in Malaysia, the tiered architecture designed for general use and the alignment ontology created to formalise the architecture of K-HAS and the data standard that it uses.

8.1.1 K-HAS

In Chapter 4, we present a novel tiered architecture, K-HAS, for WSNs that uses local knowledge. We explored existing networks, those related to our motivating scenario in Malaysia and also those that use knowledge or context-awareness, and routing protocols that use the sensed data to determine how it is routed. We explain the purpose of each tier in the network and show how sensed data is enriched and routed as it progresses through the network. Using Darwin Core as a data standard, each node communicates in a common format and metadata is packaged with data in archives that can be read by any node in the network. We also show how rules can be used to infer the content of sensed data and the ability to run rule engines in the network allows sensed data to be prioritised based on its value; not just the time it was captured.

8.1.2 Ontology

In Chapter 5 we explain the aligning ontology created to formally represent the K-HAS architecture and the data standard used. We show how, while there are existing ontologies that join observation-centric and sensor-centric ontologies, K-HAS allows for the representation of knowledge exchange in the network and show nodes performing tasks similar to humans. We also show how the ontology is extensible and does not need to be specific to K-HAS; it can be used with any WSN that deals with scientific observations.

8.1.3 LORIS

In Chapter 6, we present the Local-knowledge Ontology-based Remote-Sensing Informatics System (LORIS), a system developed for our motivating scenario when we discovered that current sensing technology means that K-HAS is not ready to be implemented in its current form. DC nodes are not a type of sensor that is readily available,

and this is especially true for image capturing sensor nodes. Because of this, we modified the architecture to use commercial hardware with fewer capabilities that would allow us to deploy a sensor network in the Malaysian rainforest that does use local knowledge. This solution combines the DA and DP node to hold all of the knowledge-processing capabilities and then uses commercially available sensors to replace DC nodes.

This architecture is easier to implement but does not provide in-network processing, although it automates the delivery of sensed data and uses the increased knowledge-processing capabilities of modern PCs to process sensed data on arrival and inform users. We show how our deployment of LORIS was successful and highlighted how sensed data was delivered within minutes of being captured in some cases, and processed shortly after.

8.1.4 Simulations

In Chapter 7, we explain the implementation and results of our simulations to model an ideal deployment of K-HAS. We model every variable of the network on existing data collected from our motivating scenario and show that the delivery of interesting observations can be effectively halved from almost 120 hours (for central and no processing) down to 64 hours for HK-HK-NK or 80 hours for K-HAS (see Figure 7.5a). We also outline how the network is able to prioritise data that it believes to be interesting, using a priority queue mechanism that delays data it believes to be empty.

8.2 Future Work

The focus of this thesis is to show that local knowledge can improve the timeliness of interesting data by making inferences based on previously sensed data and knowledge of the environment. We have shown, using simulations, that this can be done with

our tiered approach. However, another aim was to deploy such an architecture for our collaborators. We have explained that current technology means that this is not feasible, but sensing, and microcomputer technology has moved forward significantly since the beginning of this PhD and it would be possible with a longer time period. Our deployment in Malaysia was limited due to time constraints and we would have liked to leave a functional deployment active in Danau Girang for a six month period.

With more time, we could create custom DC nodes using webcams and micro-computers, like the Raspberry Pi, encased in a watertight enclosure; allowing us to use higher levels of knowledge-processing at the edge of the network. A deployment of this length would also allow nodes to act on sensed data classified by users and update their knowledge base, responding to changes in the network dynamically.

K-HAS uses a combination of open source projects and some software created during the course of this work, but it does require specialist knowledge to be deployed. We would like to create an installation candidate that would be usable by those without any expertise that could provide basic information on the purpose of the network and the installation of all necessary packages would be automated.

On top of this, our software's user interfaces have been tested by researchers at Danau Girang but we have no metrics on the usability of the software. Testing the software on users to determine how they would score the different areas of the software, such as usability, response time and learning curve, would help us to improve the software and ensure that it can be used by those without any technical knowledge of K-HAS or its underlying architecture.

Experiments have shown us that the range of wireless transmissions is often hard to predict and can be heavily influenced by weather and obstacles. If sensors were placed at the edge of range for a neighbouring sensor, then it is not guaranteed that transmissions could be made every time. We believe that using humans as an intermediate hop could be an interesting research opportunity. Using a mobile app to transfer data between nodes as they are within walking range could speed up data transfers and their

knowledge could be used to prioritise data without the need to process it.

If a human views a recent observation when they are in range of a node, they could make an instant assessment on whether it is interesting or not. If they mark it as such, then it could be passed through the network with a higher priority and reach a DA node within a short period as it would not require processing, updating knowledge bases for future, similar observations as it is forwarded.

One important goal is to create a deployment of K-HAS for a different purpose than our motivating scenario in Malaysia. There is a need in Malaysia to track hunters in the forest and alert authorities, this would not require any change in how K-HAS is currently implemented but would show how it can dynamically adapt based to changes in its sensing requirements. We would also like to test K-HAS in a situation where power is no longer a limiting factor, but delivery time of sensed data would be. A building security network would be one such example, deploying K-HAS across a number of floors and processing video feeds to alert users within seconds about suspicious activity.

However, using K-HAS for a WSN that, for example, uses text based sensed data to monitor the temperature and lava level of a volcano is a completely different implementation, but the local knowledge could be used to prevent an emergency and predict eruptions. A deployment such as this would show that the benefit of using local knowledge in a WSN is not limited to our motivating scenario, but is versatile enough to benefit almost any network.

Using heterogeneous sensor nodes within a K-HAS network would also show how local knowledge can be used in different streams, as well as how they can be combined to make more detailed inferences. Using motion sensors with microphones could be used to determine what person/animal is near the sensor and this information can then be used to infer patterns as fallback sources when an image based classification cannot be made.

Our simulations in Chapter 7 show how K-HAS can use local knowledge, but it is by no means a complete implementation. We need to implement a more modular simulation that allows K-HAS to be applied to any form of WSN. We would like to simulate individual knowledge bases on every node and experiment with networks that contain multiple central nodes to visualise the flow of sensed data through the network and see if there is an effect, positive or negative, on the speed of interesting sensed data. A more immediate goal is to test K-HAS when all of the network uses a communication medium with lower range but a higher transfer rate, such as Wi-Fi, to determine how much Zigbee slows the delivery of sensed data when compared with processing.

We would also like to experiment with different ratios of knowledge processing capabilities on nodes to determine if there is an ideal ratio that maximises the flow of sensed data, delivering interesting data first and quickly but also delivering data that has been classified as not interesting within a shorter time period that would allow humans to act on the data if it had been misclassified.

Bibliography

- [1] Reconyx. URL <http://www.reconyx.com/>.
- [2] Danau Girang Field Centre. URL <http://www.cardiff.ac.uk/biosi/facilities/danaugirangfieldcentre/>.
- [3] Dweet.io. URL <http://dweet.io/>.
- [4] Fitbit. URL <http://www.fitbit.com/uk>.
- [5] Global biodiversity information facility. URL <http://www.gbif.org/>.
- [6] If This Then That. URL <https://ifttt.com/myrecipes/personal>.
- [7] Texas Instruments Pandaboard. URL <http://pandaboard.org/>.
- [8] Nest. URL <https://nest.com/uk/>.
- [9] Sen.se. URL <http://open.sen.se/>.
- [10] Xively. URL <https://xively.com>.
- [11] Karl Aberer, Manfred Hauswirth, and Ali Salehi. The Global Sensor Networks middleware for efficient and flexible deployment and interconnection of sensor networks. Technical report, 2006.
- [12] Karl Aberer, Manfred Hauswirth, and Ali Salehi. A middleware for fast and flexible sensor network deployment. In *Proceedings of the 32nd International*

- Conference on Very large data bases*, pages 1199–1202. VLDB Endowment, 2006.
- [13] Jorge A Ahumada, Carlos EF Silva, Krisna Gajapersad, Chris Hallam, Johanna Hurtado, Emanuel Martin, Alex McWilliam, Badru Mugerwa, Tim O’Brien, and Francesco Rovero. Community structure and diversity of tropical forest mammals: data from a global camera trap network. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 366(1578):2703–2711, 2011.
 - [14] Kemal Akkaya and Mohamed Younis. A survey on routing protocols for wireless sensor networks. *Ad hoc networks*, 3(3):325–349, 2005.
 - [15] Ian F Akyildiz, Weilian Su, Yogesh Sankarasubramaniam, and Erdal Cayirci. Wireless Sensor Networks: A Survey. *Computer networks*, 38(4):393–422, 2002.
 - [16] Ian F Akyildiz, Weilian Su, Yogesh Sankarasubramaniam, and Erdal Cayirci. A survey on sensor networks. *IEEE Communications magazine*, 40(8):102–114, 2002.
 - [17] Laurentius Nayan Ambu and Benoit Goossens. Sabah Wildlife Department and 10 years of research: Towards a better conservation of Sabah’s wildlife. *Journal of Oil Palm, Environment and Health (JOPEH)*, 3, 2012.
 - [18] Steven J Baskauf. Darwin-SW Examples, 2011. URL <https://code.google.com/p/darwin-sw/wiki/Examples>.
 - [19] Steven J Baskauf and Campbell O Webb. Darwin-SW: Darwin Core-based terms for expressing biodiversity data as RDF. *Semantic Web Journal*, 2014.
 - [20] Daryoush Bayat, Daryoush Habibi, and Iftikhar Ahmad. Development of a Wireless Sensor Node for Environmental Monitoring. In *SENSORCOMM 2012, The Sixth International Conference on Sensor Technologies and Applications*, pages 1–5, 2012.

- [21] G. Bradski. OpenCV library. *Dr. Dobb's Journal of Software Tools*, 2000.
- [22] Janez Brank, Marko Grobelnik, and Dunja Mladenić. A survey of ontology evaluation techniques. In *Proceedings of the Conference on Data Mining and Data Warehouses*, 2005.
- [23] Dan Brickley and Libby Miller. FOAF Vocabulary Specification 0.99. Namespace document, 2014. URL <http://xmlns.com/foaf/spec/>.
- [24] Arne Bröring, Johannes Echterhoff, Simon Jirka, Ingo Simonis, Thomas Everding, Christoph Stasch, Steve Liang, and Rob Lemmens. New generation sensor web enablement. *Sensors*, 11(3):2652–2699, 2011.
- [25] Buckeye. Buckeye Camera. URL <http://www.buckeyecam.com/site/products.html>.
- [26] Benjie Chen, Kyle Jamieson, Hari Balakrishnan, and Robert Morris. Span: An energy-efficient coordination algorithm for topology maintenance in ad hoc wireless networks. *Wireless networks*, 8(5):481–494, 2002.
- [27] JooYoung Choi, Melissa J Davis, Andrew F Newman, and Mark A Ragan. A semantic web ontology for small molecules and their biological targets. *Journal of chemical information and modeling*, 50(5):732–741, 2010.
- [28] Christopher Brewster and Harith Alani and An Dasmahapatra. Data driven ontology evaluation. In *Int. Conf. on Language Resources and Evaluation*, 2004.
- [29] Nick Collier. Repast: An extensible framework for agent simulation. Technical report, Social Science Research Computing, University of Chicago, Chicago, Illinois, 2000.
- [30] Michael Compton, Cory Henson, Laurent Lefort, Holger Neuhaus, and Amit Sheth. A survey of the semantic specification of sensors. In *2nd International Semantic Sensor Networks Workshop*, 2009.

- [31] Michael Compton, Payam Barnaghi, Luis Bermudez, RaúL GarcíA-Castro, Oscar Corcho, Simon Cox, John Graybeal, Manfred Hauswirth, Cory Henson, and Arthur Herzog. The SSN ontology of the W3C semantic sensor network incubator group. *Web Semantics: Science, Services and Agents on the World Wide Web*, 17:25–32, 2012.
- [32] P. Corke, T. Wark, R. Jurdak, Wen Hu, P. Valencia, and D. Moore. Environmental wireless sensor networks. *Proceedings of the IEEE*, 98(11):1903–1917, Nov 2010. ISSN 0018-9219. doi: 10.1109/JPROC.2010.2068530.
- [33] V Dedoose. Web Application for Managing, Analyzing, and Presenting Qualitative and Mixed Method Data. *Los Angeles, CA Sociocult. Res. Consult. LLC*, 2012.
- [34] Murat Demirbas, Ken Yian Chow, and Chieh Shyan Wan. INSIGHT: Internet-sensor integration for habitat monitoring. In *Proceedings of the International Symposium on on World of Wireless, Mobile and Multimedia Networks*, pages 553–558. IEEE Computer Society, 2006.
- [35] Billie R DeWalt. Using indigenous knowledge to improve agriculture and natural resource management. *Human organization*, 53(2):123–131, 1994.
- [36] Sanjeev Dhawan. Analogy of promising wireless technologies on different frequencies: Bluetooth, wifi, and wimax. In *The 2nd International Conference on Wireless Broadband and Ultra Wideband Communications, 2007. AusWireless 2007*, pages 14–14. IEEE, 2007.
- [37] Mohamad Eid, Ramiro Liscano, and Abdulmotaleb El Saddik. A universal ontology for sensor networks data. In *International Conference on Computational Intelligence for Measurement Systems and Applications (CISMA)*, pages 59–62. IEEE, 2007.
- [38] Deborah Estrin, Lewis Girod, Greg Pottie, and Mani Srivastava. Instrumenting the world with wireless sensor networks. In *Proceedings of International Con-*

- ference on Acoustics, Speech, and Signal Processing (ICASSP)*, volume 4, pages 2033–2036. IEEE, 2001.
- [39] Mariano Fernández-López. Overview of methodologies for building ontologies. In *Proceedings of IJCAI-99, Workshop KRR5, Sweden*. CEUR Publications (1999), 1999.
- [40] Mariano Fernández-López, Asunción Gómez-Pérez, and Natalia Juristo. Methontology: from ontological art towards ontological engineering. *AAAI Symposium on Ontological Engineering, Stanford*, 1997.
- [41] Carlos MS Figueiredo, Eduardo F Nakamura, Afonso D Ribas, Thales RB de Souza, and Raimundo S Barreto. Assessing the communication performance of wireless sensor networks in rainforests. In *Wireless Days (WD), 2nd IFIP*, pages 1–6. IEEE, 2009.
- [42] Antonio-Javier Garcia-Sánchez, Felipe Garcia-Sánchez, Fernando Losilla, Paweł Kulakowski, Joan García-Haro, Alejandro Rodríguez, José-Vicente López-Bao, and Francisco Palomares. Wireless sensor network deployment for monitoring wildlife passages. *Sensors*, 10(8):7236–7262, 2010.
- [43] Gartner. Gartner’s 2013 Hype Cycle for Emerging Technologies, 2013. URL <http://www.gartner.com/newsroom/id/2575515>.
- [44] Daniel Giusto, Antonio Lera, Giacomo Morabito, and Luigi Atzori. *The Internet of Things*. Springer, 2010.
- [45] Carles Gomez, Joaquim Oller, and Josep Paradells. Overview and evaluation of bluetooth low energy: An emerging low-power wireless technology. *Sensors*, 12(9):11734–11753, 2012.
- [46] Michael Grüninger and Mark S Fox. Methodology for the Design and Evaluation of Ontologies. In *Proceedings of the Workshop on Basic Ontological Issues in Knowledge Sharing, IJCAI-95*. Citeseer, 1995.

- [47] Ying Guo, Peter Corke, Geoff Poulton, Tim Wark, Greg Bishop-Hurley, and Dave Swain. Animal behaviour understanding using wireless sensor networks. In *Proceedings of 31st Conference on Local Computer Networks*, pages 607–614. IEEE, 2006.
- [48] Gaurav Gupta and Mohamed Younis. Performance evaluation of load-balanced clustering of wireless sensor networks. In *10th International Conference on Telecommunications (ICT)*, volume 2, pages 1577–1583. IEEE, 2003.
- [49] Chris Gwilliams, Alun Preece, Alex Hardisty, Benoit Goossens, and Laurentius N Ambu. Local and Global Knowledge to Improve the Quality of Sensed Data. *International Journal of Digital Information and Wireless Communications (IJDIWC)*, 2(2):164–180, 2012.
- [50] Christopher Gwilliams. Triton. <https://github.com/encima/triton>, 2013.
- [51] Christopher Gwilliams, Alun Preece, Alex Hardisty, Benoit Goossens, and Laurentius N Ambu. K-HAS: An Architecture for Using Local and Global Knowledge in Wireless Sensor Networks. In *The International Conference on Informatics and Applications (ICIA2012)*, pages 331–341. The Society of Digital Information and Wireless Communication, 2012.
- [52] Christopher Gwilliams, Alun D Preece, and Alex R Hardisty. Poster: using local and global knowledge in wireless sensor networks. In *Proceedings of the 10th International Conference on Mobile systems, applications, and services*, pages 515–516. ACM, 2012.
- [53] Salem Hadim and Nader Mohamed. Middleware: Middleware challenges and approaches for wireless sensor networks. *IEEE distributed systems online*, 7(3):1, 2006.
- [54] Mark Hall, Eibe Frank, Geoffrey Holmes, Bernhard Pfahringer, Peter Reutemann, and Ian H Witten. The weka data mining software: an update. *ACM SIGKDD explorations newsletter*, 11(1):10–18, 2009.

- [55] Wendi B. Heinzelman, Amy L. Murphy, Hervaldo S. Carvalho, and Mark A. Perillo. Middleware to support sensor network applications. *IEEE Network*, 18(1):6–14, 2004.
- [56] Wendi Rabiner Heinzelman, Joanna Kulik, and Hari Balakrishnan. Adaptive protocols for information dissemination in wireless sensor networks. In *Proceedings of the 5th annual ACM/IEEE International Conference on Mobile computing and networking*, pages 174–185. ACM, 1999.
- [57] Mark Hempstead, Nikhil Tripathi, Patrick Mauro, Gu-Yeon Wei, and David Brooks. An ultra low power system architecture for sensor network applications. In *ACM SIGARCH Computer Architecture News*, volume 33, pages 208–219. IEEE Computer Society, 2005.
- [58] Chalermek Intanagonwiwat, Ramesh Govindan, and Deborah Estrin. Directed diffusion: a scalable and robust communication paradigm for sensor networks. In *Proceedings of the 6th annual International Conference on Mobile computing and networking*, pages 56–67. ACM, 2000.
- [59] Krzysztof Janowicz and Michael Compton. The Stimulus-Sensor-Observation Ontology Design Pattern and its Integration into the Semantic Sensor Network Ontology. In *Proceedings of the 3rd International Workshop on Semantic Sensor Networks (SSN)*, 2010. URL <http://ceur-ws.org/Vol-668/paper12.pdf>.
- [60] Ernesto Jiménez-Ruiz, Bernardo Cuenca Grau, Ulrike Sattler, Thomas Schneider, and Rafael Berlanga. *Safe and economic re-use of ontologies: A logic-based methodology and tool support*. Springer, 2008.
- [61] L Joshi, M van Noordwijk, and F Sinclair. Bringing local knowledge into perspective – a case of sustainable technology development in jungle rubber agro forests in jambi. In *Participatory Technology Development and Local Knowledge for Sustainable Land Use in Southeast Asia*, 2001.

- [62] Joseph M Kahn, Randy H Katz, and Kristofer SJ Pister. Next century challenges: mobile networking for ‘Smart Dust’. In *Proceedings of the 5th annual ACM/IEEE International Conference on Mobile computing and networking*, pages 271–278. ACM, 1999.
- [63] Roland Kays, Bart Kranstauber, Patrick Jansen, Chris Carbone, Marcus Rowcliffe, Tony Fountain, and Sameer Tilak. Camera traps as sensor networks for monitoring animal communities. In *34th Conference on Local Computer Networks (LCN)*, pages 811–818. IEEE, 2009.
- [64] Ron Kohavi. The power of decision tables. In *8th European Conference on Machine Learning*, pages 174–189. Springer, 1995.
- [65] Jin-Shyan Lee, Yu-Wei Su, and Chung-Chou Shen. A comparative study of wireless protocols: Bluetooth, UWB, ZigBee, and Wi-Fi. In *33rd Annual Conference of the IEEE Industrial Electronics Society (IECON)*, pages 46–51. IEEE, 2007.
- [66] Philip Levis, Sam Madden, Joseph Polastre, Robert Szewczyk, Kamin Whitehouse, Alec Woo, David Gay, Jason Hill, Matt Welsh, and Eric Brewer. TinyOS: An operating system for sensor networks. In *Ambient intelligence*, pages 115–148. Springer, 2005.
- [67] John P Lewis. Fast template matching. In *Vision interface*, volume 95, pages 15–19, 1995.
- [68] Libelium. Wasp mote - Open Source Sensor Device | Libelium. URL <http://www.libelium.com/products/wasp mote/>.
- [69] Ting Liu and Margaret Martonosi. Impala: A middleware system for managing autonomic, parallel sensor systems. In *ACM SIGPLAN Notices*, volume 38, pages 107–118. ACM, 2003.

- [70] Sean Luke, Claudio Cioffi-Revilla, Liviu Panait, Keith Sullivan, and Gabriel Balan. Mason: A multiagent simulation environment. *Simulation*, 81(7):517–527, July 2005. ISSN 0037-5497. doi: 10.1177/0037549705058073. URL <http://dx.doi.org/10.1177/0037549705058073>.
- [71] Joshua Madin, Shawn Bowers, Mark Schildhauer, Sergeui Krivov, Deana Pennington, and Ferdinando Villa. An ontology for describing and synthesizing ecological observation data. *Ecological informatics*, 2(3):279–296, 2007.
- [72] Joshua Madin, Shawn Bowers, Mark Schildhauer, Sergeui Krivov, Deana Pennington, and Ferdinando Villa. An ontology for describing and synthesizing ecological observation data. *Ecological informatics*, 2(3):279–296, 2007.
- [73] Alan Mainwaring, David Culler, Joseph Polastre, Robert Szewczyk, and John Anderson. Wireless sensor networks for habitat monitoring. In *Proceedings of the 1st ACM international workshop on Wireless sensor networks and applications*, pages 88–97. ACM, 2002.
- [74] Arati Manjeshwar and Dharma P Agrawal. TEEN: a routing protocol for enhanced efficiency in wireless sensor networks. In *Parallel and Distributed Processing Symposium, International*, volume 3, pages 30189a–30189a. IEEE Computer Society, 2001.
- [75] Andrew Markham, Niki Trigoni, and Stephen Ellwood. Effect of rainfall on link quality in an outdoor forest deployment. In *Proceedings of the International Conference on Wireless Information Networks and Systems (WINSYS)*, pages 1–6. IEEE, 2010.
- [76] Kirk Martinez, Royan Ong, and Jane Hart. Glacsweb: a sensor network for hostile environments. In *1st Annual Communications Society Conference on Sensor and Ad Hoc Communications and Networks (SECON)*, pages 81–87. IEEE, 2004.

- [77] Uwe Maurer, Anthony Rowe, Asim Smailagic, and Daniel P Siewiorek. eWatch: a wearable sensor and notification platform. In *International Workshop on Wearable and Implantable Body Sensor Networks (BSN)*, pages 4–pp. IEEE, 2006.
- [78] Steven McCanne, Sally Floyd, Kevin Fall, Kannan Varadhan, and Others. Network simulator ns-2, 1997. URL <http://www.isi.edu/nsnam/ns>.
- [79] Deborah L McGuinness and Frank van Harmelen. Web ontology language (owl) overview. *World Wide Web Consortium Recommendation*, 10:60–71, 2004.
- [80] Luc Moreau, Ben Clifford, Juliana Freire, Joe Futrelle, Yolanda Gil, Paul Groth, Natalia Kwasnikowska, Simon Miles, Paolo Missier, and Jim Myers. The open provenance model core specification (v1. 1). *Future Generation Computer Systems*, 27(6):743–756, 2011.
- [81] Lama Nachman, Jonathan Huang, Junaith Shahabdeen, Robert Adler, and Ralph Kling. Imote2: Serious computation at the edge. In *Wireless Communications and Mobile Computing Conference (IWCMC)*, pages 1118–1123. IEEE, 2008.
- [82] Cristóbal Carnero Li nán. cvBlob. <http://cvblob.googlecode.com>. URL <http://cvblob.googlecode.com>.
- [83] Ian Niles and Adam Pease. Towards a standard upper ontology. In *Proceedings of the international conference on Formal Ontology in Information Systems-Volume 2001*, pages 2–9. ACM, 2001.
- [84] Natalya F Noy, Michael Sintek, Stefan Decker, Monica Crubézy, Ray W Ferger-son, and Mark A Musen. Creating semantic web contents with protege-2000. *IEEE intelligent systems*, 16(2):60–71, 2001.
- [85] Chris Otto, Aleksandar Milenkovic, Corey Sanders, and Emil Jovanov. System architecture of a wireless body area sensor network for ubiquitous health monitoring. *Journal of Mobile Multimedia*, 1(4):307–326, 2006.

- [86] Diego Pizzocaro, Alun Preece, Fangfei Chen, Thomas La Porta, and Amotz Bar-Noy. A distributed architecture for heterogeneous multi sensor-task allocation. In *International Conference on Distributed Computing in Sensor Systems and Workshops (DCOSS)*, pages 1–8. IEEE, 2011.
- [87] María Poveda-Villalón, Mari Carmen Suárez-Figueroa, and Asunción Gómez-Pérez. Validating ontologies with oops! In *Knowledge Engineering and Knowledge Management*, pages 267–281. Springer, 2012.
- [88] Florian Probst, Andrew Gordon, and Ivan Dornelas. OGC Discussion Paper: Ontology-based Representation of the OGC Observations and Measurements Model. *Institute for Geoinformatics (ifgi)*, 2006.
- [89] Mark Proctor, M Neale, P Lin, and M Frandsen. Drools. URL <http://www.drools.org>.
- [90] J. Ross Quinlan. *C4.5: Programs for Machine Learning*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 1993. ISBN 1-55860-238-0.
- [91] R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2014. URL <http://www.R-project.org/>.
- [92] Vijay Raghunathan, Curt Schurges, Sung Park, and Mani B Srivastava. Energy-aware wireless microsensor networks. *Signal Processing Magazine, IEEE*, 19(2):40–50, 2002.
- [93] AK Rahman, MS Anuar, SA Aljunid, MN Junita, A Sahadah, MN Ismail, and BS Naimullah. Study of rain attenuation with diversity of wavelengths propagation transmission in tropical rainforest region. In *International Conference on Electronic Design (ICED)*, pages 1–4. IEEE, 2008.

- [94] Md Atiqur Rahman. Middleware for wireless sensor networks: Challenges and Approaches. In *Seminar on Internetworking, Helsinki University of Technology, Finland*, volume 124, 2009.
- [95] Reconyx. Reconyx PC900C Camera. URL http://www.reconyx.com/shop/PC900C_Cellular_HyperFire_Professional_HO_Covert_IR/d/358/56.
- [96] C. Reed, M. Botts, J. Davidson, and G. Percivall. OGC sensor web enablement: overview and high level architecture. In *Autotestcon*, pages 372–380, Sept 2007. doi: 10.1109/AUTEST.2007.4374243.
- [97] Bernd Resch, Thomas Blaschke, and Manfred Mittlboeck. Live geography: Interoperable geo-sensor webs facilitating the vision of digital earth. *International Journal on Advances in Networks and Services*, 3(3 and 4):323–332, 2011.
- [98] Alexandre Robin and Michael E Botts. Creation of Specific SensorML Process Models. *White Paper Earth System Science Center (NSSTC). University of Alabama in Huntsville (UAH)*, 2006.
- [99] Narayanan Sadagopan, Bhaskar Krishnamachari, and Ahmed Helmy. The ACQUIRE mechanism for efficient querying in sensor networks. In *Proceedings of the 1st International Workshop on Sensor Network Protocols and Applications*, pages 149–155. IEEE, 2003.
- [100] Curt Schurgers and Mani B Srivastava. Energy efficient routing in wireless sensor networks. In *Military Communications Conference MILCOM. Communications for Network-Centric Operations: Creating the Information Force*, volume 1, pages 357–361. IEEE, 2001.
- [101] Michael Segal. Improving lifetime of wireless sensor networks. *Network Protocols and Algorithms*, 1(2):48–60, 2010.
- [102] Ladislaus M Semali. *What is indigenous knowledge?: Voices from the academy*. Routledge, 1999.

- [103] SONet. Darwin Core and OBOE Representation Examples. URL
[https://sonet.ecoinformatics.org/observational-data-use-cases/
obo-representation-examples](https://sonet.ecoinformatics.org/observational-data-use-cases/obo-representation-examples).
- [104] Joseph E Stiglitz. Peer monitoring and credit markets. *The world bank economic review*, 4(3):351–366, 1990.
- [105] Melanie Swan. Sensor mania! the internet of things, wearable computing, objective metrics, and the quantified self 2.0. *Journal of Sensor and Actuator Networks*, 1(3):217–253, 2012.
- [106] Robert Szewczyk, Alan Mainwaring, Joseph Polastre, John Anderson, and David Culler. An analysis of a large scale habitat monitoring application. In *Proceedings of the 2nd International Conference on Embedded networked sensor systems*, pages 214–226. ACM, 2004.
- [107] Robert Szewczyk, Eric Osterweil, Joseph Polastre, Michael Hamilton, Alan Mainwaring, and Deborah Estrin. Habitat monitoring with sensor networks. *Communications of the ACM*, 47(6):34–40, 2004.
- [108] Robert Szewczyk, Joseph Polastre, Alan Mainwaring, and David Culler. Lessons from a sensor network expedition. In *Wireless Sensor Networks*, pages 307–322. Springer, 2004.
- [109] Allan Tengg, Andreas Klausner, and Bernhard Rinner. I-SENSE: A light-weight middleware for embedded multi-sensor data-fusion. In *5th Workshop on Intelligent Solutions in Embedded Systems*, pages 165–177. IEEE, 2007.
- [110] Kirsten Terfloth, Georg Wittenburg, and Jochen H Schiller. FACTS-A rule-based middleware architecture for wireless sensor networks. In *Comsware*, 2006.

- [111] Mike Uschold. Building Ontologies: Towards a Unified Methodology. In *16th Annual Conf. of the British Computer Society Specialist Group on Expert Systems*, 1996.
- [112] Mike Uschold and Martin King. Towards a Methodology for Building Ontologies. In *Workshop on Basic Ontological Issues in Knowledge Sharing, held in conjunction with IJCAI-95*, 1995.
- [113] Bruno Valente and Francisco Martins. A middleware framework for the Internet of Things. In *3rd International Conference on Advances in Future Internet (AFIN)*, pages 139–144, 2011.
- [114] Vladimir Vujovic and Mirjana Maksimovic. Raspberry Pi as a Wireless Sensor node: Performances and constraints. In *37th International Convention on Information and Communication Technology, Electronics and Microelectronics (MIPRO)*, pages 1013–1018. IEEE, 2014.
- [115] Rory PD Walsh. Drought frequency changes in Sabah and adjacent parts of northern Borneo since the late nineteenth century and possible implications for tropical rain forest dynamics. *Journal of Tropical Ecology*, 12(03):385–407, 1996.
- [116] Miao-Miao Wang, Jian-Nong Cao, Jing Li, and Sajal K Dasi. Middleware for wireless sensor networks: A survey. *Journal of computer science and technology*, 23(3):305–326, 2008.
- [117] Tim Wark, Wen Hu, Peter Corke, Jonathan Hodge, Aila Keto, Ben Mackey, Glenn Foley, Pavan Sikka, and Michael Brunig. Springbrook: Challenges in developing a long-term, rainforest wireless sensor network. In *International Conference on Intelligent Sensors, Sensor Networks and Information Processing (ISSNIP)*, pages 599–604. IEEE, 2008.

- [118] Geoffrey Werner-Allen, Konrad Lorincz, Mario Ruiz, Omar Marcillo, Jeff Johnson, Jonathan Lees, and Matt Welsh. Deploying a wireless sensor network on an active volcano. *Internet Computing, IEEE*, 10(2):18–25, 2006.
- [119] John Wieczorek, David Bloom, Robert Guralnick, Stan Blum, Markus Döring, Renato Giovanni, Tim Robertson, and David Vieglais. Darwin Core: An evolving community-developed biodiversity data standard. *PLoS One*, 7(1):e29715, 2012.
- [120] Edward Woodrow and Wendi Heinzelman. SPIN-IT: a data centric routing protocol for image retrieval in wireless networks. In *International Conference on Image Processing. Proceedings*, volume 3, pages 913–916. IEEE, 2002.
- [121] Joel Wright, Christopher Gibson, Flavio Bergamaschi, Kelvin Marcus, Ryan Pressley, Gunjan Verma, and Gene Whipps. A dynamic infrastructure for interconnecting disparate isr/istar assets (the ita sensor fabric). In *12th International Conference on Information Fusion. FUSION'09*, pages 1393–1400. IEEE, 2009.
- [122] Xue Yang and NF Vaidya. A wakeup scheme for sensor networks: Achieving balance between energy saving and end-to-end delay. In *10th Real-Time and Embedded Technology and Applications Symposium (RTAS). Proceedings*, pages 19–26. IEEE, 2004.
- [123] Yong Yao and Johannes Gehrke. The cougar approach to in-network query processing in sensor networks. *ACM Sigmod Record*, 31(3):9–18, 2002.
- [124] Fan Ye, Alvin Chen, Songwu Lu, and Lixia Zhang. A scalable solution to minimum cost forwarding in large sensor networks. In *10th International Conference on Computer Communications and Networks. Proceedings*, pages 304–309. IEEE, 2001.
- [125] Yan Yu, Ramesh Govindan, and Deborah Estrin. Geographical and energy aware routing: A recursive data dissemination protocol for wireless sensor networks. Technical report, 2001.

- [126] Yang Yu, Bhaskar Krishnamachari, and Viktor K. Prasanna. Issues in designing middleware for wireless sensor networks. *IEEE Network*, 18:15–21, 2003.
- [127] Marco Zennaro, Antoine Bagula, David Gascon, and Alberto Bielsa Noveleta. Planning and deploying long distance wireless sensor networks: The integration of simulation and experimentation. In *Ad-Hoc, Mobile and Wireless Networks*, pages 191–204. Springer, 2010.
- [128] Zoran Zivkovic. Improved adaptive Gaussian mixture model for background subtraction. In *Proceedings of the 17th International Conference on Pattern Recognition (ICPR)*, volume 2, pages 28–31. IEEE, 2004.
- [129] ZSL. Instant Wild: Live Wildlife Photographs. URL <http://www.edgeofexistence.org/instantwild/>.

Appendices

Appendix A

Example Darwin Core Archive

```

1 <?xml version='1.0' encoding='utf-8'?>
2 <eml:eml xmlns:eml="eml://ecoinformatics.org/eml-2.1.1"
3   xmlns:md="eml://ecoinformatics.org/methods-2.1.1"
4   xmlns:proj="eml://ecoinformatics.org/project-2.1.1"
5   xmlns:d="eml://ecoinformatics.org/dataset-2.1.1"
6   xmlns:res="eml://ecoinformatics.org/resource-2.1.1"
7   xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
8   xmlns:dc="http://purl.org/dc/terms/"
9   packageId="e71fdalc-dcb9-4eae-81a9-183114978e44/eml-1.xml" system="GBIF-
  IPT" scope="system">
10 <dataset>
11   <alternateIdentifier>e71fdalc-dcb9-4eae-81a9-183114978e44</
12     alternateIdentifier>
13   <title>Images from Danau Girang during the PTY Project 2011-12</title>
14   <creator>
15     <individualName>
16       <givenName>Christopher</givenName>
17       <surName>Gwilliams</surName>
18     </individualName>
19     <organizationName>Cardiff University</organizationName>
20     <positionName>PhD</positionName>
21     <address>
22       <city>Cardiff</city>
23       <administrativeArea>Cardiff</administrativeArea>
24       <postalCode>CF24 3AA</postalCode>
25       <country>Wales</country>
26     </address>
27     <phone></phone>
28     <electronicMailAddress>C.Gwilliams@cs.cf.ac.uk</
29       electronicMailAddress>
30     <onlineUrl>christopher-gwilliams.com</onlineUrl>
31   </creator>
32   <pubDate>2012-07-26</pubDate>
33   <language>en</language>
34   <abstract>
35     <para></para>
36   </abstract>
37   <keywordSet>
38     <keyword>Danau Girang</keyword>
39     <keyword>Darwin Core</keyword>
40   </keywordSet>
41   <intellectualRights>
42     <para></para>
43   </intellectualRights>
44   <coverage>
45     <geographicCoverage>
46       <geographicDescription></geographicDescription>
47       <boundingCoordinates>
48         <westBoundingCoordinate>-180</westBoundingCoordinate>
49         <eastBoundingCoordinate>180</eastBoundingCoordinate>
50         <northBoundingCoordinate>90</northBoundingCoordinate>
51         <southBoundingCoordinate>-90</southBoundingCoordinate>
52       </boundingCoordinates>
53     </geographicCoverage>
54     <taxonomicCoverage>
55       <generalTaxonomicCoverage></generalTaxonomicCoverage>
56       <taxonomicClassification>
         <taxonRankName></taxonRankName>
         <taxonRankValue></taxonRankValue>

```

```

57           <commonName></commonName>
58       </taxonomicClassification>
59
60   </taxonomicCoverage>
61 </coverage>
62 <contact>
63   <individualName>
64     <surName>Gwilliams</surName>
65   </individualName>
66   <electronicMailAddress>c.gwilliams@cs.cf.ac.uk</
67     electronicMailAddress>
68 </contact>
69 <methods>
70 </methods>
71 <project>
72   <title>Using Local and Global Knowledge in Wireless Sensor
73     Networks</title>
74 </project>
75 </dataset>
76 </additionalMetadata>
77 </additionalMetadata>
78
79 </eml:eml>
```

Listing A.1: EML File

```

1 <?xml version="1.0" encoding='utf-8'?>
2 <archive xmlns="http://rs.tdwg.org/dwc/text/" metadata="eml.xml">
3   <core encoding="UTF-8" linesTerminatedBy="\n" fieldsTerminatedBy=","
4     fieldsEnclosedBy=''" ignoreHeaderLines="1" rowType="http://rs.tdwg.org/
5       dwc/terms/Occurrence">
6     <files>
7       <location>set.csv</location>
8     </files>
9     <id index="0"/>
10    <field index="0" term="http://rs.tdwg.org/dwc/terms/eventID"/>
11    <field index="1" term="http://rs.tdwg.org/dwc/terms/basisOfRecord"/>
12    <field index="2" term="http://rs.tdwg.org/dwc/terms/recordedBy"/>
13    <field index="3" term="http://rs.tdwg.org/dwc/terms/eventDate"/>
14    <field index="4" term="http://rs.tdwg.org/dwc/terms/locationID"/>
15    <field index="5" term="http://rs.tdwg.org/dwc/terms/scientificName"/>
16    <field index="6" term="http://rs.tdwg.org/dwc/terms/identifiedBy"/>
17    <field index="7" term="http://rs.tdwg.org/dwc/terms/dateIdentified"/>
18  </core>
19  <extension encoding="UTF-8" linesTerminatedBy="\n" fieldsTerminatedBy=","
20    fieldsEnclosedBy=''" ignoreHeaderLines="1" rowType="http://rs.gbif.org/
21      terms/1.0/Image">
22    <files>
23      <location>images.csv</location>
24    </files>
25  </extension>
26 </archive>
```

Listing A.2: Metadata File

```

1 eventID, basisOfRecord, recordedBy, eventDate, locationID, scientificName,
2   identifiedBy, dateIdentified
3   '1', 'MovingImage', '1', '2012-07-26 14:21', '1', 'Neofelis Nebulosa', '
4   Camera3', '2012-07-26 15:06'
```

Listing A.3: Example Set File (set.csv)

```

1 eventID, identifier
2 '1', '/Users/encima/Dropbox/Projects/PhD/Java/GSN_DGFC/webapp/IMG_0082.JPG'
3 '1', '/Users/encima/Dropbox/Projects/PhD/Java/GSN_DGFC/webapp/IMG_0083.JPG'
4 '1', '/Users/encima/Dropbox/Projects/PhD/Java/GSN_DGFC/webapp/IMG_0084.JPG'
```

Listing A.4: File Describing Image Locations (images.csv)

Appendix B

Darwin-SW Occurrence Example

```

1  <?xml version="1.0" encoding="UTF-8"?>
2  <rdf:RDF xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
3      xmlns:rdfs="http://www.w3.org/2000/01/rdf-schema#"
4      xmlns:dcterms="http://purl.org/dc/terms/"
5      xmlns:dwc="http://rs.tdwg.org/dwc/terms/"
6      xmlns:dsw="http://purl.org/dsw/"
7      xmlns:mrtg="http://rs.tdwg.org/mrtg/trunk/RDF/mrtg.n3#"
8      xmlns:xmp="http://ns.adobe.com/xap/1.0/"
9      xmlns:foaf="http://xmlns.com/foaf/0.1/"
10     >
11
12 <rdf:Description rdf:about="http://bioimages.vanderbilt.edu/vanderbilt
13 /12-126">
14     <rdf:type rdf:resource="http://purl.org/dsw/IndividualOrganism" />
15     <rdf:type rdf:resource="http://purl.org/dsw/LivingSpecimen" />
16     <mrtg:metadataLanguage>en</mrtg:metadataLanguage>
17
18     <!--Basic information about the individual organism-->
19     <dcterms:identifier>http://bioimages.vanderbilt.edu/vanderbilt/12-126<
20         /dcterms:identifier>
21     <dcterms:description>Field individual of Ginkgo biloba L. with GUID:
22         http://bioimages.vanderbilt.edu/vanderbilt/12-126</
23         dcterms:description>
24     <dsw:individualOrganismRemarks>This tree was miraculously spared from
25         destruction in the building of the 21st Ave. crosswalk.</
26         dsw:individualOrganismRemarks>
27
28     <!-- Properties that the tree has by virtue of its LivingSpecimen type
29         -->
30     <dwc:collectionID rdf:resource="http://biocol.org/urn:lsid:biocol.
31         org:col:35259" />
32     <dwc:collectionCode>vanderbilt</dwc:collectionCode>
33     <dwc:catalogNumber>12-126</dwc:catalogNumber>
34
35     <!-- Relationships of the individual to other resources. -->
36     <foaf:isPrimaryTopicOf rdf:resource="http://bioimages.vanderbilt.edu/
37         vanderbilt/12-126.rdf" />
38     <foaf:isPrimaryTopicOf rdf:resource="http://bioimages.vanderbilt.edu/
39         vanderbilt/12-126.htm" />
40
41     <!-- Images that are derived from the individual -->
42     <foaf:depiction rdf:resource="http://bioimages.vanderbilt.edu/baskauf
43         /10502"/>
44     <foaf:depiction rdf:resource="http://bioimages.vanderbilt.edu/baskauf
45         /10557"/>
46     <foaf:depiction rdf:resource="http://bioimages.vanderbilt.edu/baskauf
47         /10556"/>
48     <foaf:depiction rdf:resource="http://bioimages.vanderbilt.edu/baskauf
49         /10554"/>
50
51     <!-- Documented occurrences of the individual -->
52     <dsw:hasOccurrence rdf:resource="http://bioimages.vanderbilt.edu/
53         baskauf/10502#occ" />
54     <dsw:hasOccurrence rdf:resource="http://bioimages.vanderbilt.edu/
55         baskauf/10557#occ" />
56     <dsw:hasOccurrence rdf:resource="http://bioimages.vanderbilt.edu/
57         baskauf/10556#occ" />

```

```

42      <dsw:hasOccurrence rdf:resource="http://bioimages.vanderbilt.edu/
43                                baskauf/10554#occ" />
44
45      <!-- Identifications applied to the individual-->
46      <dsw:hasIdentification>
47          <rdf:Description rdf:about="http://bioimages.vanderbilt.edu/
48                                vanderbilt/12-126#2002-04-10baskauf">
49              <dcterms:description>Determination of Ginkgo biloba L. sensu
50                  Flora of North America (1993) for the individual http://
51                  bioimages.vanderbilt.edu/vanderbilt/12-126</
52                  dcterms:description>
53              <rdf:type rdf:resource="http://rs.tdwg.org/dwc/terms/
54                                Identification" />
55
56              <!-- In lieu of stable external identifiers for taxon concepts
57                  , Im defining some onsite -->
58              <dsw:toTaxon rdf:resource="http://bioimages.vanderbilt.edu/
59                                taxonConcepts#183269-fna1993" />
60
61          <dwc:identifiedBy>Steven J. Baskauf</dwc:identifiedBy>
62              <dsw:idBy rdf:resource="http://bioimages.vanderbilt.edu/
63                                contact/baskauf" />
64
65          <dwc:dateIdentified>2002-04-10</dwc:dateIdentified>
66
67          <!-- Relationship of the identification to other resources -->
68          <dsw:idBasedOn rdf:resource="http://bioimages.vanderbilt.edu/baskauf
69                                /10554"/>
70              </rdf:Description>
71          </dsw:hasIdentification>
72      </rdf:Description>
73
74      <!-- Information about the metadata document itself
75          -->
76      <rdf:Description rdf:about="http://bioimages.vanderbilt.edu/vanderbilt
77                                /12-126.rdf">
78          <dcterms:type rdf:resource ="http://xmlns.com/foaf/0.1/Document" />
79          <dcterms:identifier>http://bioimages.vanderbilt.edu/vanderbilt/12-126.
        rdf</dcterms:identifier>
          <dcterms:description>RDF formatted description of the living organism
          http://bioimages.vanderbilt.edu/vanderbilt/12-126</
          dcterms:description>
          <dcterms:creator>Bioimages http://bioimages.vanderbilt.edu/</
          dcterms:creator>
          <foaf:maker rdf:resource="http://biocol.org/urn:lsid:biocol.
            org:col:35115" />
          <dcterms:language>en</dcterms:language>
          <dcterms:modified>2010-11-08T11:13:49</dcterms:modified>
          <xmp:MetadataDate>2010-11-08T11:13:49</xmp:MetadataDate>
          <dcterms:references rdf:resource="http://bioimages.vanderbilt.edu/
            vanderbilt/12-126"/>
          <foaf:primaryTopic rdf:resource="http://bioimages.vanderbilt.edu/
            vanderbilt/12-126"/>
      </rdf:Description>
    </rdf:RDF>

```

Listing B.1: Darwin-SW Representation of a Living Specimen

Appendix C

K-HAS Ontology

```

1 <?xml version="1.0"?>
2
3
4 <!DOCTYPE Ontology [
5   <!ENTITY xsd "http://www.w3.org/2001/XMLSchema#" >
6   <!ENTITY xml "http://www.w3.org/XML/1998/namespace" >
7   <!ENTITY rdfs "http://www.w3.org/2000/01/rdf-schema#" >
8   <!ENTITY rdf "http://www.w3.org/1999/02/22-rdf-syntax-ns#" >
9 ]>
10
11
12 <Ontology xmlns="http://www.w3.org/2002/07/owl#"
13   xmlns:base="http://www.semanticweb.org/encima/ontologies/2013/4/alignment"
14   xmlns:rdfs="http://www.w3.org/2000/01/rdf-schema#"
15   xmlns:xsd="http://www.w3.org/2001/XMLSchema#"
16   xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#"
17   xmlns:xml="http://www.w3.org/XML/1998/namespace"
18   ontologyIRI="http://www.semanticweb.org/encima/ontologies/2013/4/
19     alignment">
20   <Prefix name="" IRI="http://www.w3.org/2002/07/owl#" />
21   <Prefix name="owl" IRI="http://www.w3.org/2002/07/owl#" />
22   <Prefix name="rdf" IRI="http://www.w3.org/1999/02/22-rdf-syntax-ns#" />
23   <Prefix name="xsd" IRI="http://www.w3.org/2001/XMLSchema#" />
24   <Prefix name="rdfs" IRI="http://www.w3.org/2000/01/rdf-schema#" />
25   <Import>http://purl.oclc.org/NET/ssnx/ssn</Import>
26   <Import>http://purl.org/dsw/</Import>
27   <Import>http://ecoinformatics.org/oboe/oboe.1.0/oboe.owl</Import>
28   <Import>http://www.semanticweb.org/encima/ontologies/2013/4/k_has</Import>
29   <Declaration>
30     <Class IRI="http://rs.tdwg.org/dwc/terms/DecimalLatitude"/>
31   </Declaration>
32   <Declaration>
33     <Class IRI="http://rs.tdwg.org/dwc/terms/DecimalLongitude"/>
34   </Declaration>
35   <Declaration>
36     <Class IRI="http://rs.tdwg.org/dwc/terms/LocationRemarks"/>
37   </Declaration>
38   <Declaration>
39     <ObjectProperty IRI="#hasLat"/>
40   </Declaration>
41   <Declaration>
42     <ObjectProperty IRI="#hasLng"/>
43   </Declaration>
44   <Declaration>
45     <ObjectProperty IRI="#hasRemark"/>
46   </Declaration>
47   <EquivalentClasses>
48     <Class IRI="http://ecoinformatics.org/oboe/oboe.1.0/oboe-core.owl#
49       ObservationCollection"/>
50     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k_has#
51       Project"/>
52   </EquivalentClasses>
53   <EquivalentClasses>
54     <Class IRI="http://purl.oclc.org/NET/ssnx/ssn#FeatureOfInterest"/>
55     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k_has#
      Identification"/>
56   </EquivalentClasses>
57   <EquivalentClasses>
58     <Class IRI="http://purl.oclc.org/NET/ssnx/ssn#Observation"/>

```

```

56     <Class IRI="http://rs.tdwg.org/dwc/terms/Occurrence" />
57   </EquivalentClasses>
58   <EquivalentClasses>
59     <Class IRI="http://purl.oclc.org/NET/ssnx/ssn#Sensor" />
60     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Node" />
61   </EquivalentClasses>
62   <EquivalentClasses>
63     <Class IRI="http://purl.org/dc/terms/Location" />
64     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Location" />
65   </EquivalentClasses>
66   <EquivalentClasses>
67     <Class IRI="http://purl.org/dsw/Token" />
68     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Token" />
69   </EquivalentClasses>
70   <EquivalentClasses>
71     <Class IRI="http://rs.tdwg.org/dwc/terms/Event" />
72     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Event" />
73   </EquivalentClasses>
74   <EquivalentClasses>
75     <Class IRI="http://rs.tdwg.org/dwc/terms/Identification" />
76     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Identification" />
77   </EquivalentClasses>
78   <EquivalentClasses>
79     <Class IRI="http://rs.tdwg.org/dwc/terms/Occurrence" />
80     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Occurrence" />
81   </EquivalentClasses>
82   <EquivalentClasses>
83     <Class IRI="http://www.loa-cnr.it/ontologies/DUL.owl#Agent" />
84     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Agent" />
85   </EquivalentClasses>
86   <EquivalentClasses>
87     <Class IRI="http://www.loa-cnr.it/ontologies/DUL.owl#Event" />
88     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Event" />
89   </EquivalentClasses>
90   <EquivalentClasses>
91     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Agent" />
92     <Class IRI="http://xmlns.com/foaf/0.1/Agent" />
93   </EquivalentClasses>
94   <SubClassOf>
95     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Location" />
96     <ObjectSomeValuesFrom>
97       <ObjectProperty IRI="#hasRemark" />
98       <Class IRI="http://rs.tdwg.org/dwc/terms/LocationRemarks" />
99     </ObjectSomeValuesFrom>
100   </SubClassOf>
101   <SubClassOf>
102     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Location" />
103     <ObjectAllValuesFrom>
104       <ObjectProperty IRI="#hasLat" />
105       <Class IRI="http://rs.tdwg.org/dwc/terms/DecimalLatitude" />
106     </ObjectAllValuesFrom>
107   </SubClassOf>
108   <SubClassOf>
109     <Class IRI="http://www.semanticweb.org/encima/ontologies/2013/4/k\_has#Location" />
110     <ObjectAllValuesFrom>
111       <ObjectProperty IRI="#hasLng" />
112       <Class IRI="http://rs.tdwg.org/dwc/terms/DecimalLongitude" />
113     </ObjectAllValuesFrom>
114   </SubClassOf>
115   <AnnotationAssertion>
116     <AnnotationProperty abbreviatedIRI="rdfs:comment" />
117     <IRI>http://rs.tdwg.org/dwc/terms/DecimalLatitude</IRI>
118     <Literal datatypeIRI="#rdf:PlainLiteral">dwc</Literal>
119   </AnnotationAssertion>
120   <AnnotationAssertion>
121     <AnnotationProperty abbreviatedIRI="rdfs:comment" />
122     <IRI>http://rs.tdwg.org/dwc/terms/DecimalLongitude</IRI>
123     <Literal datatypeIRI="#rdf:PlainLiteral">dwc</Literal>

```

```
124 | </Literal>
125 |   </AnnotationAssertion>
126 |   <AnnotationAssertion>
127 |     <AnnotationProperty abbreviatedIRI="rdfs:comment" />
128 |     <IRI>http://rs.tdwg.org/dwc/terms/LocationRemarks</IRI>
129 |     <Literal datatypeIRI="&rdf;PlainLiteral">dwc</Literal>
130 |   </AnnotationAssertion>
131 | </Ontology>
132 |
133 |
134 |
135 | <!-- Generated by the OWL API (version 3.4.2) http://owlapi.sourceforge.net -->
```

Listing C.1: K-HAS Ontology Source Code

Appendix D

Example Drools Rules Used in LORIS

```

1 rule "Otter River"
2   when
3     $dwc:DarwinCore()
4       DarwinCore(DBTools.getLocationDescription($dwc.getOcc().getLocationID())
5         contains "River")
6       eval(TimeTools.checkTimePeriod($dwc.getOcc().getEventTime()) == "Afternoon
7         ")
8     then
9       System.out.println("Otter");
10      int id = $dwc.getOcc().getEventID();
11      Date d = new Date();
12      Identification ident = new Identification(id, 1, d, 1);
13      $dwc.setId(ident);
14      System.out.println($dwc.getId());
15    end
16
17 rule "Plantation Night"
18   when
19     $dwc:DarwinCore()
20       DarwinCore(DBTools.getLocationDescription($dwc.getOcc().getLocationID())
21         contains "Plantation")
22       eval(TimeTools.checkTimePeriod($dwc.getOcc().getEventTime()) == "Night")
23     then
24       System.out.println("Sun Bear or Civet");
25     end
26
27 rule "Clouded Leopard Ridgeline"
28   when
29     $dwc:DarwinCore()
30       DarwinCore(DBTools.getLocationDescription($dwc.getOcc().getLocationID())
31         contains "Ridge Line")
32       eval(TimeTools.checkTimePeriod($dwc.getOcc().getEventTime()) == "Night")
33     then
34       System.out.println("Clouded Leopard - Most likely male");
35
36 rule "Location Description"
37   when
38     $dwc:DarwinCore()
39     eval($dwc.getOcc().getLocationID() == 1)
40   then
41     System.out.println("Occurrence in site 1, animals here are: Clouded Leopard
42       , Malay Civets and Small Mammals");
43   end

```

Listing D.1: Example Drools File

Appendix E

Interview Transcript: Gill Bolongan

1 Interviewer: Chris Gwilliams
 2 Interviewee: Roshan
 3
 4 Interviewer: Can you state your name please?
 5 Interviewee: My name is Roshan Guharajan
 6 Interviewer: The species you are looking at is?
 7 Interviewee: Sun bear.
 8 Interviewer: What is your project title?
 9 Interviewee: Sun bear (*Helarctos malayanus*) ecology within a fragmented landscape: a case study from the lower kinabatangan wildlife riparian (land between river) ecosystem.
 10 Interviewer: What kind of collection or analysis do you do? Trapping, physical samples etc.
 11 Interviewee: So, initially camera traps, now we are going into trapping. Once that is done: tracking, scat collection and feeding sign surveys.
 12 Interviewer: Scat is faeces?
 13 Interviewee: Yes, sorry, faeces.
 14 Interviewer: What made you move on from cameras to physical trapping?
 15 Interviewee: The cameras were more of a preliminary investigation into what areas they were frequenting.
 16 Interviewer: How did you choose the sites for the cameras?
 17 Interviewee: They were already placed by DG.
 18 Interviewer: So, you used the existing network.
 19 Interviewee: Yes, the existing camera locations.
 20 Interviewer: So, you used the existing locations, looked for pictures of Sun Bears and placed your physical traps where those pictures were taken?
 21 Interviewee: Yes. Close to. We also used cameras facing the traps to monitor what animals are coming to the trap.
 22 Interviewer: So, the specific sites you look into have been chosen by the cameras?
 23 Interviewee: Yes.
 24 Interviewer: Have you looked into any sites that weren't selected by the cameras?
 25 Interviewee: Yes we have, in lot 6, we have two sites where (after finding a lot of secondary signs), we placed two traps.
 26 Interviewer: Are secondary signs things like marking on a tree?
 27 Interviewee: Predominantly claw marks.
 28 Interviewer: So, a primary sign would be the bear itself?
 29 Interviewee: Yes, seeing it.
 30 Interviewer: What times do you work on your project? Whether that is trapping or looking for signs.
 31 Interviewee: Particularly in the morning, or just before lunch. If we check traps, we check them twice a day. So, once in the morning and once in the late afternoon.
 32 Interviewer: Why do you have to check twice a day?
 33 Interviewee: At this point, we have trap monitors that are basically VHF frequency and emit when the door closes, but there are no elevation points high enough around here to pick up the signals from all of our sites. So, we have to physically be at the trap to see if it is closed or not.
 34 Interviewer: Is that if you only check once a day then the bear has less chance of surviving?
 35 Interviewee: Actually, once a day is more desirable because it means less human disturbance. But, because of the logistics here, we have to check twice; unless we can find a high enough point.
 36 Interviewer: So, would you say your project is time critical then? If you catch a bear, you need to get to it as soon as?
 37 Interviewee: Yes, that would be ideal.
 38 Interviewer: Is there anyway you could put some antennae higher in the trees?

39 Interviewee: There could be, yes. We tried going up the tower, but we could
 not get any signal. Especially in Batangan, as we have one trap there.
 40 Interviewer: Is it that the forest is too thick?
 41 Interviewee: It may be that, or the canopy, and just the distance as well.
 42 Interviewer: How far away are they?
 43 Interviewee: Batangan, by river, is at least 11km.
 44 Interviewer: Is there a reason you choose the morning?
 45 Interviewee: The bears seem to be nocturnal, so best to check in the morning.
 46 Interviewer: Are there any conditions you have learnt, or from other research,
 that are optimum? Like: weather conditions, the temperature or the moon
 phase. Are there any times that you know, on, say, a particularly bright
 night, so more likely to find a Sun Bear. Or if it is particularly humid.
 47 Interviewee: To be honest, I am not sure. I know periods of low food
 availability, but that is more of a seasonal thing.
 48 Interviewer: So they would be more enticed to go into your trap?
 49 Interviewee: Yep, more readily available food.
 50 Interviewer: Do you do anything different if the weather is suitable/not
 suitable?
 51 Interviewee: We try and go earlier if it is raining, especially at the end of
 the year when it is the monsoon.
 52 Interviewer: Does the height of the river seem to affect them at all?
 53 Interviewee: I don't know but from last year to this year, we have not had as
 many pictures on camera traps as usual. There are a few locations, like bi
 weekly, or monthly, they would pass by the camera. During the rain, they
 have, kind of, slowed down. I don't know if that is affecting them; maybe
 they move slower or the flooding made them go further inland.
 54 Interviewer: So, flooding may affect them?
 55 Interviewee: Yep, definitely.
 56 Interviewer: Are they fairly permanent in their sites then? Like, if you
 captured a bear with a camera trap, would you say they stay in that area?
 Or, do they kind of move around?
 57 Interviewee: I don't know, I know they have pretty large home ranges.
 58 Interviewer: How large, do you know?
 59 Interviewee: A study using radio telemetry found a minimum of 15 square
 kilometers. We think that is a minimum, so probably double that, 30 to 40.
 60 Interviewer: So, if I said to you, list off the specifics that you know about
 your species, i.e. certain times it does not come out, locations it may
 avoid, certain animals it may avoid, or maybe certain temperatures it does
 not come out at.
 61 Interviewee: The bears here seem to be more nocturnal, it might be because, if
 they do encroach into plantations, night time is the time to do it. This
 is just speculation, but day time is very hard with human disturbance.
 62 Interviewer: They do go into plantation then?
 63 Interviewee: There have been reports from plantation workers seeing bears on
 plantation land.
 64 Interviewer: So, because there are fewer humans?
 65 Interviewee: Fewer humans, and it is a lot cooler during the night.
 66 Interviewer: Are there any areas that, if you saw a Sun Bear, it would be out
 of place? Or is that this is all forest, so this is all suitable for them?
 67 Interviewee: The latter, I think.
 68 Interviewer: There is not really anywhere they would not go?
 69 Interviewee: Yep.
 70 Interviewer: Do they have any predators?
 71 Interviewee: There was a study in Indonesia that found that large pythons do
 prey on small bears. But, I think the biggest danger comes from other
 bears, like big males killing younger ones and fighting.
 72 Interviewer: I guess pythons and things are less common.
 73 Interviewee: Yes, extremely rare.
 74 Interviewer: How long is the duration of your project?
 75 Interviewee: I started last September (2012) and I am projected to finish in
 May 2015.
 76 Interviewer: Essentially a 3 year project?
 77 Interviewee: Yes.
 78 Interviewer: That is all within DG?
 79 Interviewee: Not exactly, there are about 18 months outside of DG.
 80 Interviewer: Using research from DG?
 81 Interviewee: Yes, using all research from DG.
 82 Interviewer: OK, have/will/do others use your research? Currently, or in the
 future? Any active collaborations?
 83 Interviewee: I think we definitely will collaborate with the Sabah Wildlife
 Department in printing some action plans, recommendations. As far as
 collaborations, there may be collaborations with some local universities.
 For genetics.
 84 Interviewer: You would share your samples with them?
 85 Interviewee: Yes. Scat samples and blood samples.
 86 Interviewer: Would you expect anything in return?
 87 Interviewee: Probably just use of their labs.
 88 Interviewer: OK, so we have discussed how time critical your project is. Is
 there a reason that you want to get the bear out as quick as possible?

89 Interviewee: The animal use protocol is keep the animal in the trap for as little time as possible.
 90 Interviewer: So, just less disturbance for the animal? What do you do once you have the animal?
 91 Interviewee: Yes. Once we have the animal, we proceed to immobilise it.
 92 Interviewer: Is that darting?
 93 Interviewee: It is with a jab stick, just like a pole with a syringe on the end. It should go down within 5-10 minutes, possibly longer. After that, the most important thing is the collaring. Then some physical measurements, samples, blood, tissue, parasites, fecal and then we let the bear recover and release it when it is fully able to walk.
 94 Interviewer: Your collars are GPS collars, are they?
 95 Interviewee: GPS through the Iridium network.
 96 Interviewer: Is that the service that you log on to access the data?
 97 Interviewee: Yes.
 98 Interviewer: Are they the same network used by the crocodile project?
 99 Interviewee: Yes, the same.
 100 Interviewer: Final question, what other knowledge (if any) have you used in your project. For example, in this project, you have used the existing cameras and the pictures from there, have you used anything else to help? Or someone else's research outside of the field centre to help with your understanding of the bears or even your site choosings?
 101 Interviewee: Another student here, Grace, has been doing secondary sign surveys, and I have been talking to her about moving traps around based on where she has found fresh Sun Bear sign.
 102 Interviewer: So, the two sites that you found, where you selected the secondary signs, were they found by Grace?
 103 Interviewee: She was not doing her project at the time. While we were hanging baits on trees, in random places, we came across this area where we saw fresh signs.
 104 Interviewer: Have you moved any of your traps to her locations?
 105 Interviewee: Not at the moment, we are still deciding.
 106 Interviewer: Do you think you may do in the future?
 107 Interviewee: We may do.
 108 Interviewer: And has she used any of your research, do you know?
 109 Interviewee: At the moment, no she has not.
 110 Interviewer: Does she also use the pictures from within DG?
 111 Interviewee: Not for research.
 112 Interviewer: Ok, thanks, that is everything.

Listing E.1: Sample Interview Transcript

Appendix F

Qualitative Analysis Extract

animals are coming to the trap.

Interviewer: So, the specific sites you look into have been chosen by the cameras?

Interviewee: Yes.

Interviewer: Have you looked into any sites that weren't selected by the cameras?

Interviewee: Yes we have, in lot 6, we have two sites where (after finding a lot of secondary signs), we placed two traps.

Interviewer: Are secondary signs things like marking on a tree?

Interviewee: Predominantly claw marks.

Interviewer: So, a primary sign would be the bear itself?

Interviewee: Yes, seeing it.

Interviewer: What time was trapping or looking for signs. Unknown 1 Jan 1, 00:00

Interviewee: Particularly them twice a day. So, Project Info Local
e check traps, we check
in the afternoon.

Interviewer: Why do you think it's important to have a strong relationship with your clients?

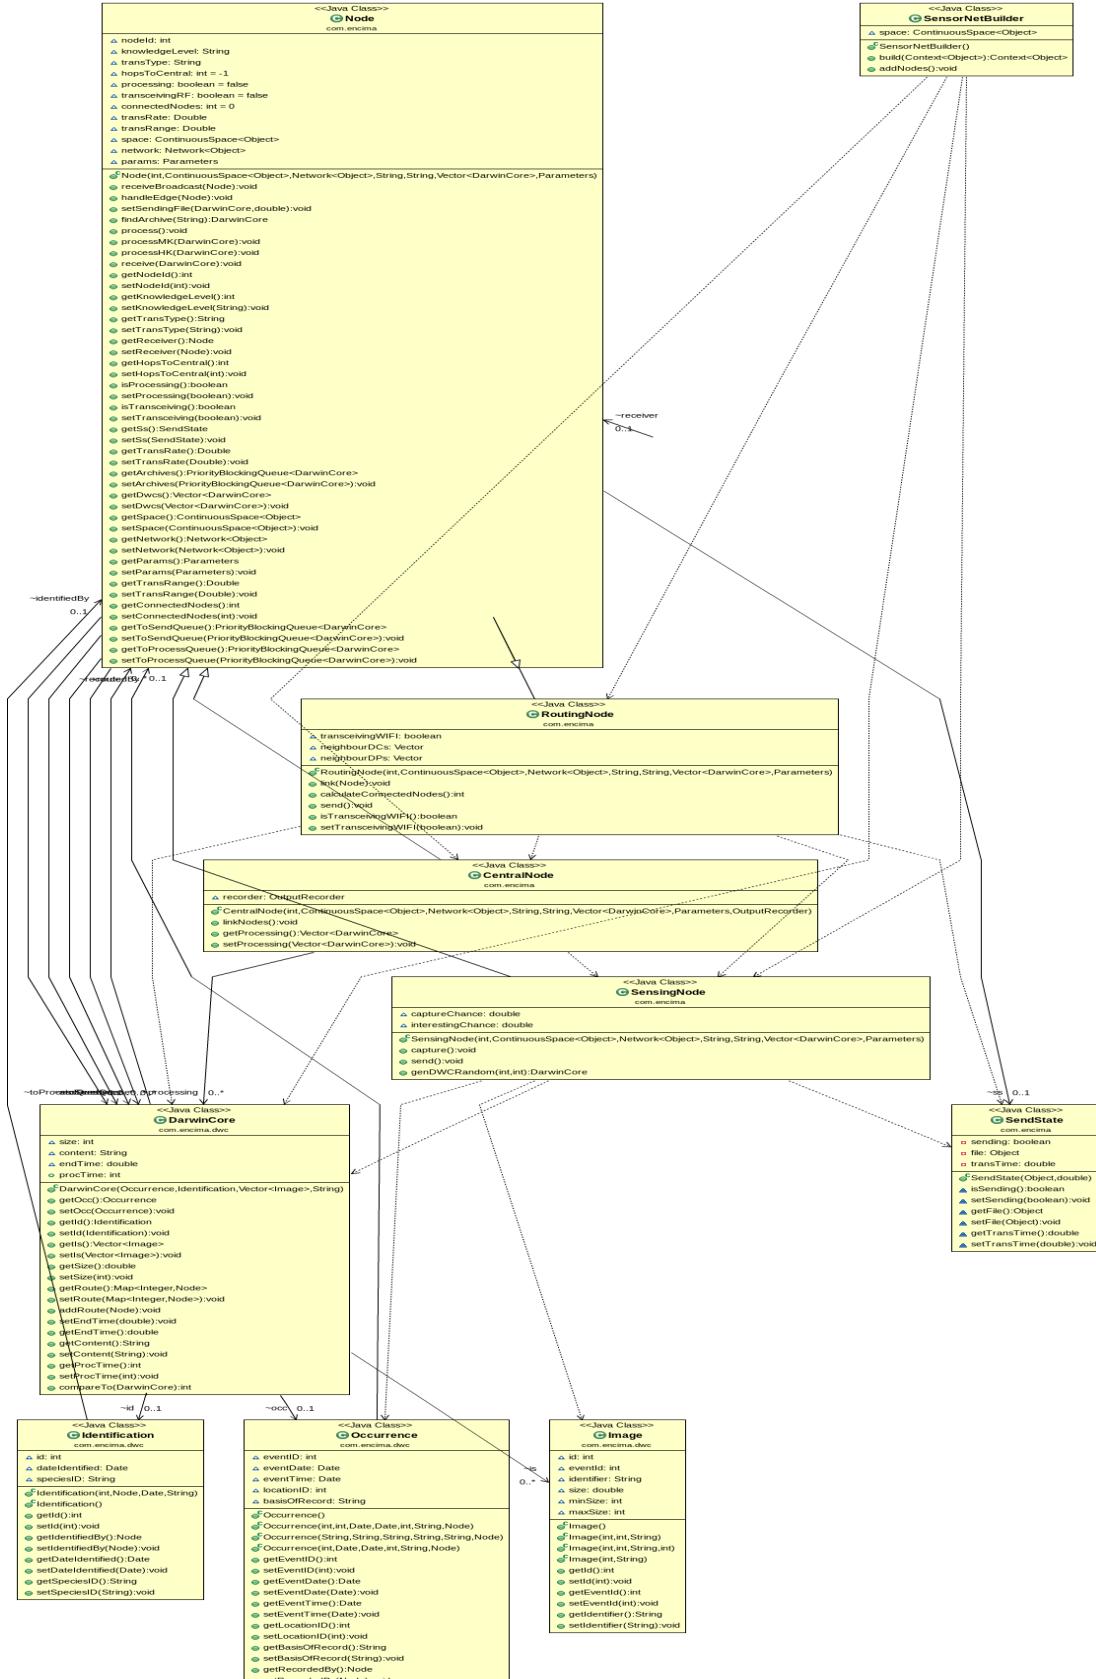
Interviewee: At this point, we have trap monitors that are basically VHF frequency and emit when the door closes, but there are no elevation points high enough around here to pick up the signals from all of our sites. So, we have to physically be at the trap to see if it is closed or not.

Interviewer: Is that if you only check once a day then the bear has less chance of surviving?

Interviewee: Actually, once a day is more desirable because it means less human disturbance. But, because of the logistics here, we have to check twice; unless we can find a high enough point.

Appendix G

Class Diagram for Repast Simulation



Appendix H

Repast Parameters File Example

```

1 <?xml version="1.0" encoding="UTF-8" ?>
2 <parameters>
3   <parameter name="randomSeed" displayName="Default_Random_Seed" type="int"
4     defaultValue="__NULL__" />
5   <parameter name="sensingNodeCount" displayName="Number_of_sensing_Nodes"
6     type="int" defaultValue="20" isReadOnly="false" converter="repast.
7       simphony.parameter.StringConverterFactory$IntConverter" />
8   <parameter name="routingNodeCount" displayName="Number_of_routing_Nodes"
9     type="int" defaultValue="4" isReadOnly="false" converter="repast.
10      simphony.parameter.StringConverterFactory$IntConverter" />
11   <parameter name="centralNodeCount" displayName="Number_of_central_Nodes"
12     type="int" defaultValue="1" isReadOnly="false" converter="repast.
13      simphony.parameter.StringConverterFactory$IntConverter" />
14   <parameter name="sensingKnowledge" displayName="Knowledge_for_sensing_nodes"
15     type="java.lang.String" defaultValue="MK" isReadOnly="false" converter=
16       "repast.simphony.parameter.StringConverterFactory$StringStringConverter"
17     />
18   <parameter name="routingKnowledge" displayName="Knowledge_for_routing_nodes"
19     type="java.lang.String" defaultValue="MK" isReadOnly="false" converter=
20       "repast.simphony.parameter.StringConverterFactory$StringStringConverter"
21     />
22   <parameter name="centralKnowledge" displayName="Knowledge_for_central_nodes"
23     type="java.lang.String" defaultValue="NK" isReadOnly="false" converter=
24       "repast.simphony.parameter.StringConverterFactory$StringStringConverter"
25     />
26   <parameter name="captureChance" displayName="chance_of_image_capture" type="
27     double" defaultValue="0.000857703189" isReadOnly="false" converter="
28       repast.simphony.parameter.StringConverterFactory$DoubleConverter"/>
29   <parameter name="interestingChance" displayName="chance_of_interesting_
30     capture" type="double" defaultValue="0.207" isReadOnly="false" converter
31       ="repast.simphony.parameter.StringConverterFactory$DoubleConverter"/>
32   <parameter name="transRate" displayName="Transmission_rate_(variable_or_
33     ideal)" type="java.lang.String" defaultValue="variable" isReadOnly="
34     false" converter="repast.simphony.parameter.StringConverterFactory$"
35     StringStringConverter" />
36   <parameter name="bandwidth" displayName="Chance_of_image_capture_(normal_or_
37     saturated)" type="java.lang.String" defaultValue="normal" isReadOnly="
38     false" converter="repast.simphony.parameter.StringConverterFactory$"
39     StringStringConverter" />
40   <parameter name="hkProcessingTime" displayName="time_to_process_archive"
41     type="int" defaultValue="43" isReadOnly="false" converter="repast.
42       simphony.parameter.StringConverterFactory$IntConverter"/>
43   <parameter name="mkProcessingTime" displayName="time_to_process_archive"
44     type="int" defaultValue="5" isReadOnly="false" converter="repast.
45       simphony.parameter.StringConverterFactory$IntConverter"/>
46 
```

```
29   <parameter name="tickStop" displayName="Stop_at_Tick_(-1_for_No)" type="int"
30     defaultValue="15552000" isReadOnly="false" converter="repast.simphony.
31       parameter.StringConverterFactory$IntConverter"/>
32   <parameter name="height" displayName="Height" type="int" defaultValue="1200"
33     isReadOnly="false" converter="repast.simphony.parameter.
34       StringConverterFactory$IntConverter" />
34 </parameters>
```

Listing H.1: Example params.xml file for a Repast simulation